

# Computer Vision

## Lec 2 – local invariant features

University of Haifa

Simon Korman

detection, description and matching  
of local invariant features

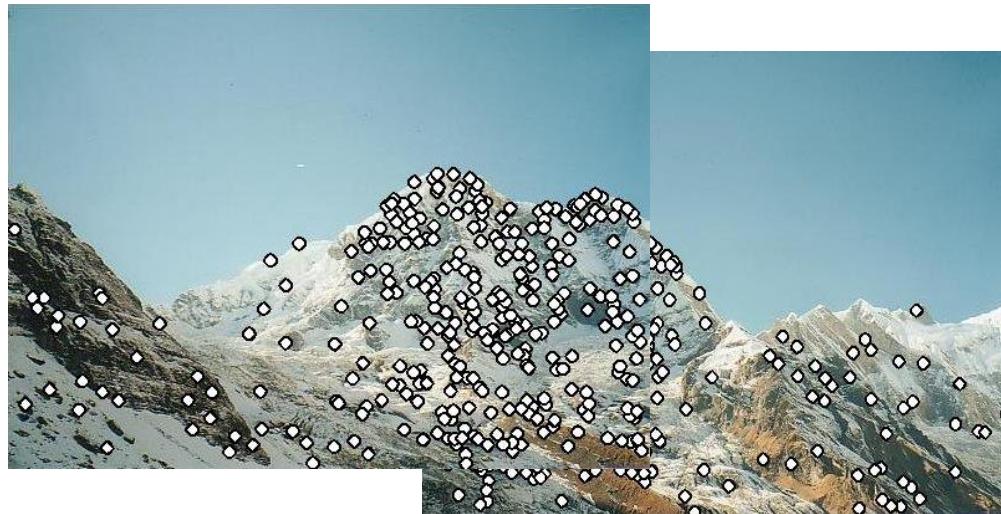
# slide credit

- Svetlana Lazebnik, Sanja Fidler, Kristen Grauman, Ioannis Gkioulekas, Kris Kitani, James Hays, Fredo Durand, Rick Szeliski, Andrew Zisserman, Kyros Kutulakos

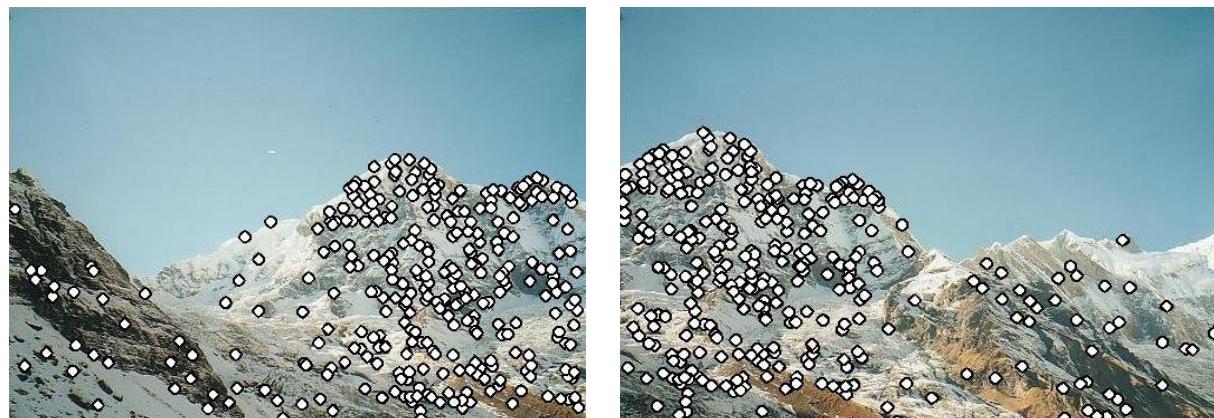
# Local invariant features



# Important tool for multiple views: Local features



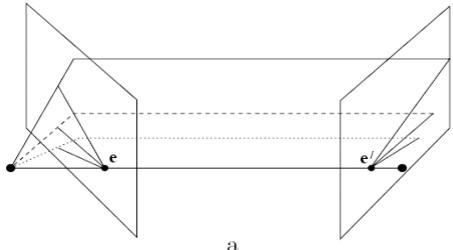
Multi-view matching relies on **local feature** correspondences.



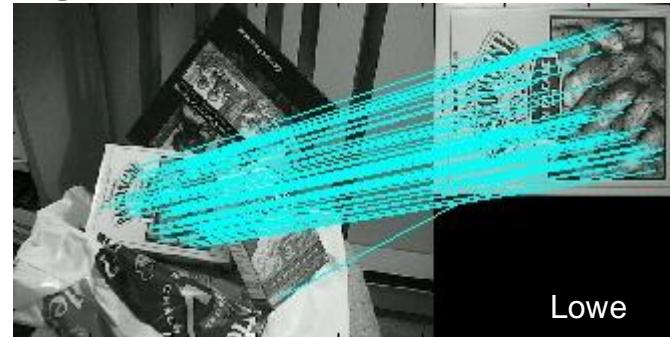
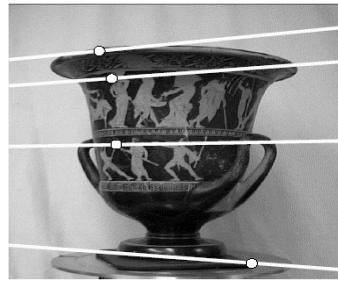
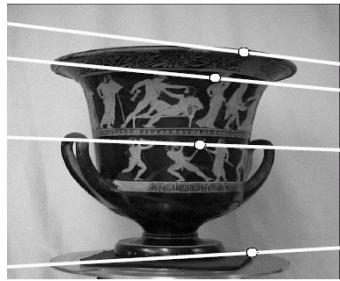
How to detect *which local features* to match?

# Multiple views

Matching, invariant features,  
stereo vision, instance  
recognition



a



Lowe



Fei-Fei Li



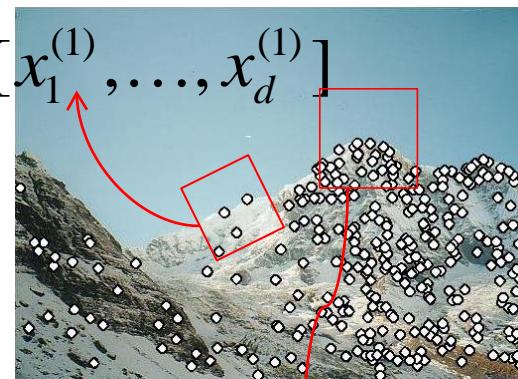
# Local features: main components

- 1) Detection: Identify the interest points



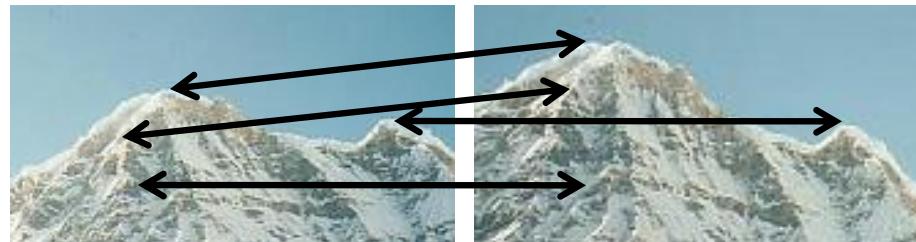
- 2) Description: Extract vector feature descriptor surrounding each interest point.

$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$



- 3) Matching: Determine correspondence between descriptors in two views

$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$



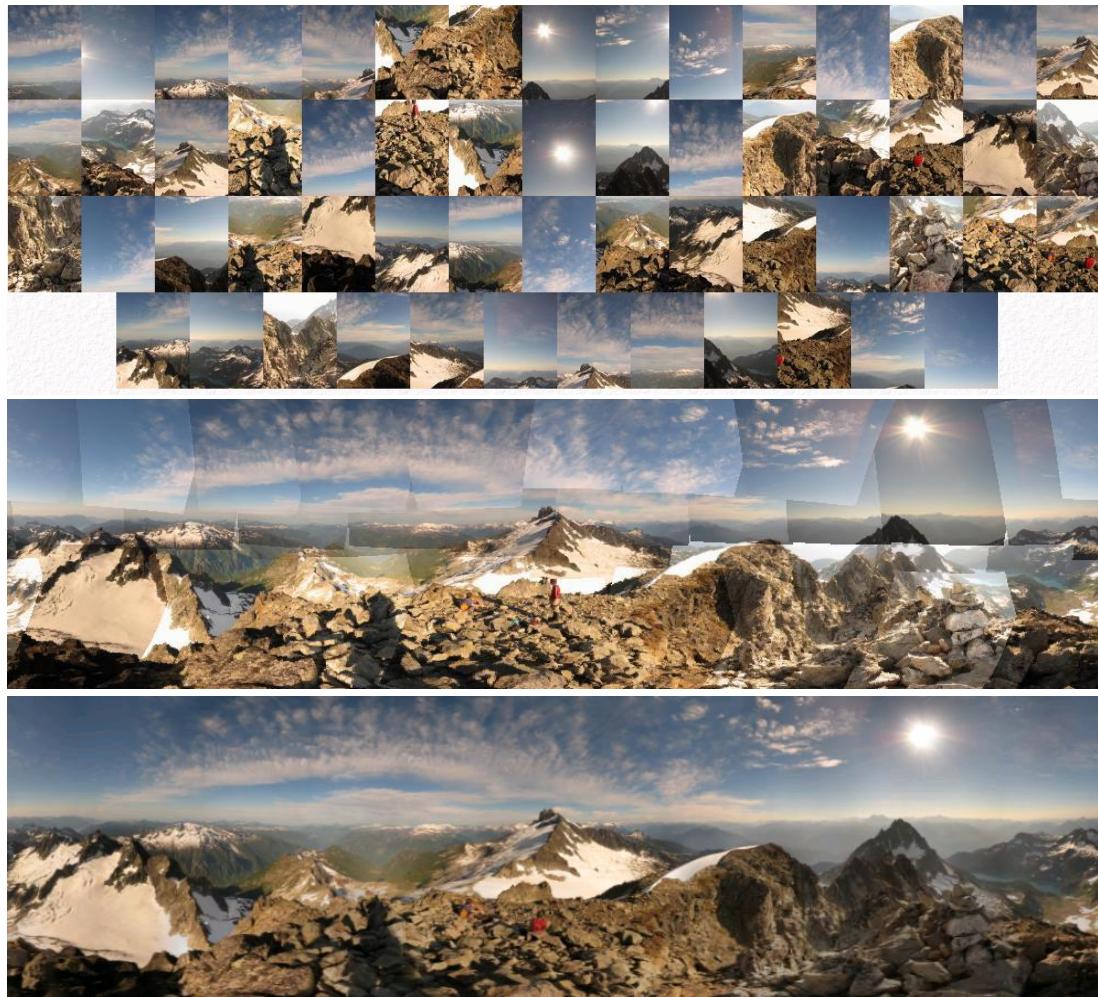
# Value of local (invariant) features

- Complexity reduction via selection of distinctive points
- Describe images, objects, parts without requiring segmentation
  - Local character means robustness to clutter, occlusion
- Robustness: similar descriptors in spite of noise, blur, etc.

# Applications of local invariant features

- Wide baseline stereo
- Motion tracking
- Panoramas
- Mobile robot navigation
- 3D reconstruction
- Recognition
- ...

# Automatic mosaicing



Matthew Brown

<http://matthewwalunbrown.com/autostitch/autostitch.html>

# Wide baseline stereo



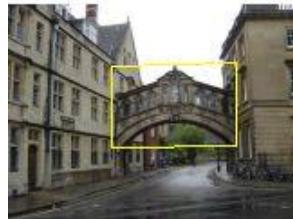
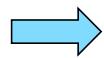
[Image from T. Tuytelaars ECCV 2006 tutorial]

# Photo tourism [Snavely et al. 2006/8]

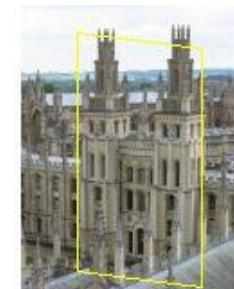
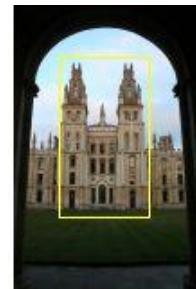
---



# Recognition of specific objects, scenes



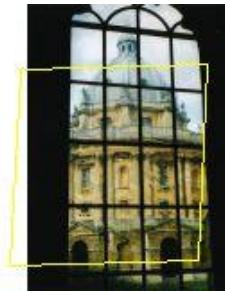
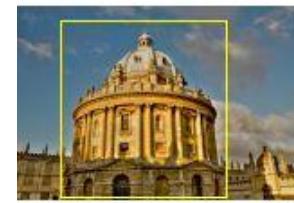
Scale



Viewpoint



Lighting



Occlusion

# Google Goggles

---



# Local features: desired properties

- Repeatability
  - The same feature can be found in several images despite geometric and photometric transformations
- Saliency
  - Each feature has a distinctive description
- Compactness and efficiency
  - Many fewer features than image pixels
- Locality
  - A feature occupies a relatively small area of the image; robust to clutter and occlusion

# Goal: interest operator repeatability

- We want to detect (at least some of) the same points in both images.

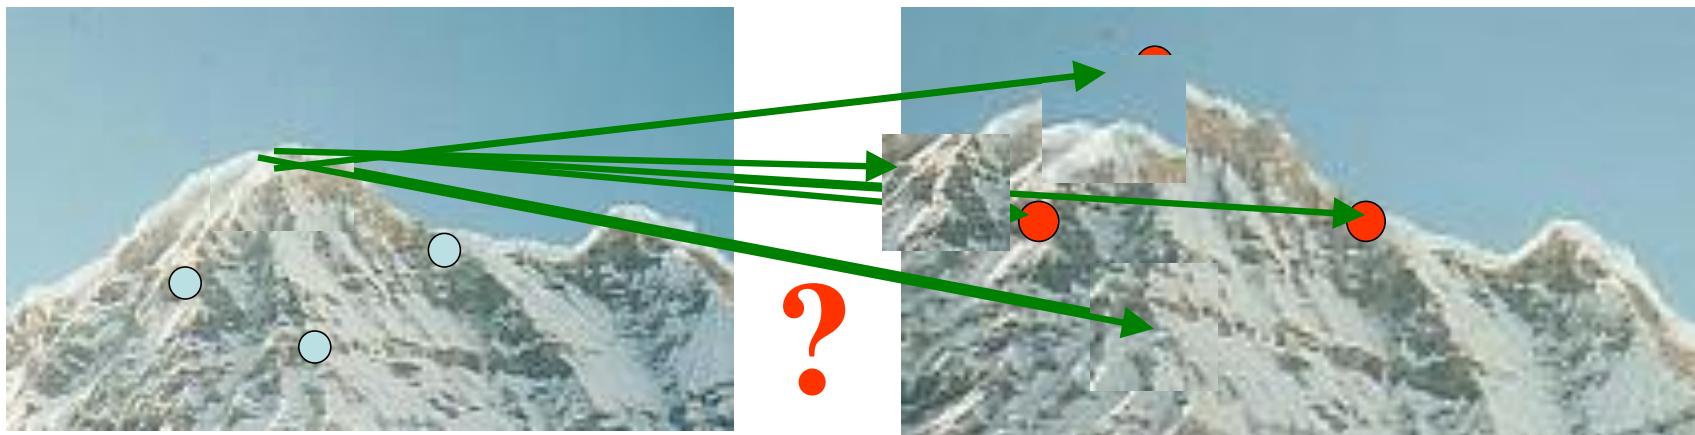


No chance to find (all and only) true matches!

- Yet we have to be able to run the detection procedure *independently* per image.

# Goal: descriptor distinctiveness

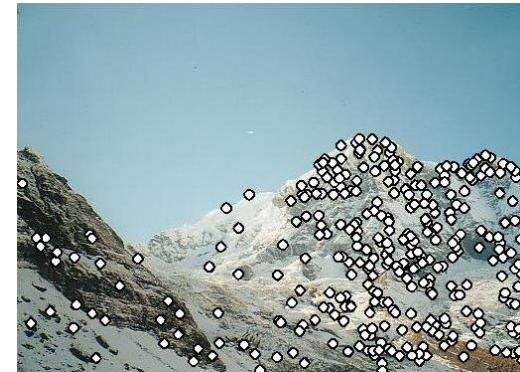
- We want to be able to reliably determine which point goes with which.



- Must provide some invariance to geometric and photometric differences between the two views.

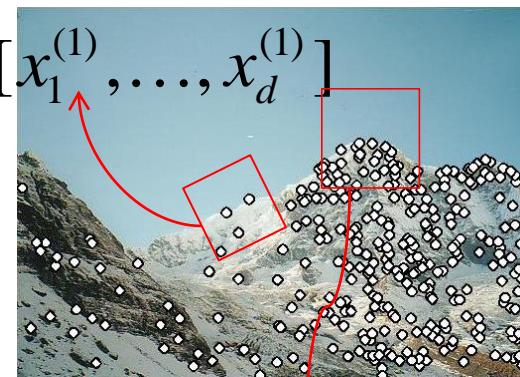
# Local features: main components

- 1) Detection: Identify the interest points



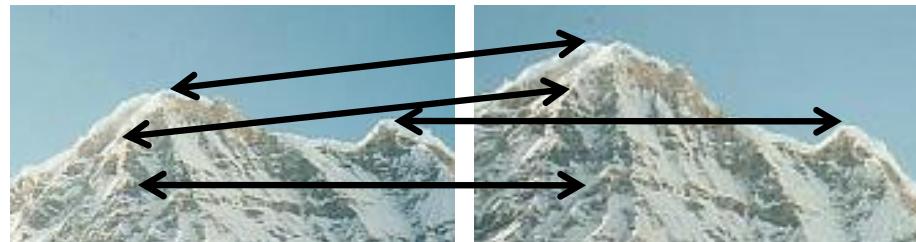
- 2) Description: Extract vector feature descriptor surrounding each interest point.

$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$

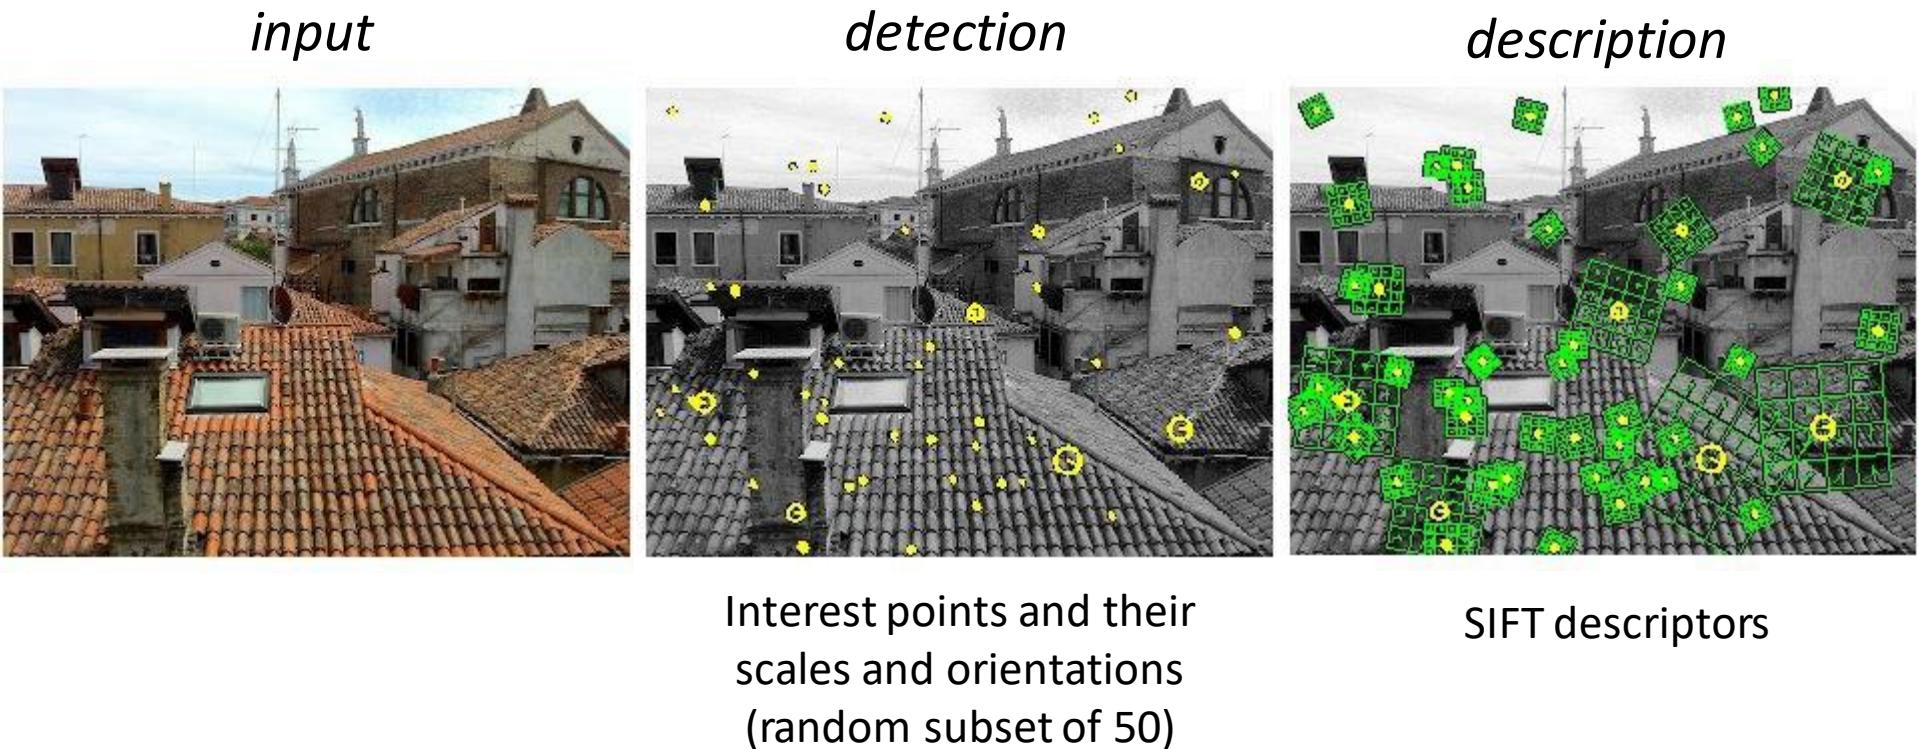


- 3) Matching: Determine correspondence between descriptors in two views

$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$



# Scale Invariant Feature Transform (SIFT) descriptor [Lowe 2004]



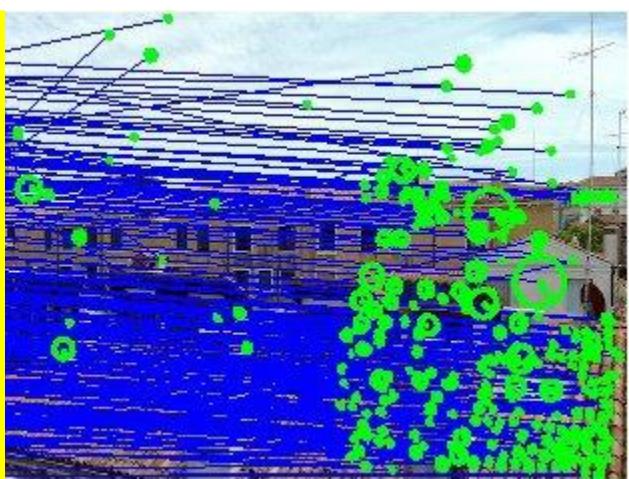
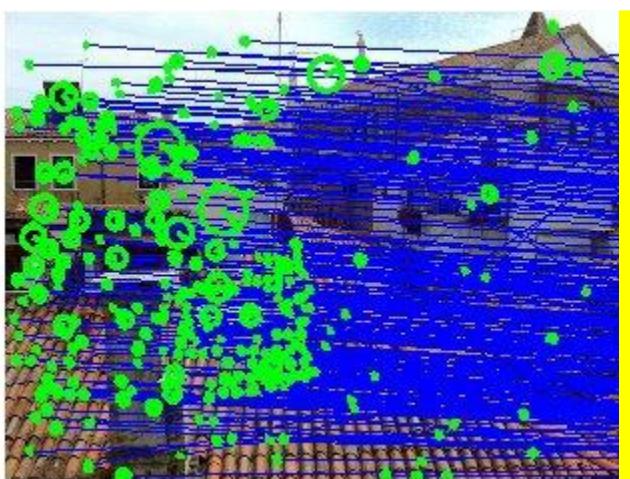
# SIFT (preliminary) matches

img1



img2

img1



img2

# SIFT descriptor [Lowe 2004]

- Extraordinarily robust matching technique
  - Can handle changes in viewpoint
    - Up to about 60 degree out of plane rotation
  - Can handle significant changes in illumination
    - Sometimes even day vs. night (below)
  - Fast and efficient—can run in real time
  - Lots of code available, e.g. <http://www.vlfeat.org/overview/sift.html>



# SIFT properties

- Invariant to
  - Scale
  - Rotation
- Partially invariant to
  - Illumination changes
  - Camera viewpoint
  - Occlusion, clutter
- ASIFT: fully affine invariant image comparison (SIAM 2009)

# Scale Invariant Feature Transform (SIFT) descriptor [Lowe 2004]

- introduction to the image matching problem
- image matching using SIFT features
- multi-scale interest-point detection
- the SIFT feature detector
- the SIFT descriptor

# Scale Invariant Feature Transform (SIFT) descriptor [Lowe 2004]

- introduction to the image matching problem
- image matching using SIFT features
- multi-scale interest-point detection
- the SIFT feature detector
- the SIFT descriptor

# Scale Invariant Feature Transform (SIFT) descriptor [Lowe 2004]

- **introduction to the image matching problem**
- image matching using SIFT features
- multi-scale interest-point detection
- the SIFT feature detector
- the SIFT descriptor

# The Image Matching Problem

Image I



Image I'



Image I''

Goal: Identify "features" or patches in image I that appear in another image, I''



# The Image Matching Problem



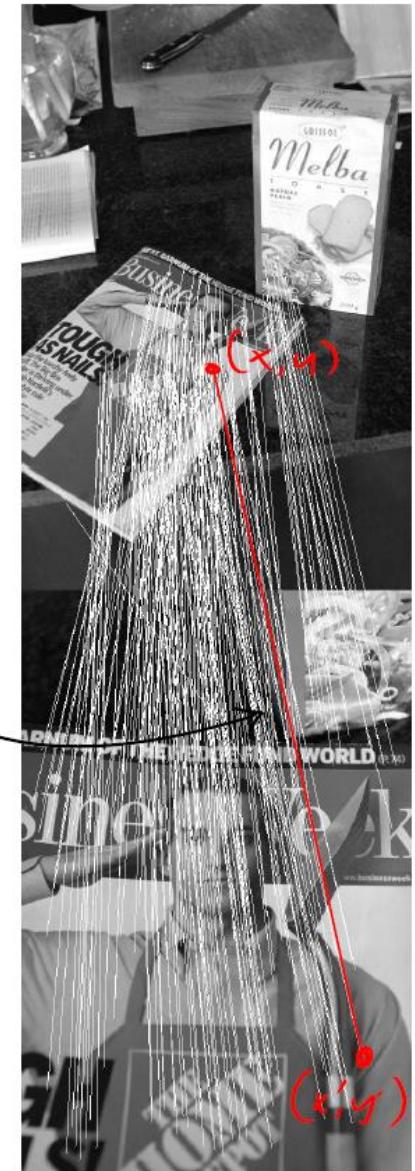
Image I



Image I'

Image I'

Indicates a correspondence between location  $(x,y)$  in image I and location  $(x',y')$  in image I'



# The Image Matching Problem

Image I



Image II



Q: Is it possible  
to solve this  
problem by  
direct template  
matching between  
the two images?



# The Image Matching Problem

Image 1



Patch P

Image 2



Corresponding patch P'

Q: Is it possible  
to solve this  
problem by  
direct template  
matching between  
the two images?



# The Image Matching Problem

Image 1



Image 2



Q: Is it possible  
to solve this  
problem by  
direct template  
matching between  
the two images?

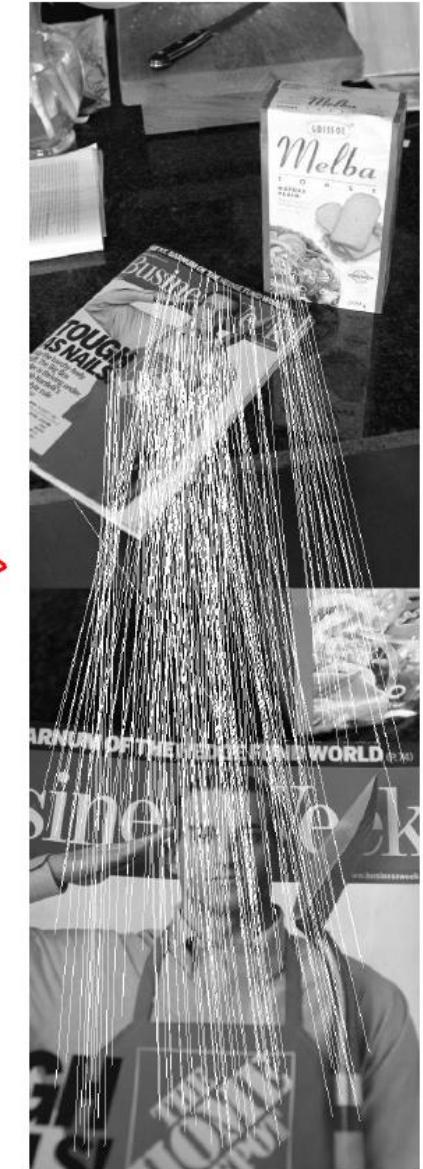
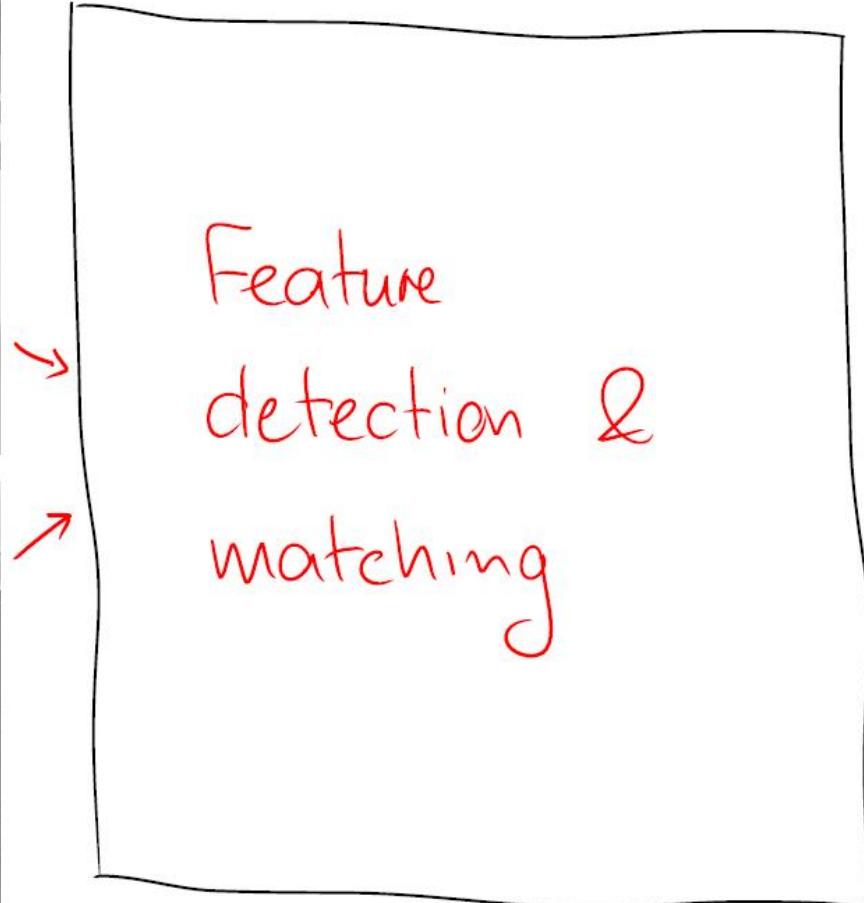
Ans: Yes, but it would be  
impossibly inefficient  
(i.e. must search over all  
possible pairs of patches)

# Feature-Based Image Matching

Image 1



Image 2

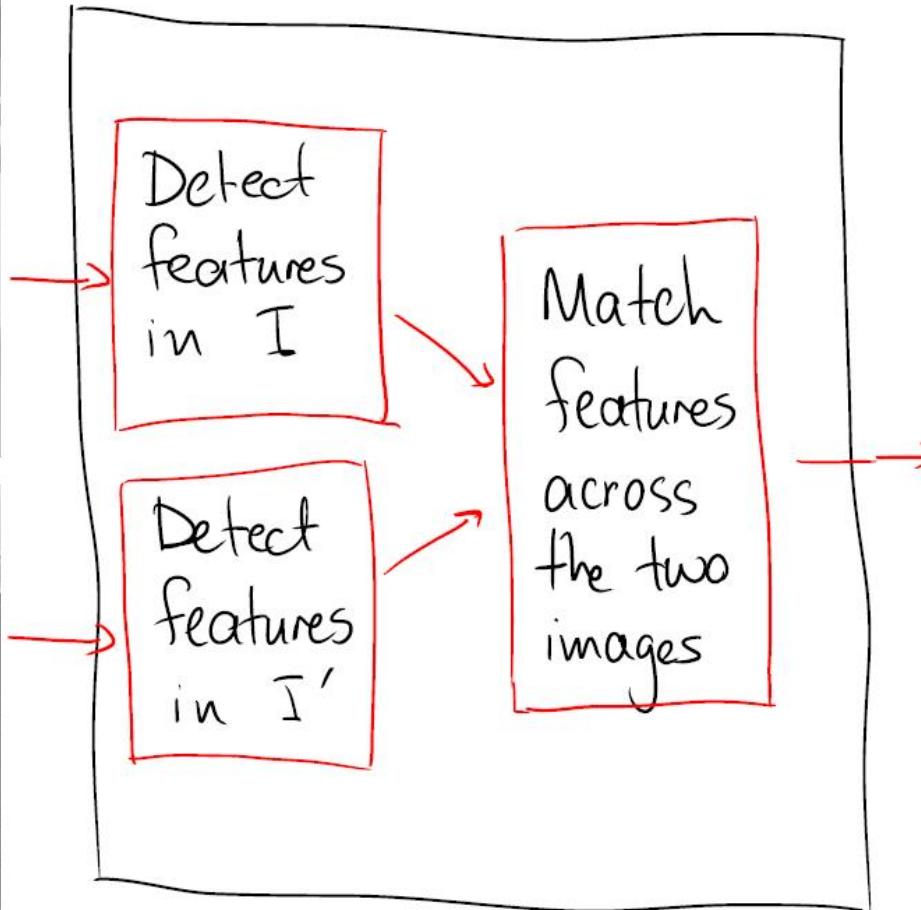


# Feature-Based Image Matching

Image I



Image I'



# Errors in Feature-Based Image Matching

Image 1



Image 2



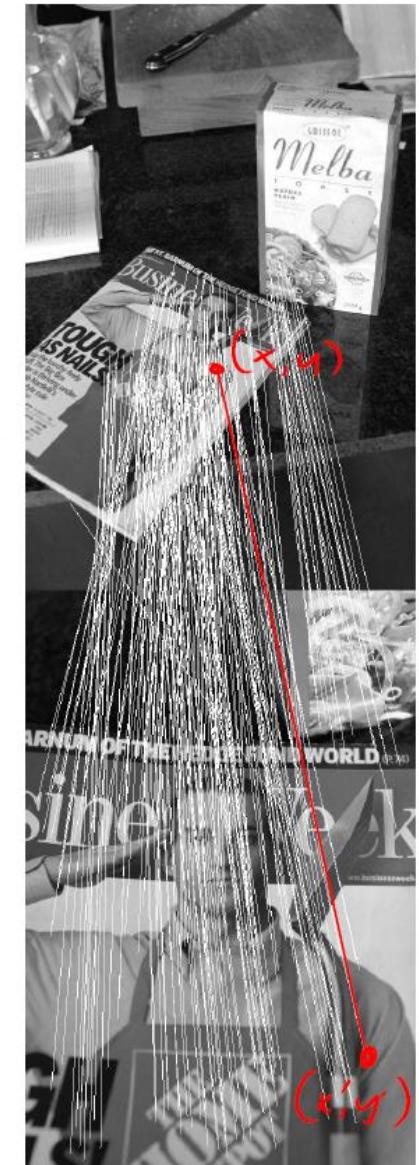
- In general some/many of these correspondences may be incorrect
- Two types of error:

## ① False positive matches

algorithm returns a correspondence between 2 locations where none exists

## ② False negative matches

algorithm fails to detect a correspondence between two instances of the same



# Errors in Feature-Based Image Matching

1

Image

2

Image



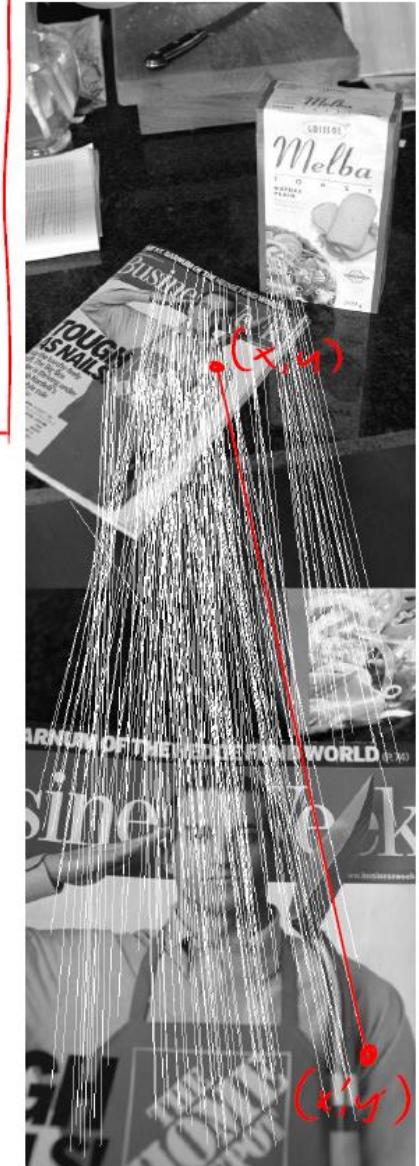
Goal: Minimize false positives  
AND false negatives  
across a wide range of  
imaging conditions

## ① False positive matches

algorithm returns a correspondence between 2 locations where none exists

## ② False negative matches

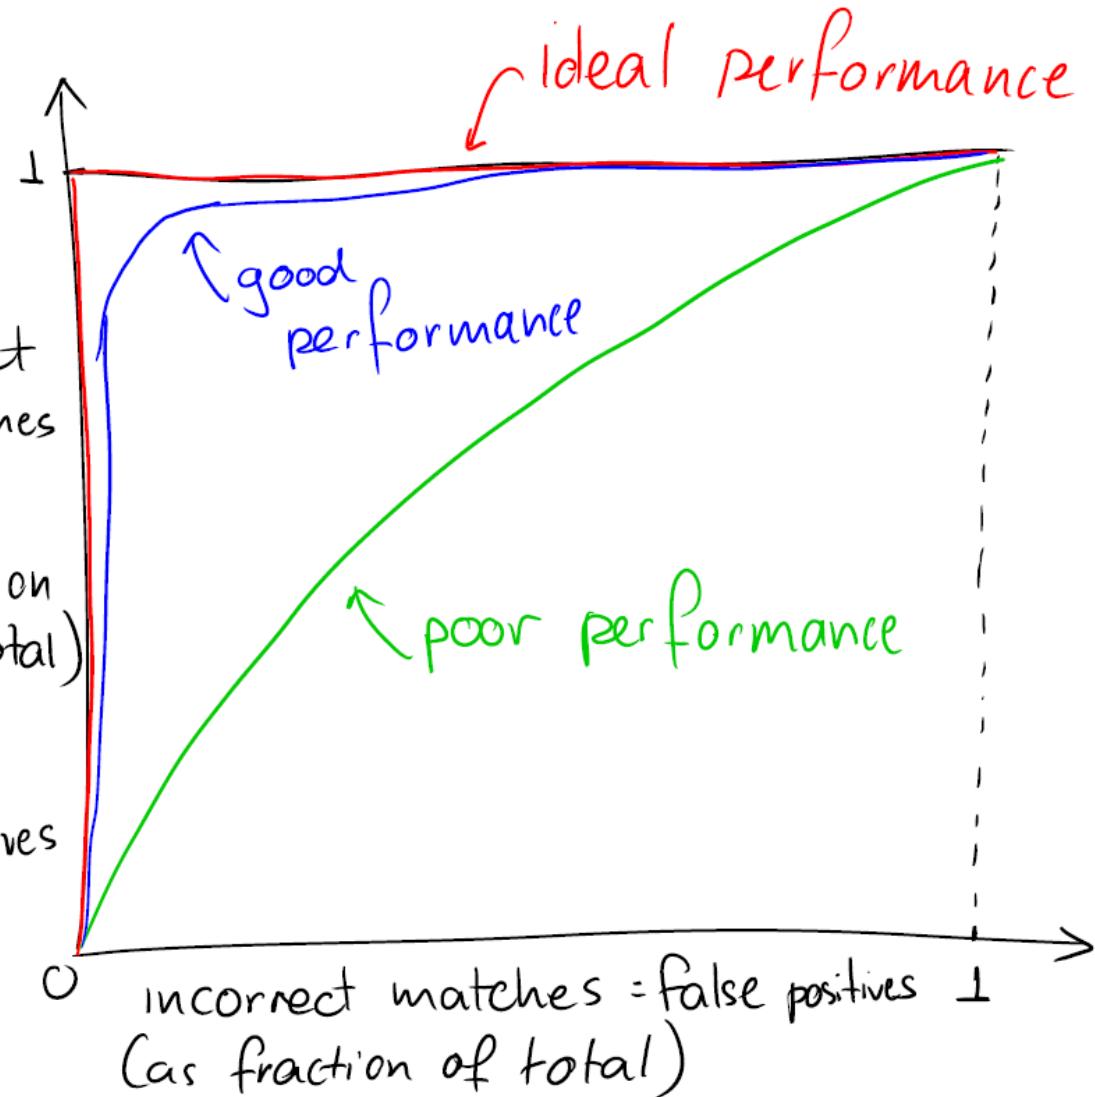
algorithm fails to detect a correspondence between two instances of the same feature/patch



# Evaluating a Feature Detector's Performance



connect  
matches  
(as  
fraction  
of total)  
=  
true  
positives



# Feature Matching & Transformation Invariance

Source image I



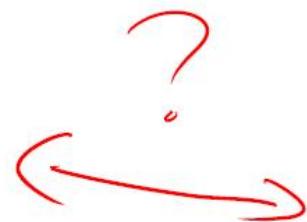
"Transformed" source images



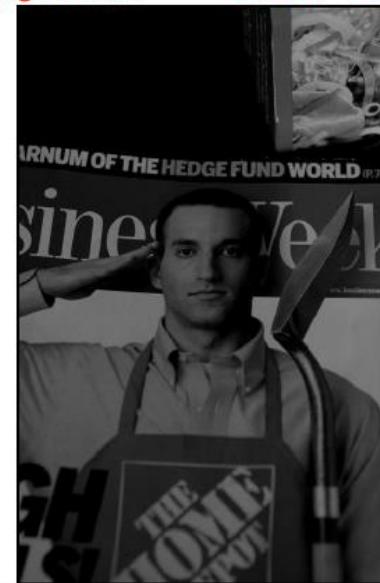
To be most useful, the feature detector & matching algorithm must be insensitive to a wide range of image transformations

# Transformation-Invariant Feature Detectors

Source image I



"Transformed" source images



A feature detector is called invariant to a certain image transformation if it can reliably detect features in a transformed version of the source image

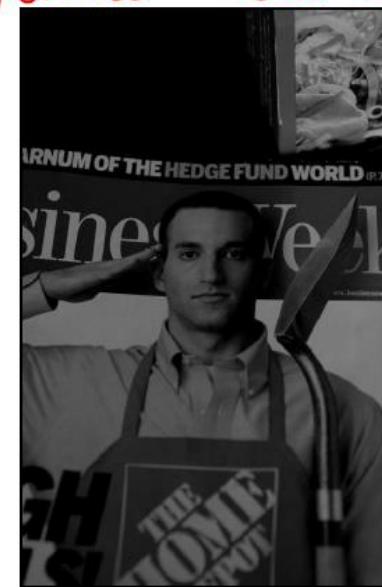


# Transformation-Invariant Feature Detectors

Source image I



"Transformed" source images



A feature detector  
is called invariant  
to a certain image  
transformation if  
it can reliably  
detect features  
in a transformed  
version of the source image

Brightness transformation

# Transformation-Invariant Feature Detectors

Source image I



"Transformed" source images



Distortion due  
to change in  
viewpoint

A feature detector  
is called invariant  
to a certain image  
transformation if  
it can reliably  
detect features  
in a transformed  
version of the source image



# Transformation-Invariant Feature Detectors

Source image I



"Transformed" source images



A feature detector is called invariant to a certain image transformation if it can reliably detect features in a transformed version of the source image

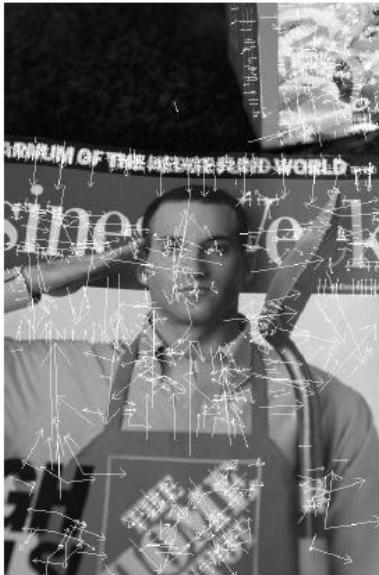


Distortion due to change in viewpoint & magnification (ie. scale)

# Scale Invariant Feature Transform (SIFT) descriptor [Lowe 2004]

- introduction to the image matching problem
- **image matching using SIFT features**
- multi-scale interest-point detection
- the SIFT feature detector
- the SIFT descriptor

# Image Matching Using SIFT Features



Scale +  
viewpoint

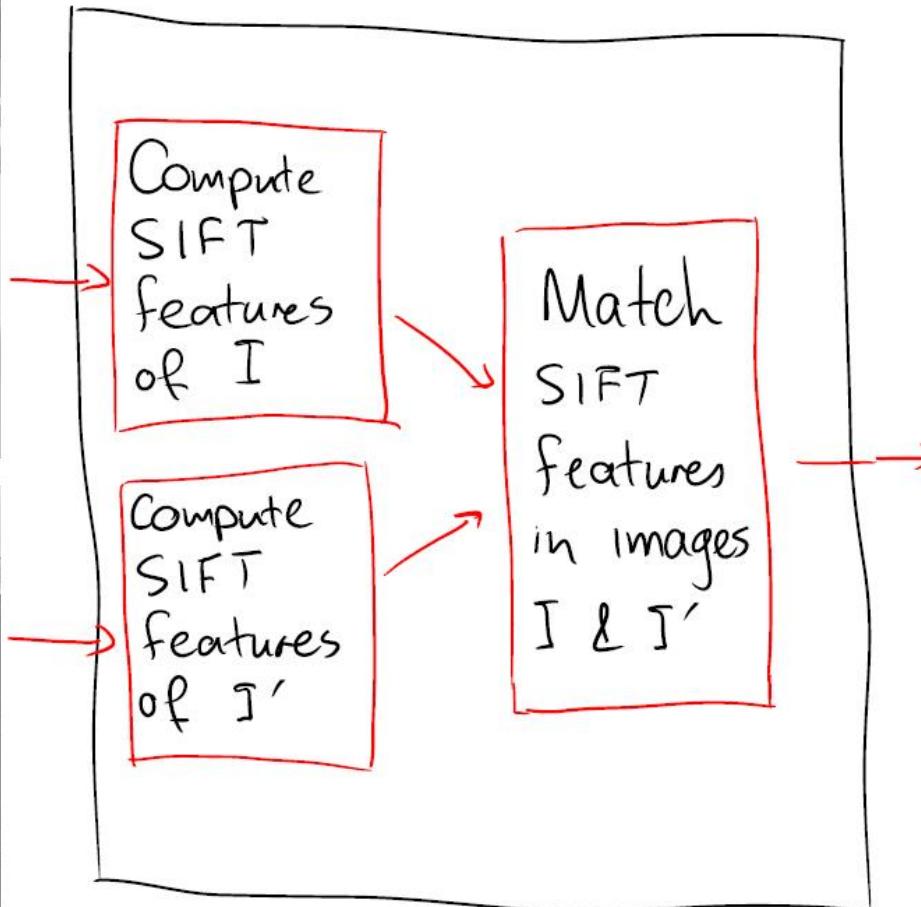


# Image Matching Using SIFT Features

Image I



Image I'



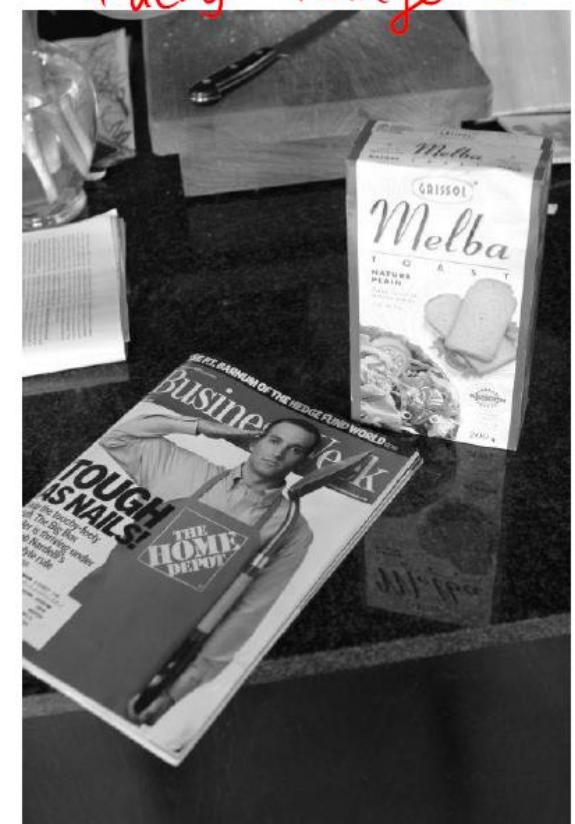
# The SIFT Feature Detection Algorithm

Goal: Represent an image  $I$  as a collection of SIFT features that can be identified reliably in other images where the same (or similar) objects are present

"Source" image  $I$



"Query" image  $I'$



# The SIFT Feature Detection Algorithm

Goal: Represent an image  $I$  as a collection of SIFT features that can be identified reliably in other images where the same (or similar) objects are present

Input:

Image  $I$

Output:

A set of  $k$  SIFT  
keypoints

$\{P_1, P_2, \dots, P_k\}$  &

feature vectors

$\{f_1, f_2, \dots, f_k\}$



# The SIFT Keypoints

**Keypoint:** A location  $(x, y)$  in the source image, with an associated orientation & scale, that is "visually distinct" from its surroundings

Input:

Image  $I$

Output:

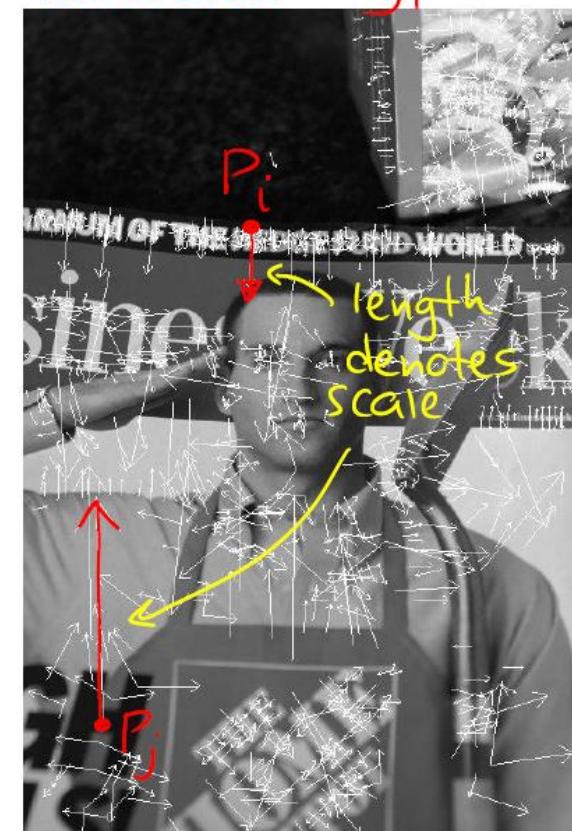
A set of  $\kappa$  SIFT  
keypoints

$$P_i = (x_i, y_i, g_i, \theta_i)$$

location  
scale  
orientation



Detected keypoints



# The SIFT Feature Vectors

Feature vector (of a keypoint  $P_i$ ): A vector of fixed length that represents the image patch centered at  $P_i$ .

Input:

Image  $I$

Output:

A set of  $K$   
keypoints

$$P_i = (x_i, y_i, g_i, \theta_i)$$

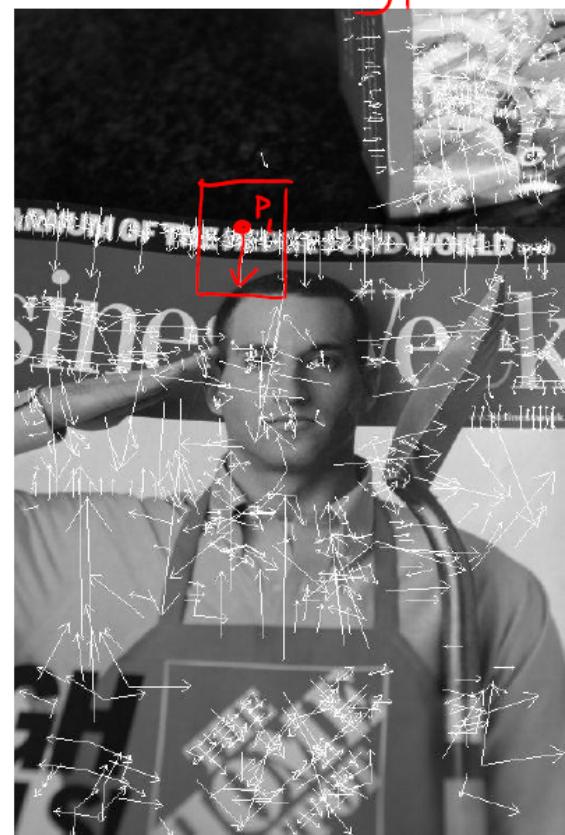
location  
scale  
orientation

A set of  $K$   
feature vectors

$$f_i = \underbrace{[ \quad | \quad \dots | \quad ]}_{\text{describes the patch centered at } P_i}$$

describes the patch centered at  $P_i$

Detected keypoints



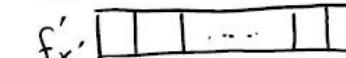
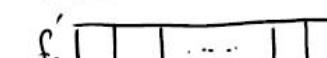
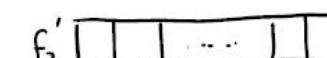
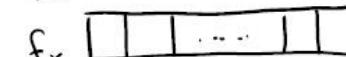
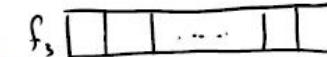
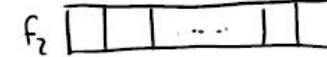
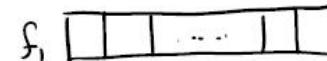
# Representing an Image Using SIFT Features: Steps

---

Source Image I



# Matching 2 Images Using SIFT Features



①

Identify keypoints

② Build feature  
vectors

# Matching 2 Images Using SIFT Features

Image 1

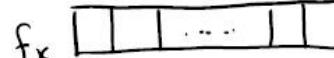
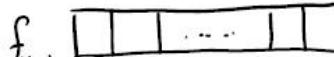
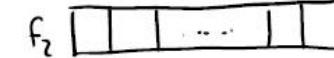
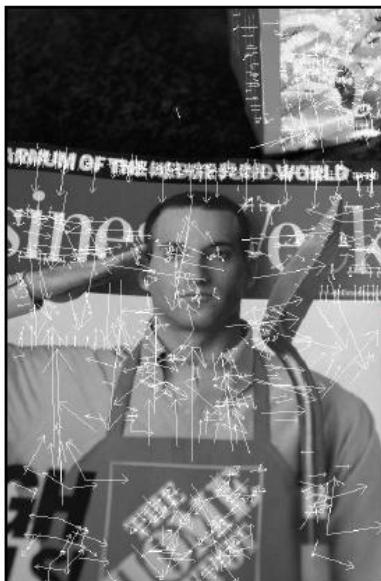
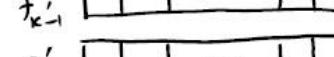
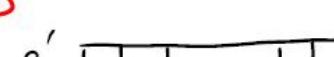
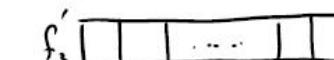
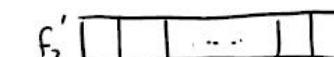


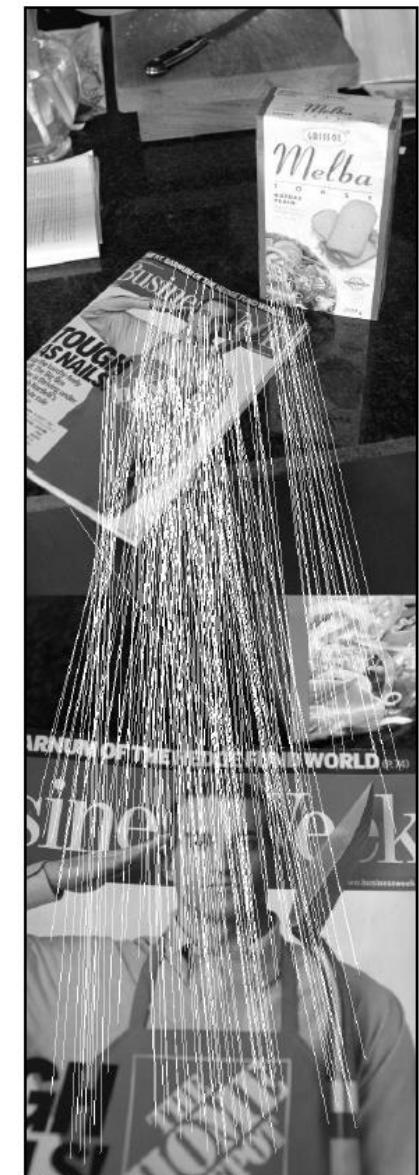
Image 2



① Identify keypoints

② Build feature vectors

③ Match vectors



# Scale Invariant Feature Transform (SIFT) descriptor [Lowe 2004]

- introduction to the image matching problem
- image matching using SIFT features
- **multi-scale interest-point detection**
- the SIFT feature detector
- the SIFT descriptor

# Detecting corners

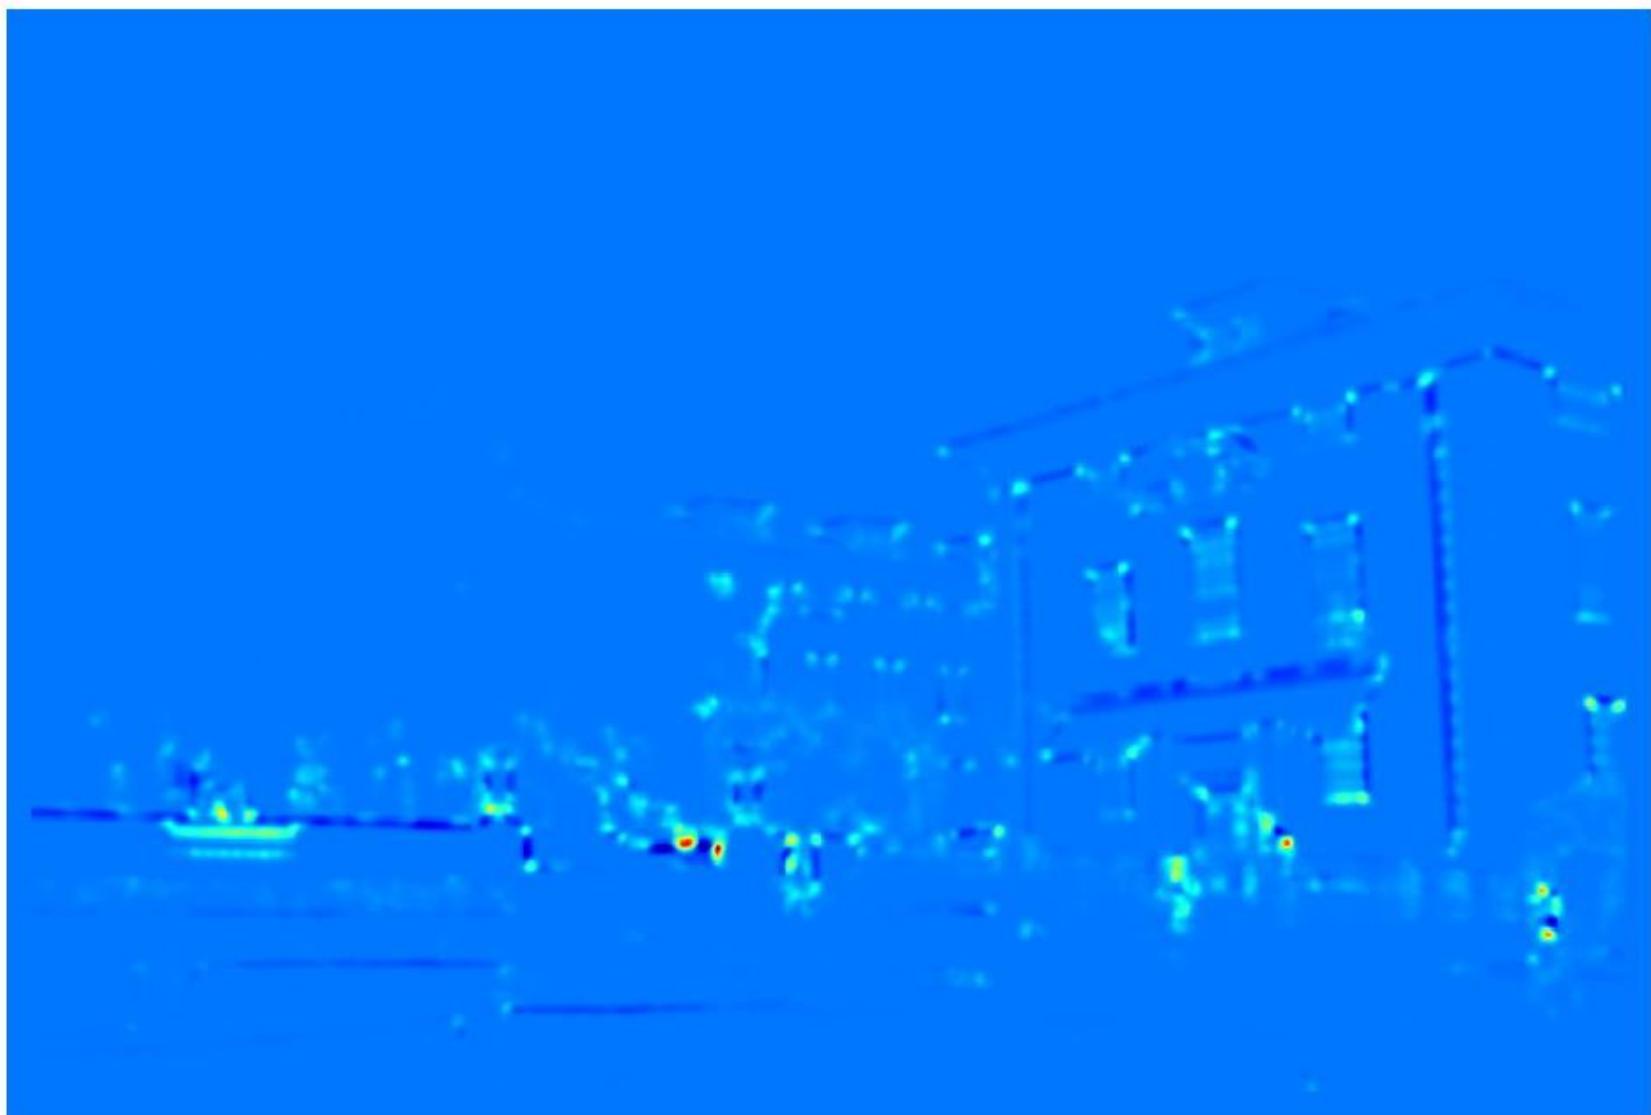
---



Slide credit: Kristen Grauman

# Detecting corners

---

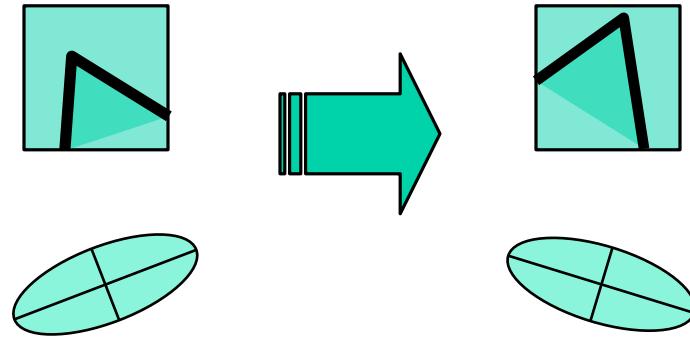


# Detecting corners



Slide credit: Kristen Grauman

# Harris corner response is invariant to rotation



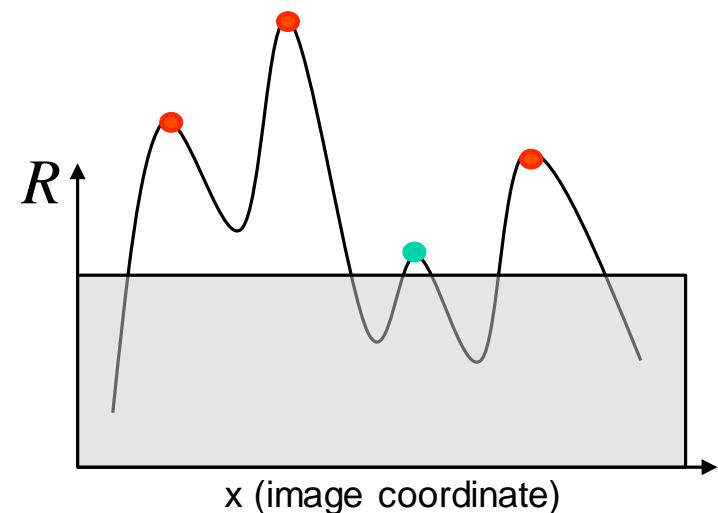
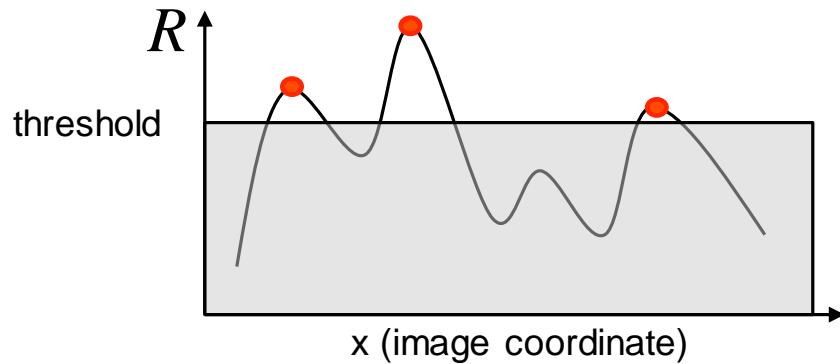
Ellipse rotates but its shape  
**(eigenvalues)** remains the same

**Corner response R is invariant to image rotation**

# Harris corner response is invariant to intensity changes

Partial invariance to *affine intensity* change

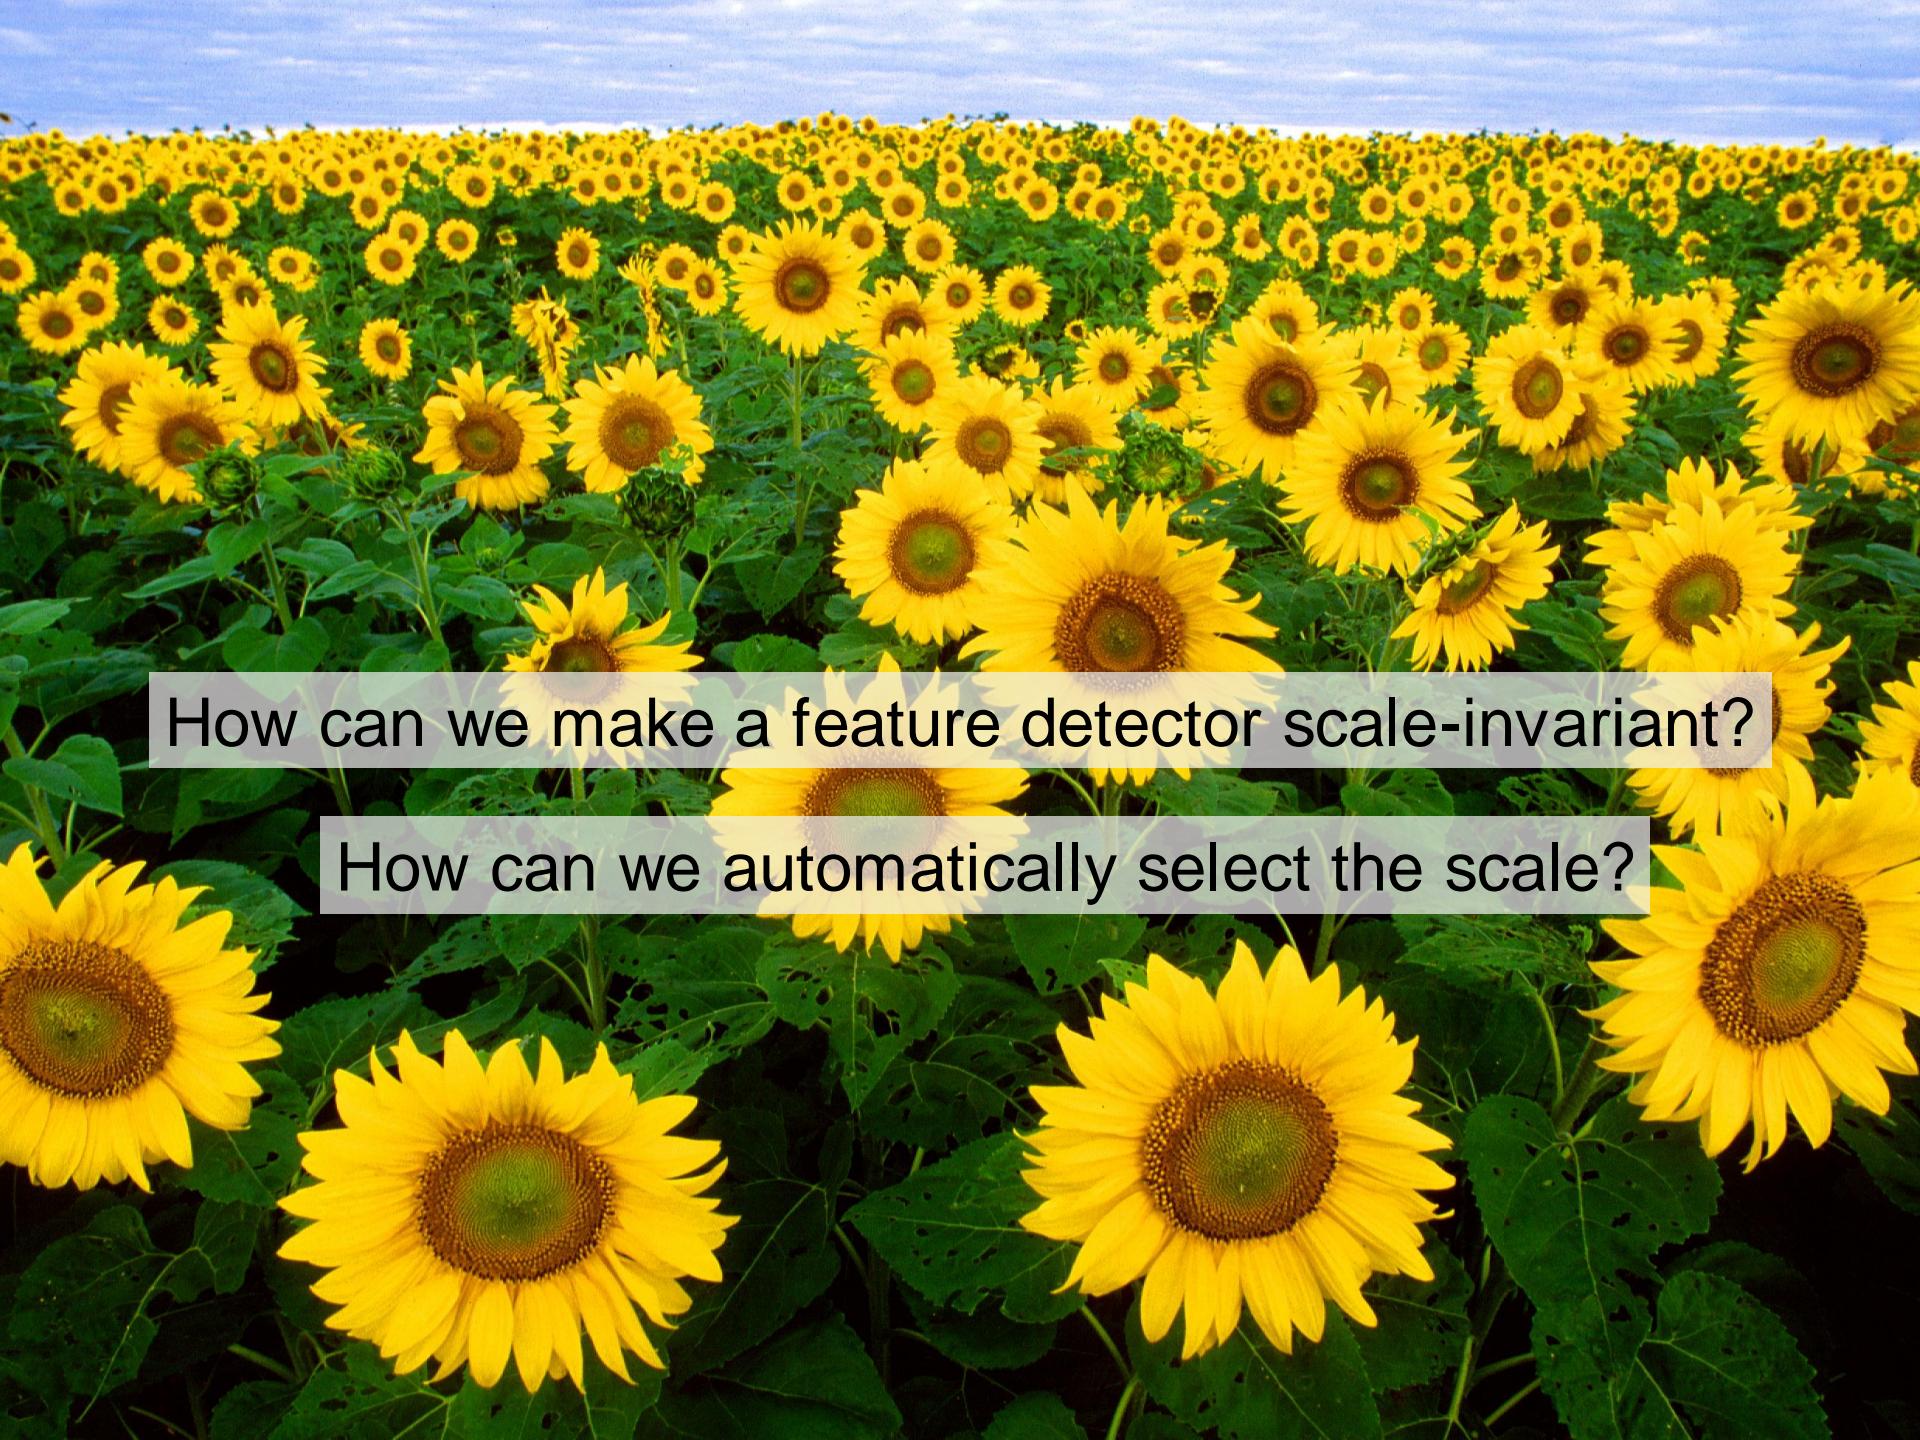
- Only derivatives are used => invariance to intensity shift  $I \rightarrow I + b$
- Intensity scale:  $I \rightarrow a I$



# The (Harris) corner detector is not invariant to scale



# Multi-scale detection

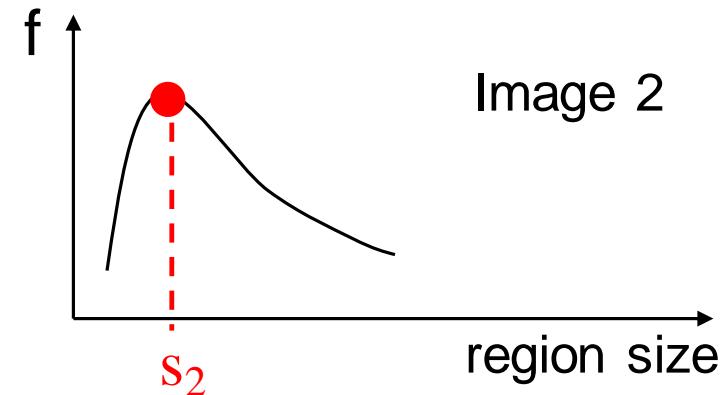
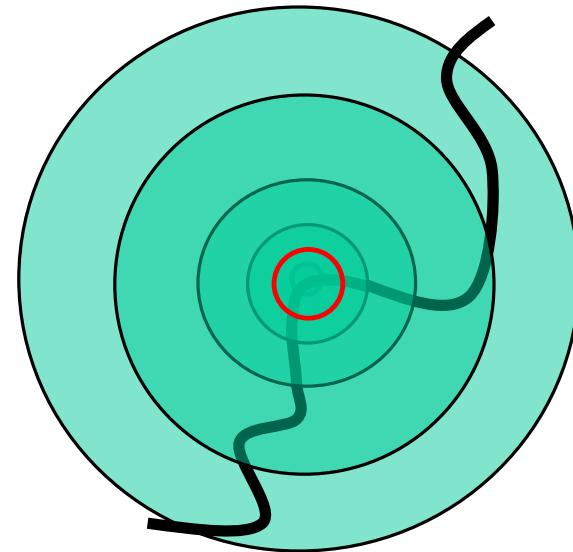
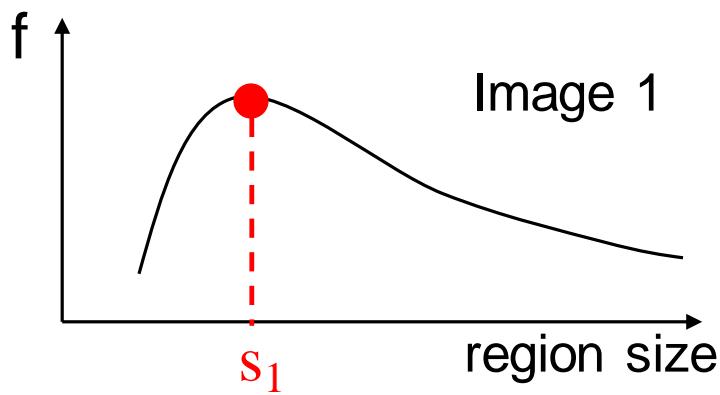
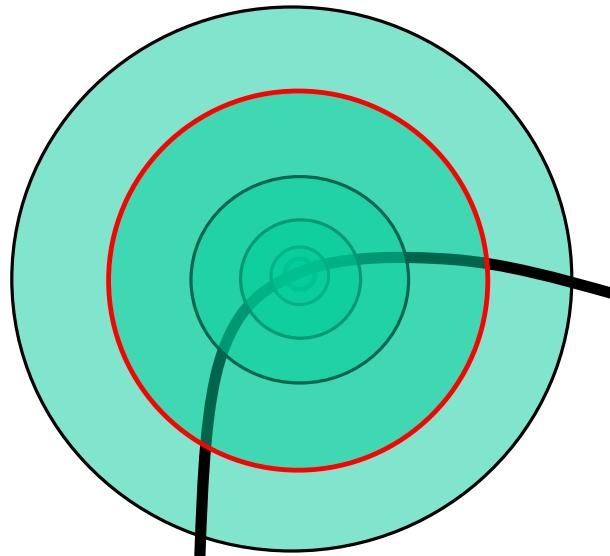
A wide-angle photograph of a vast field of sunflowers. The flowers are bright yellow with dark brown centers. They are densely packed, creating a pattern of yellow and green. The sky above is a clear, pale blue with a few wispy clouds.

How can we make a feature detector scale-invariant?

How can we automatically select the scale?

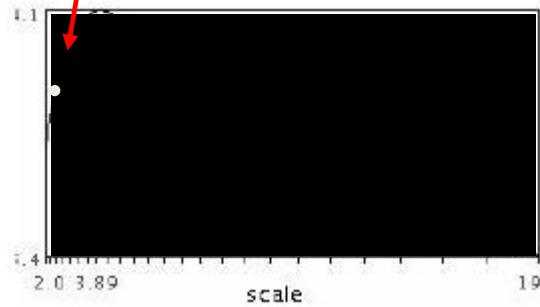
Intuitively...

Find local maxima in both **position** and **scale**

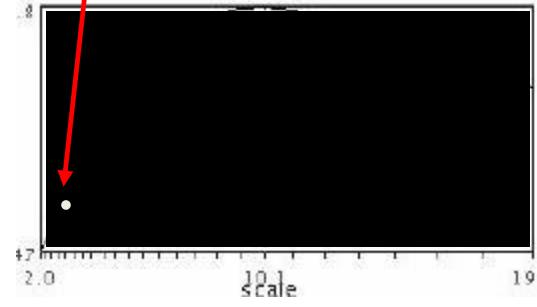


# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



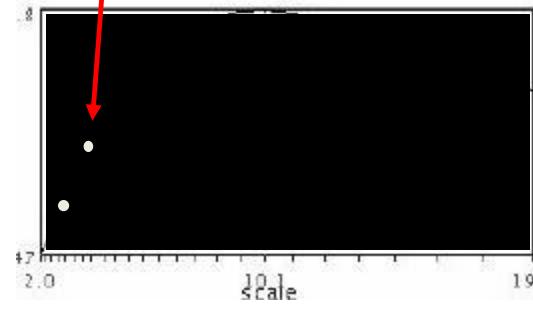
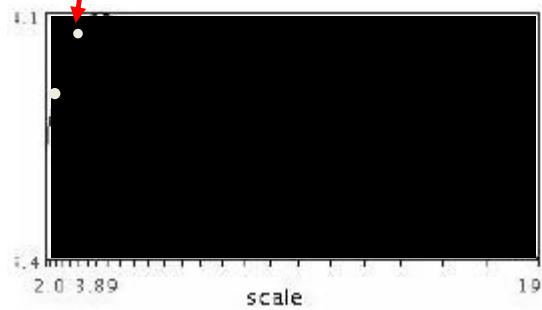
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



$$f(I_{i_1 \dots i_m}(x', \sigma))$$

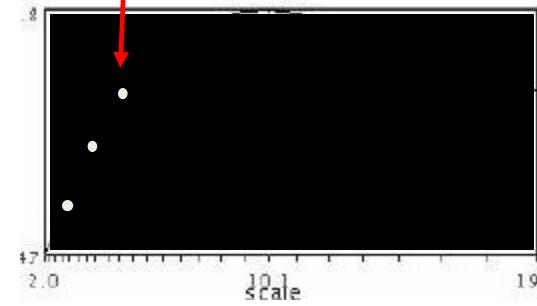
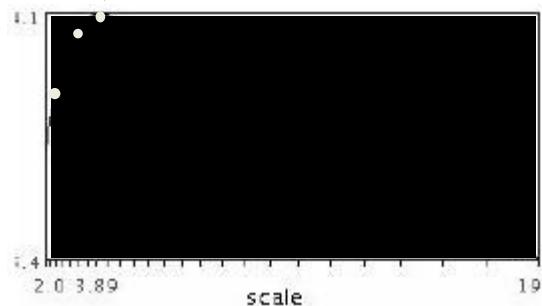
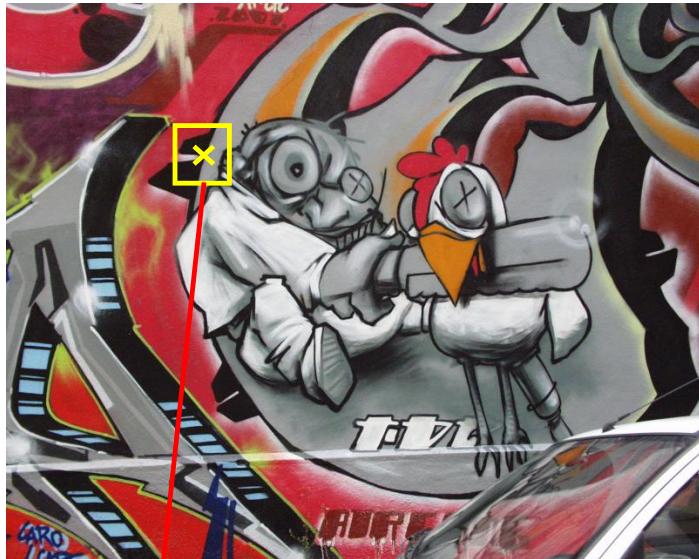
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



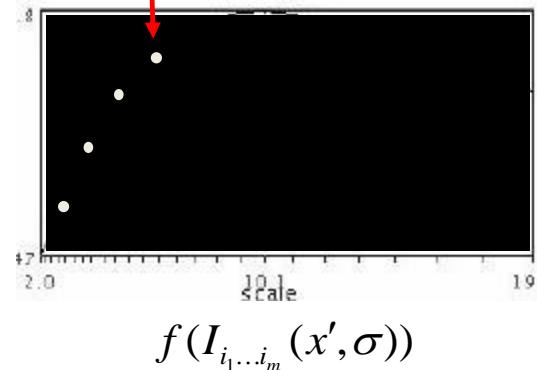
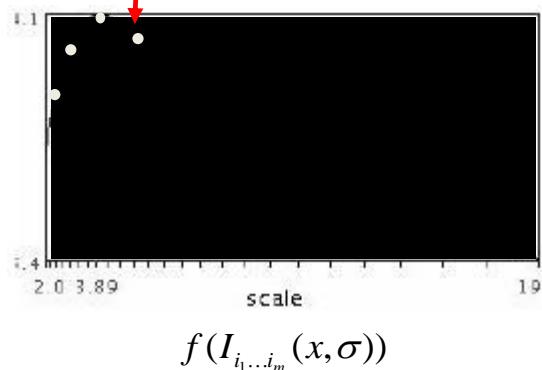
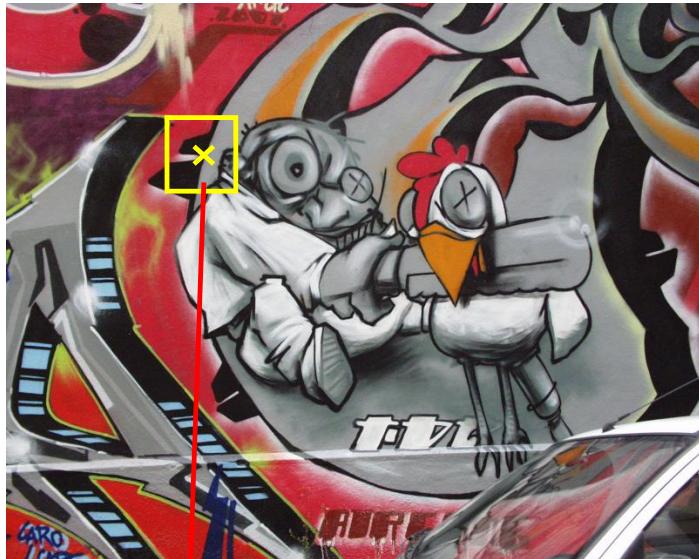
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



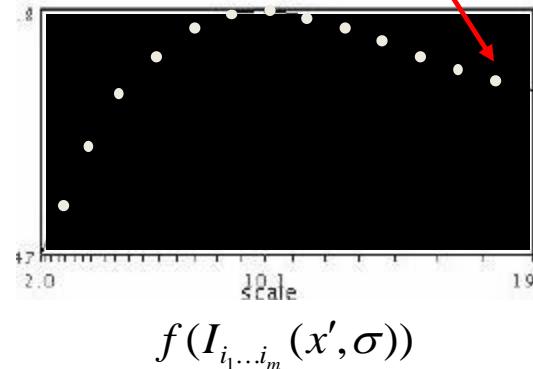
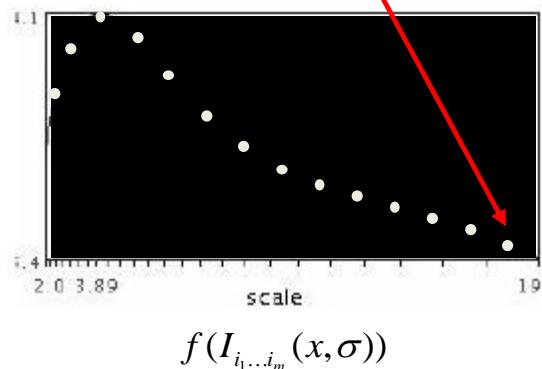
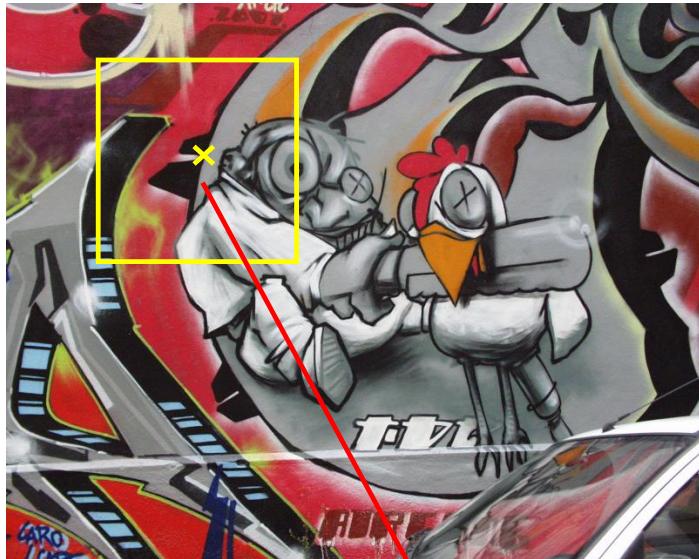
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



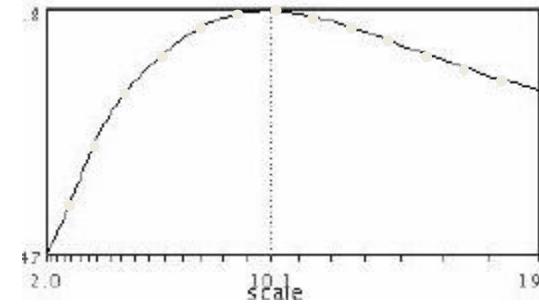
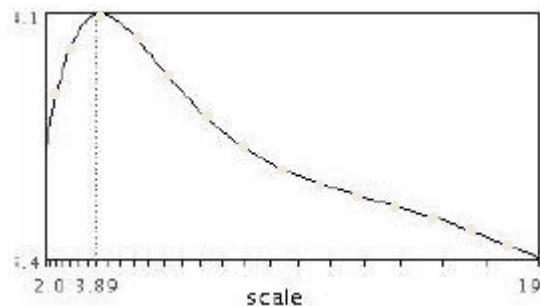
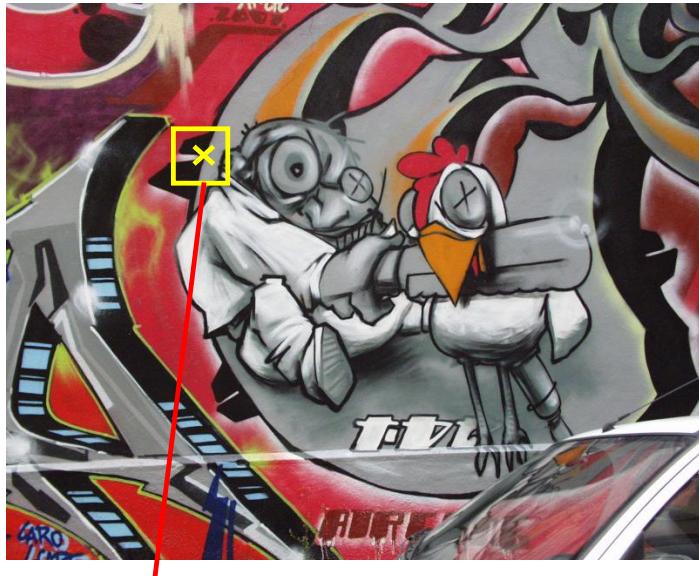
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



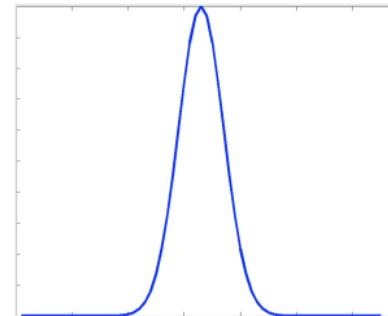
# Second Derivative Filters

---

Are applied to images to produce “signature” functions

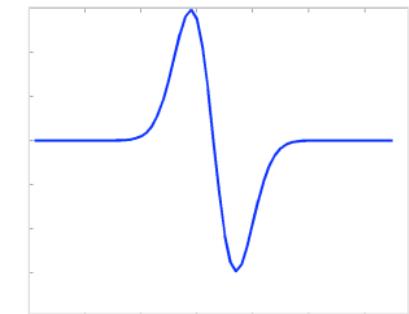
gaussian

$$g(x) = e^{-\frac{x^2}{2\sigma^2}}$$



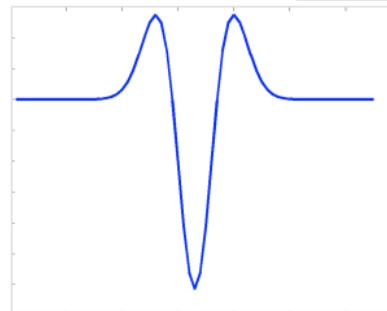
1<sup>st</sup> derivative

$$g'(x) = -\frac{1}{2\sigma^2} 2xe^{-\frac{x^2}{2\sigma^2}} = -\frac{x}{\sigma^2} e^{-\frac{x^2}{2\sigma^2}}$$



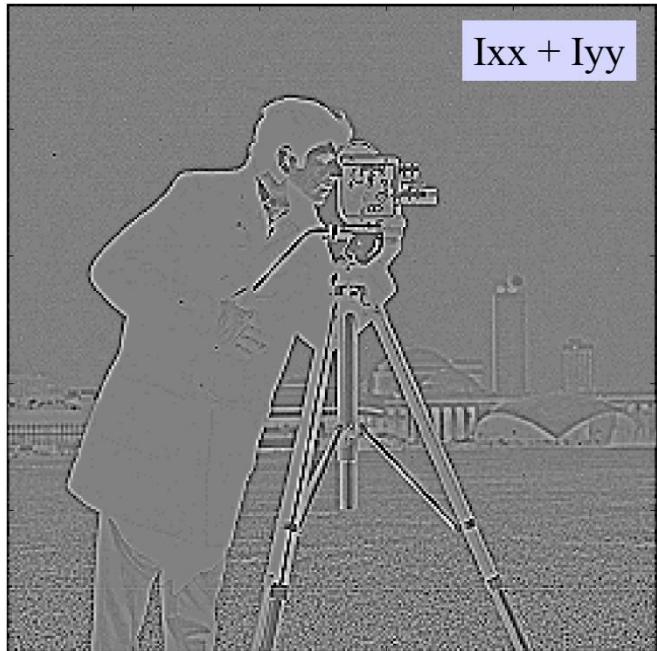
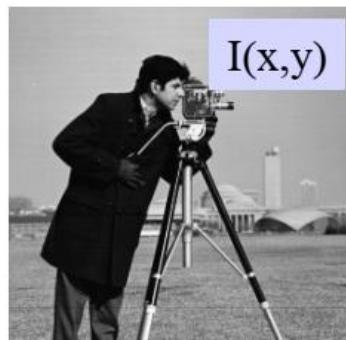
2<sup>nd</sup> derivative

$$g''(x) = \left(\frac{x^2}{\sigma^4} - \frac{1}{\sigma^2}\right) e^{-\frac{x^2}{2\sigma^2}}$$



# Second Derivative Filters

---

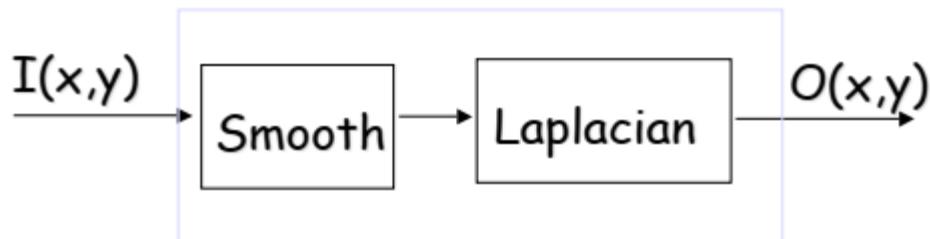


Laplacian

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix} * I$$

# Second Derivative Filters

---



$$\nabla^2(f(x, y) \otimes G(x, y)) = \nabla^2G(x, y) \otimes f(x, y)$$

$\overbrace{\qquad\qquad}^{\text{Laplacian of}} \qquad \overbrace{\qquad\qquad}^{\text{Used in practice}}$

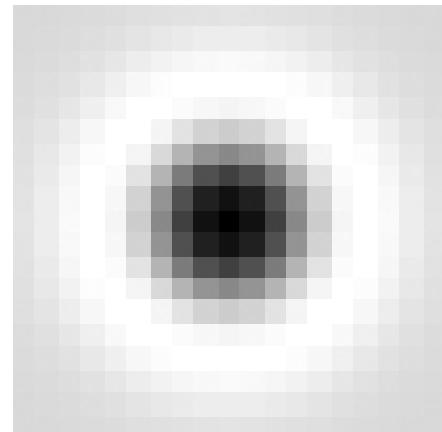
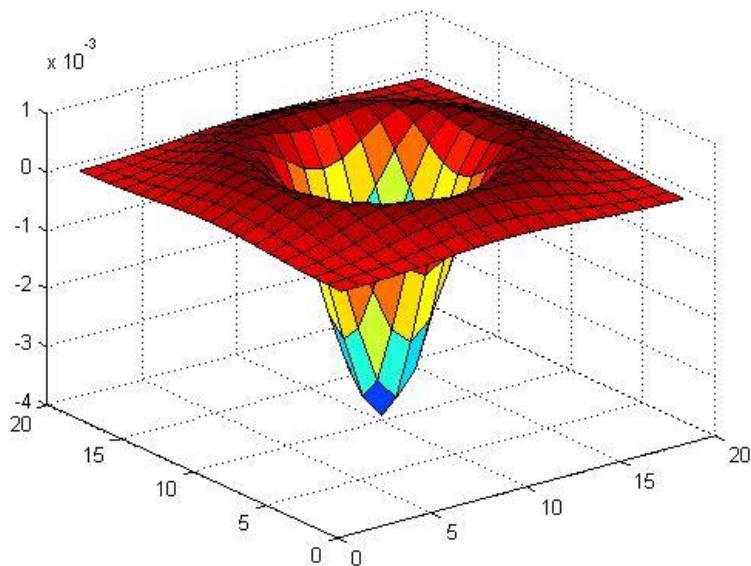
Laplacian of  
Gaussian-filtered image

Laplacian of Gaussian (LoG)  
-filtered image

# Blob detection in 2D

---

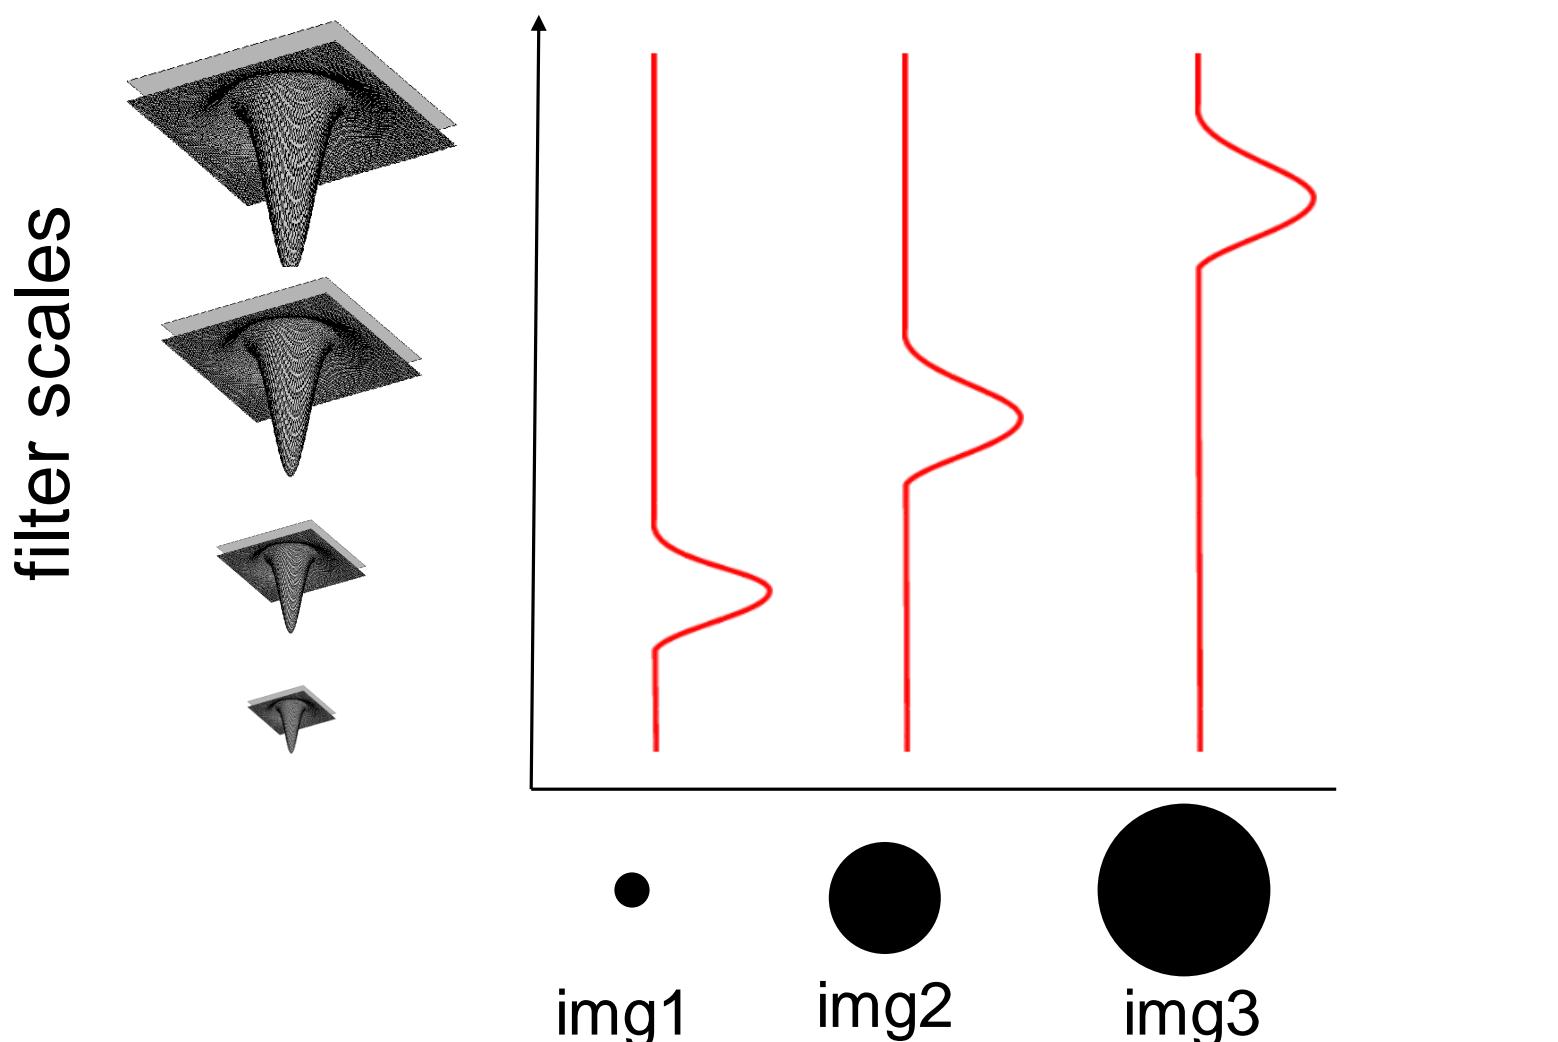
Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D



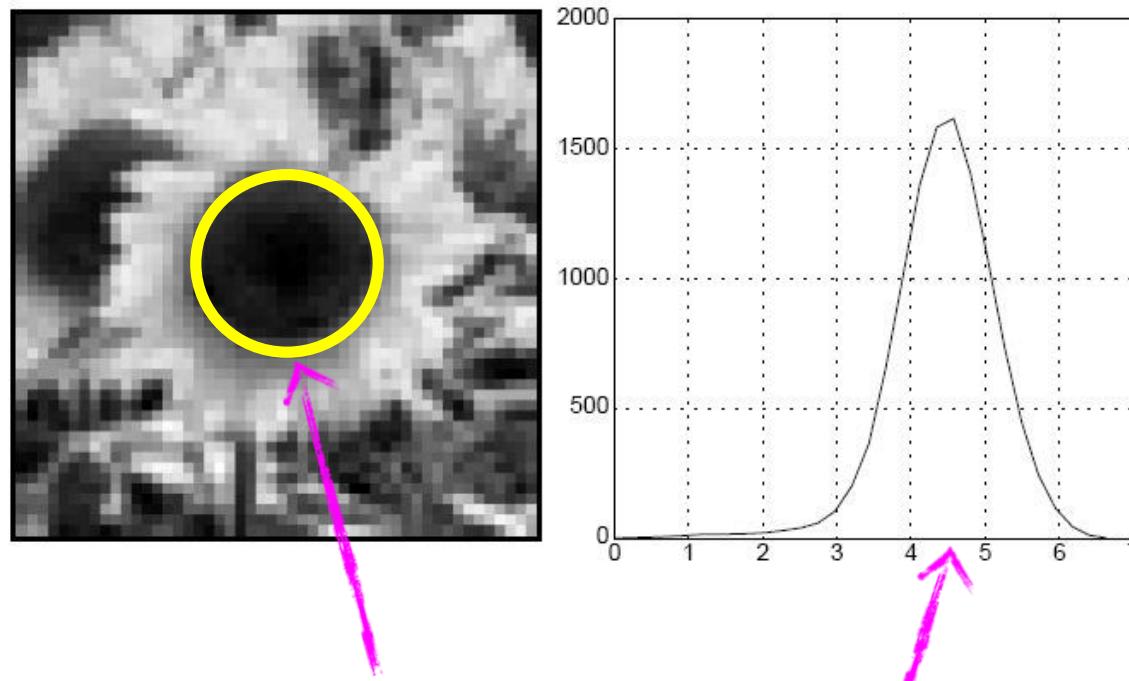
$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

# Blob detection in 2D: scale selection

Laplacian-of-Gaussian = “blob” detector     $\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$



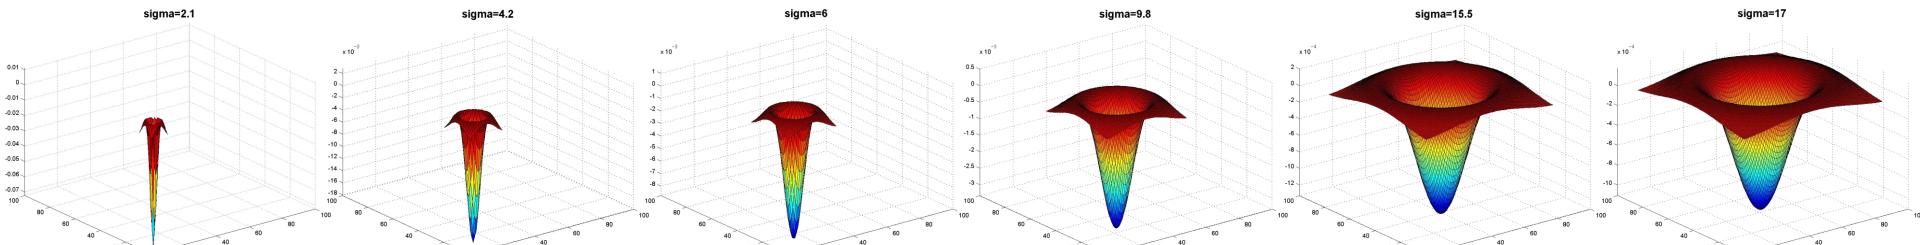
characteristic scale - the scale that produces peak filter response



characteristic scale

**we need to search over characteristic scales**

# What happens if you apply different Laplacian filters?



Full size



3/4 size



What happened when you applied different Laplacian filters?

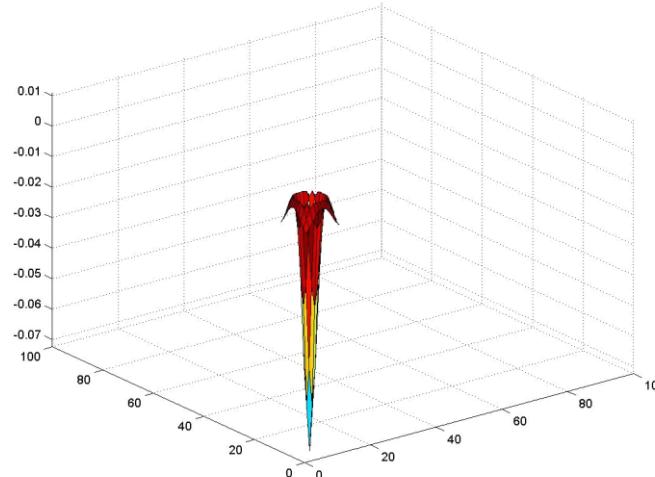
Full size



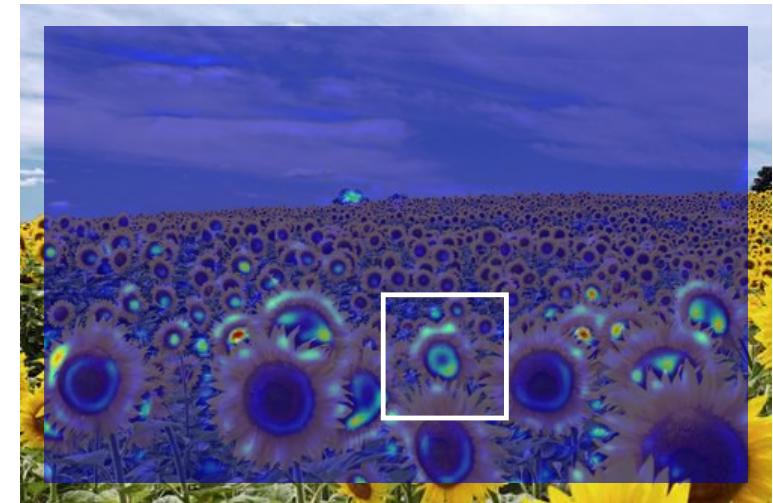
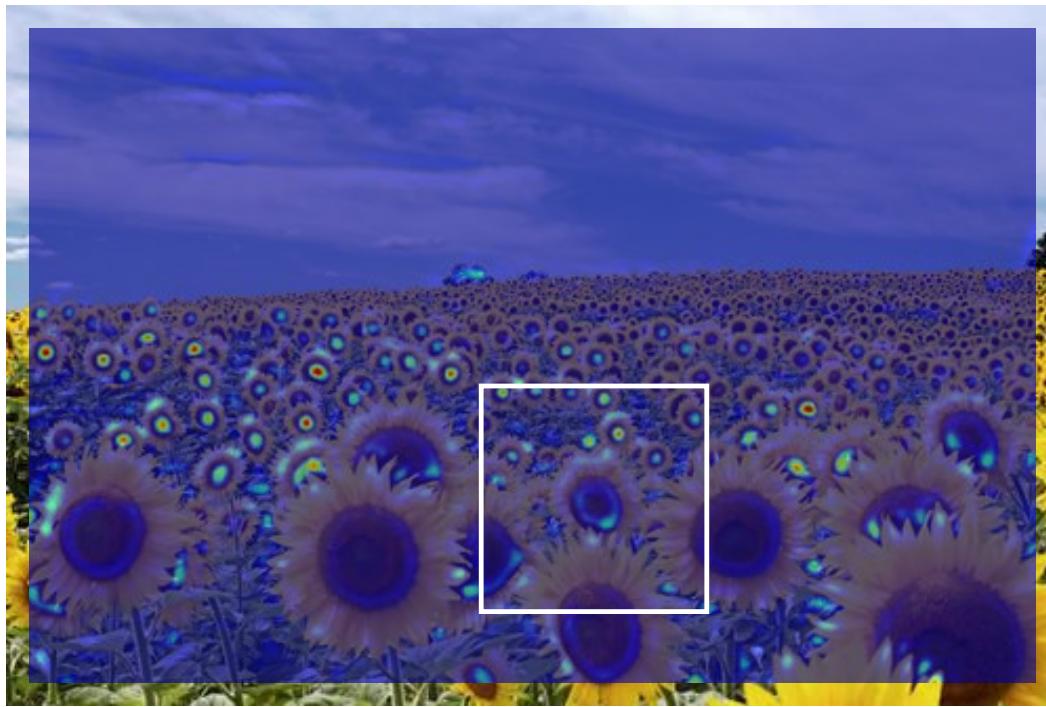
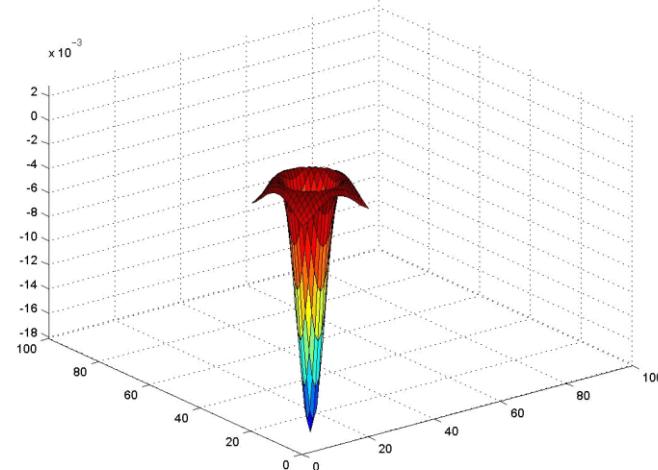
3/4 size



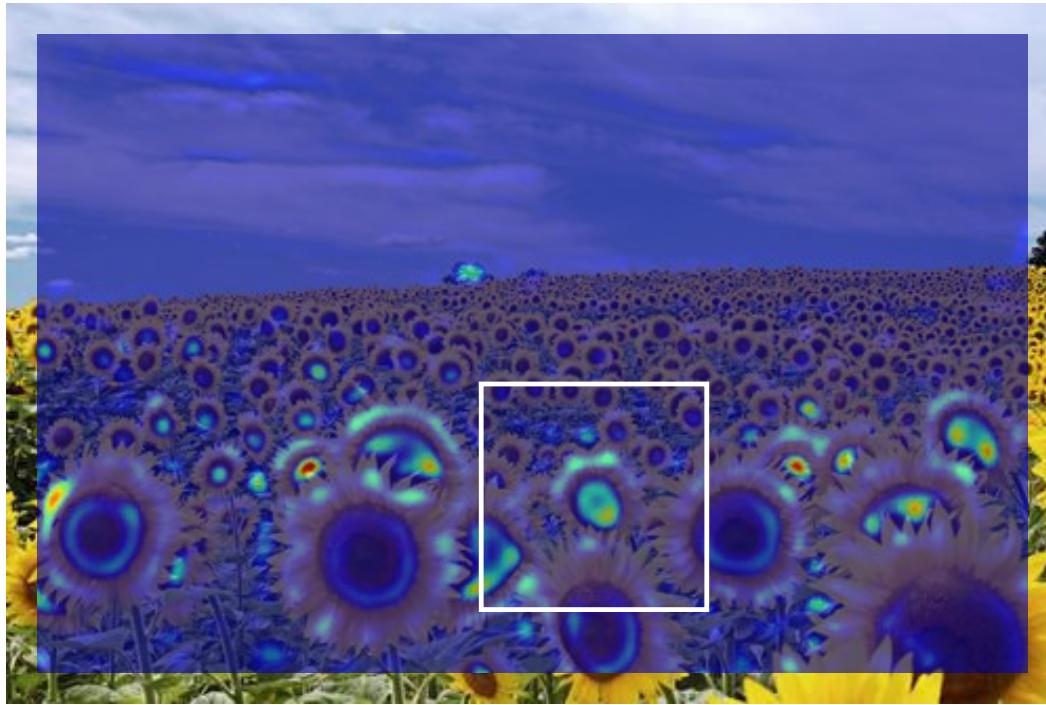
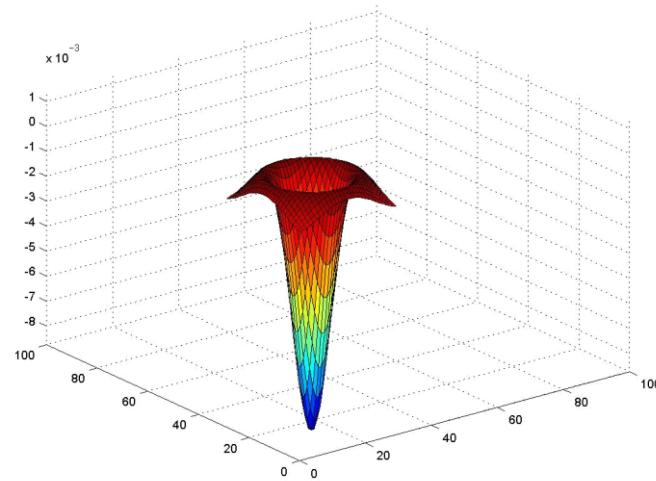
**sigma=2.1**

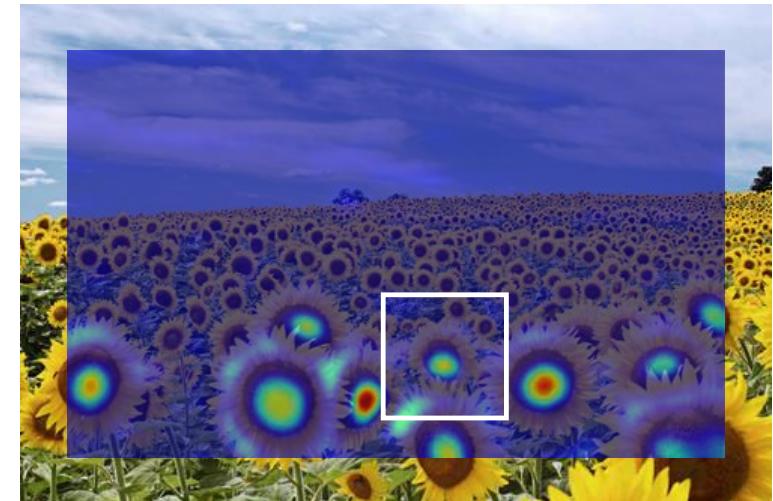
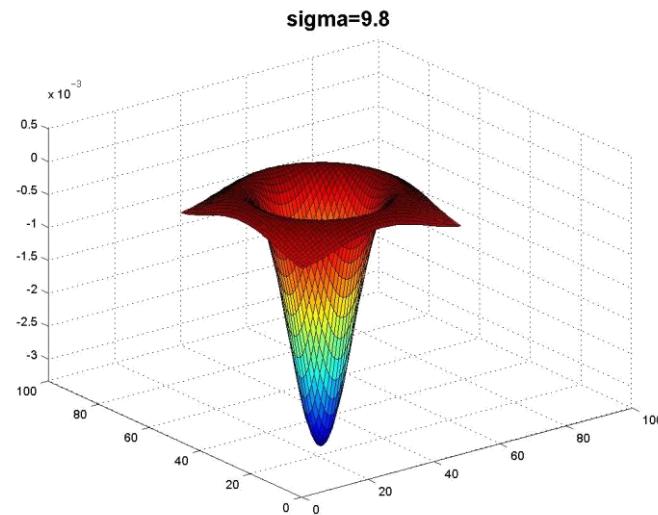


**sigma=4.2**

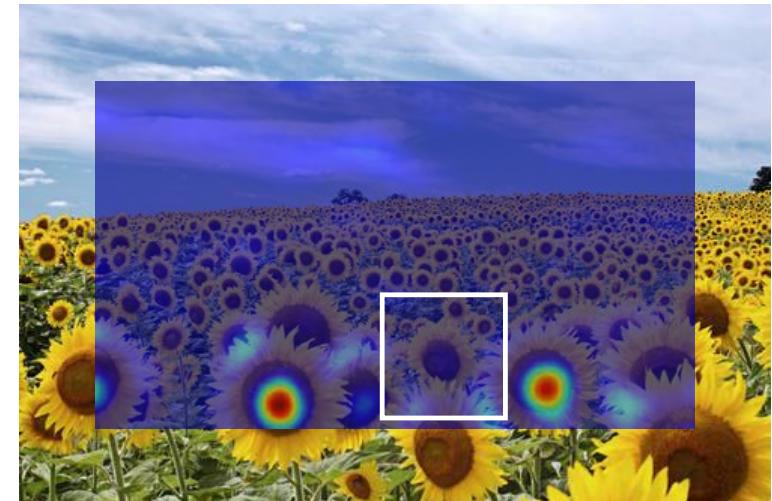
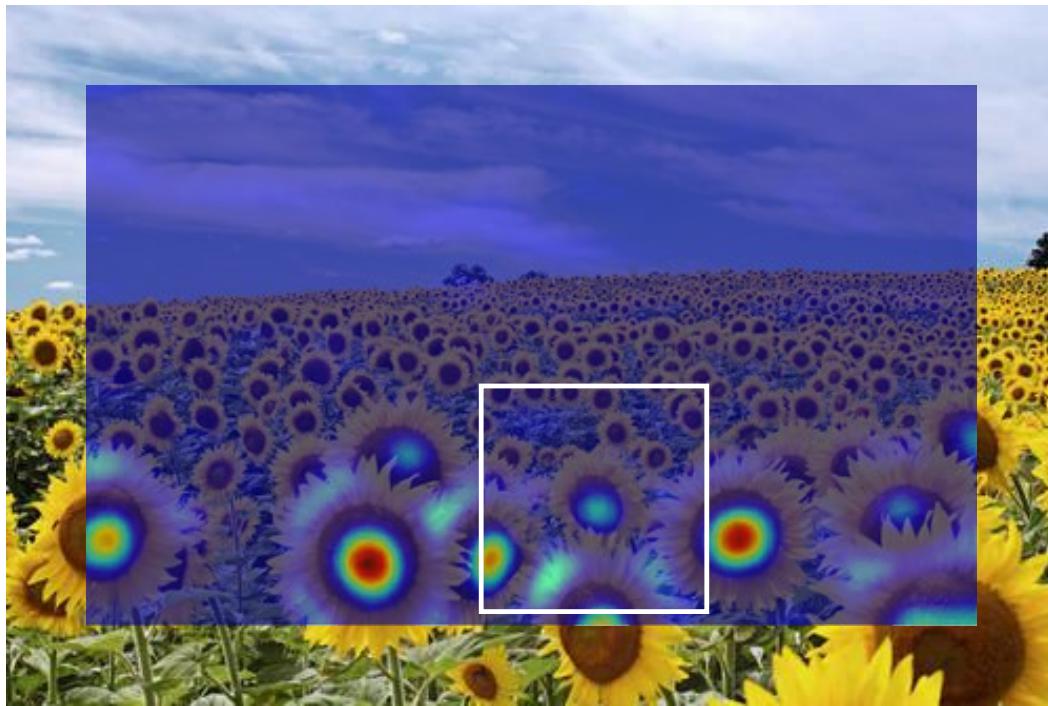
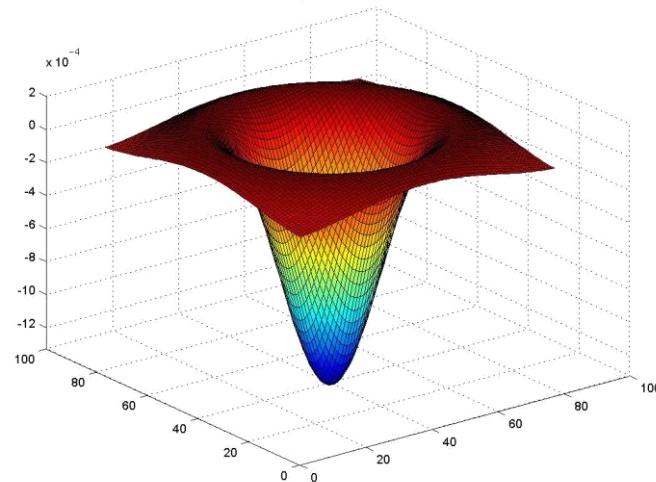


**sigma=6**

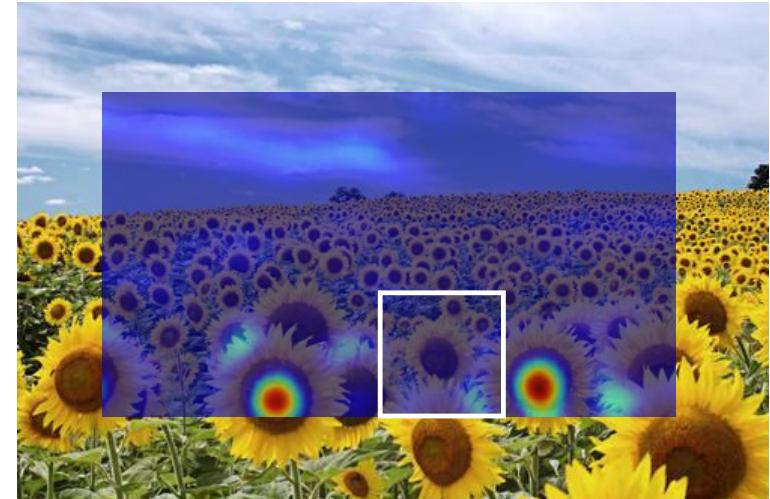
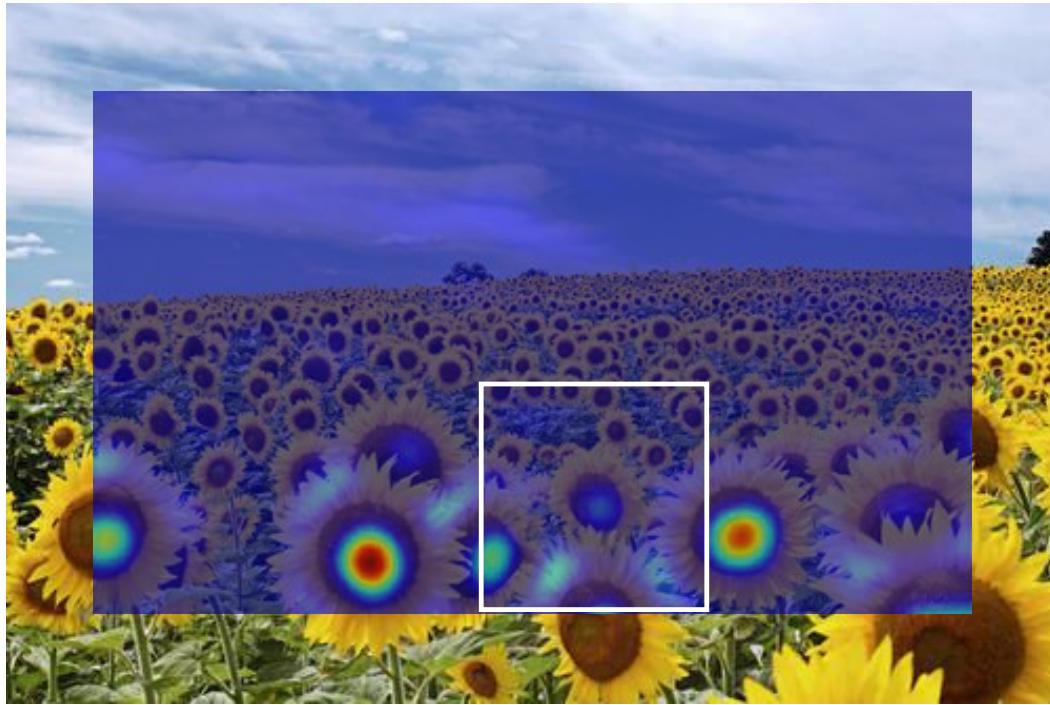
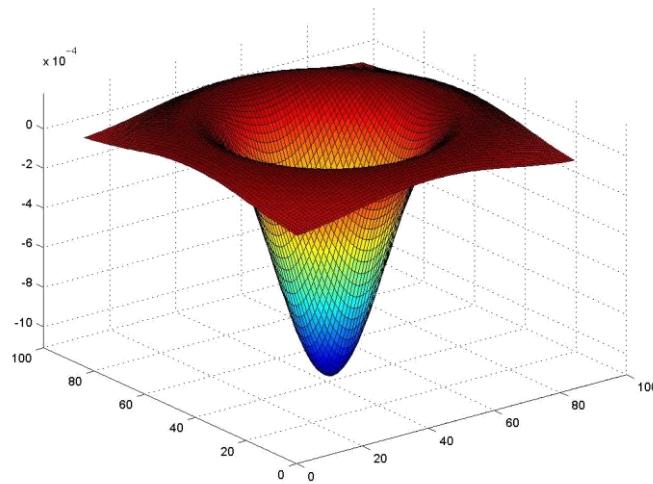




**sigma=15.5**



**sigma=17**



What happened when you applied different Laplacian filters?

Full size

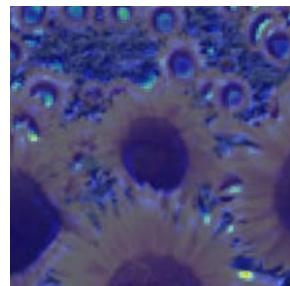


3/4 size

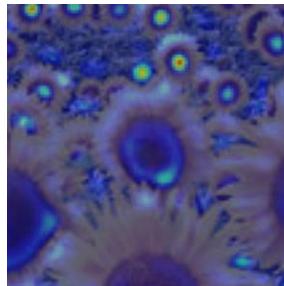


# optimal scale

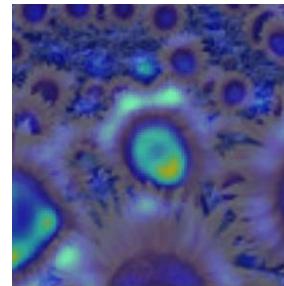
2.1



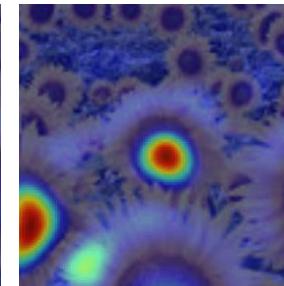
4.2



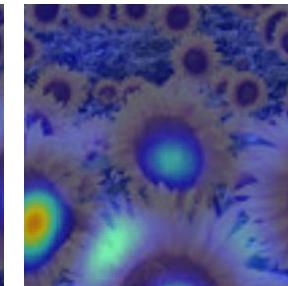
6.0



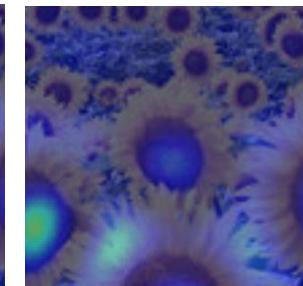
9.8



15.5

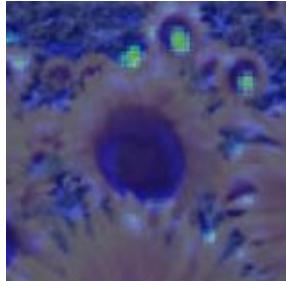


17.0

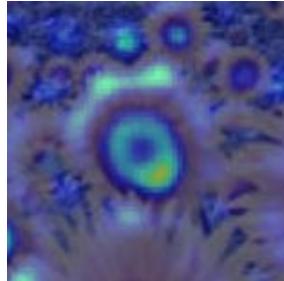


Full size image

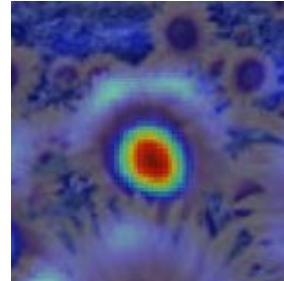
2.1



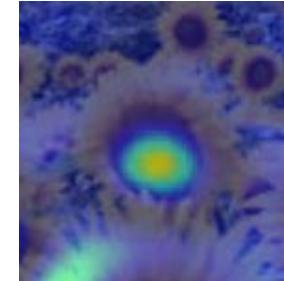
4.2



6.0



9.8



15.5



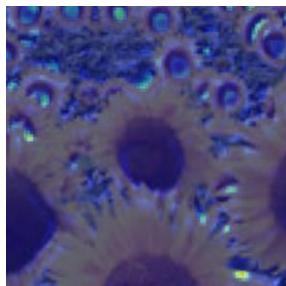
17.0



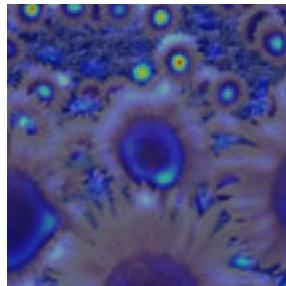
3/4 size image

# optimal scale

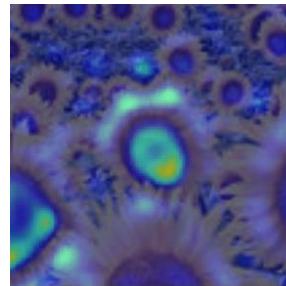
2.1



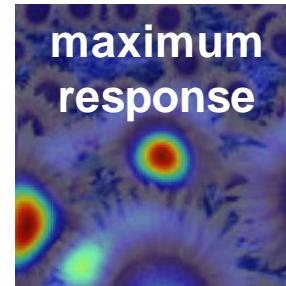
4.2



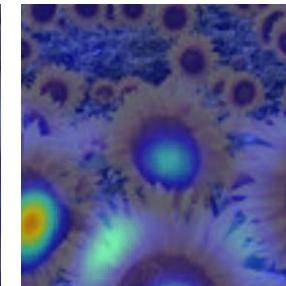
6.0



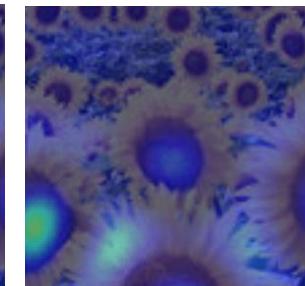
9.8



15.5

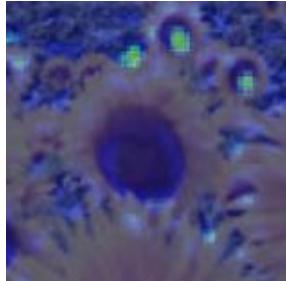


17.0

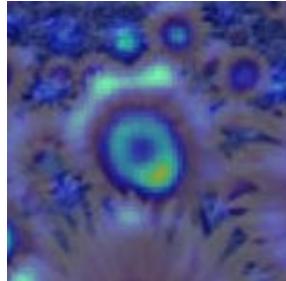


Full size image

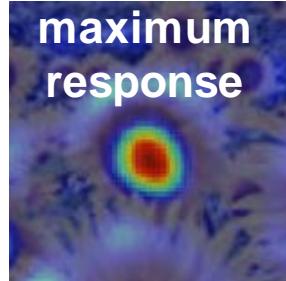
2.1



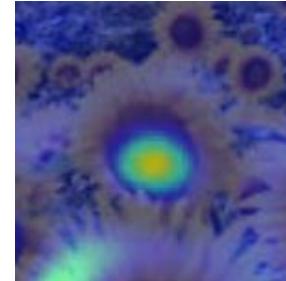
4.2



6.0



9.8



15.5

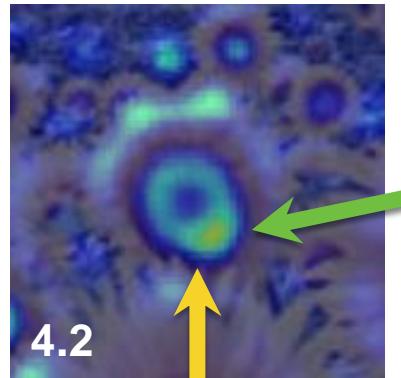


17.0

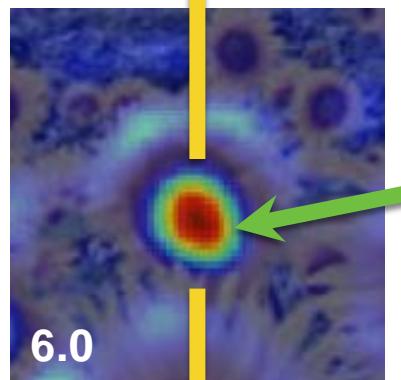


3/4 size image

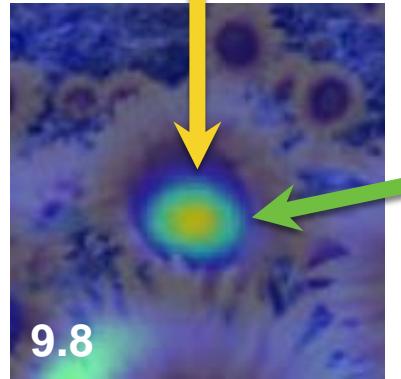
cross-scale maximum



local maximum



local maximum



local maximum

# implementation

For each level of the Gaussian pyramid:

- compute feature response (e.g. Harris, Laplacian)

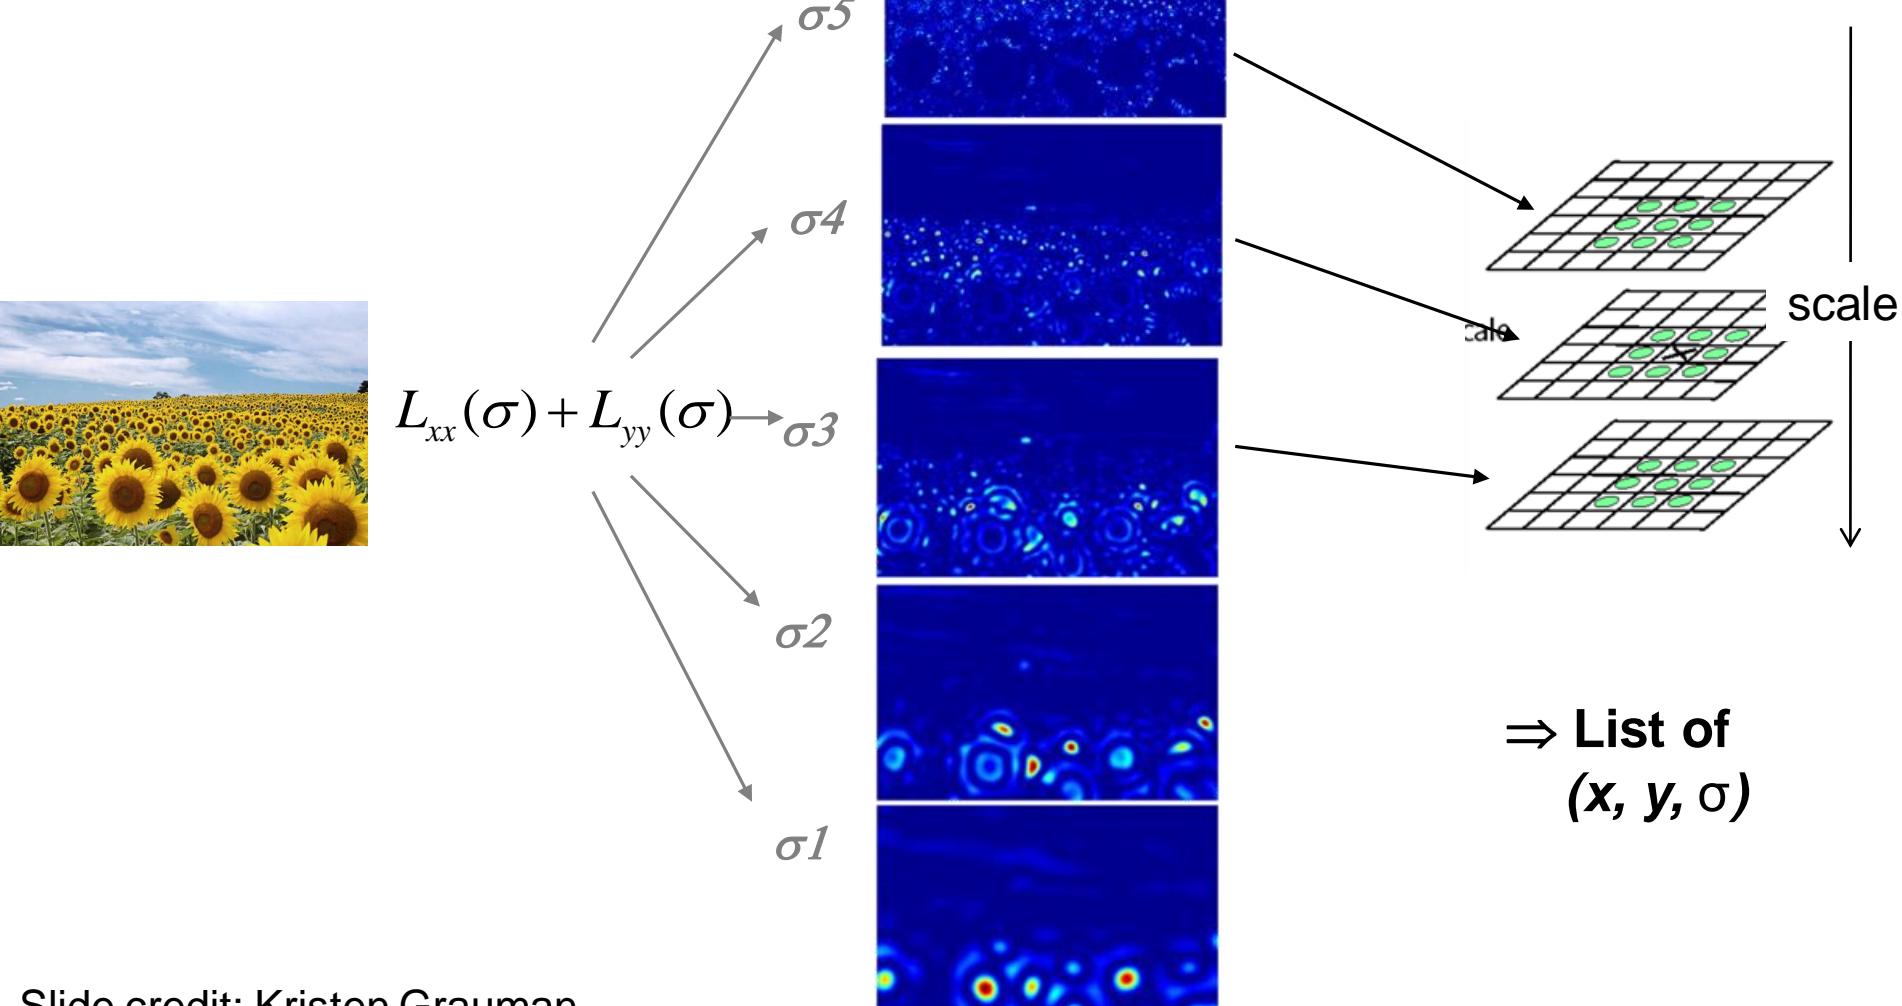
For each level of the Gaussian pyramid:

- if local and cross-scale maximum:

- save** scale and location of feature  $(x, y, s)$

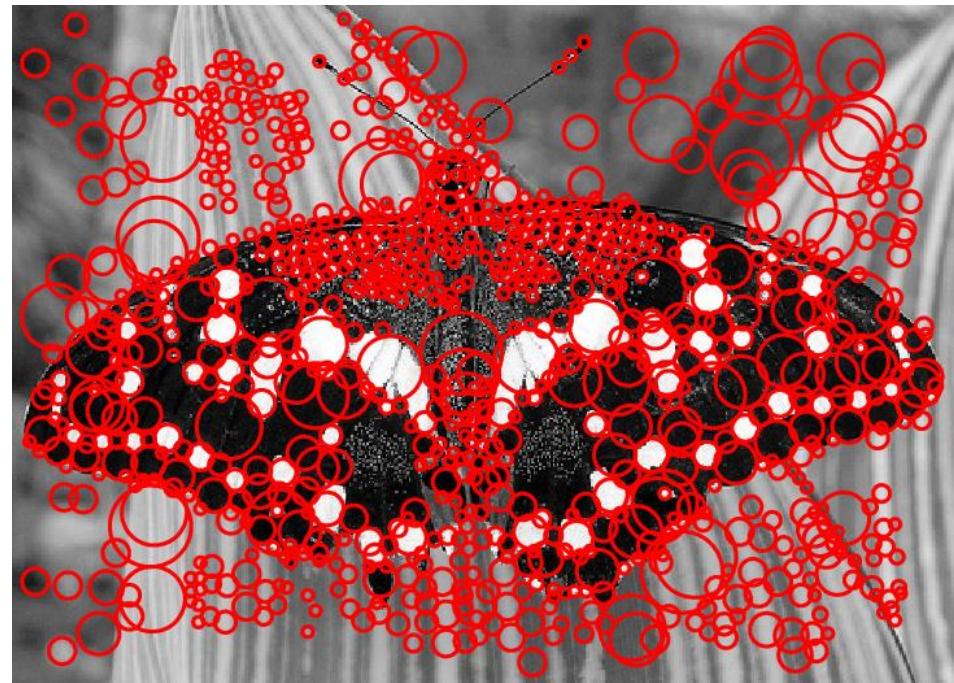
# Scale invariant interest points

Interest points are local maxima in both position and scale.



# Scale-space blob detector: Example

---



# Technical detail

---

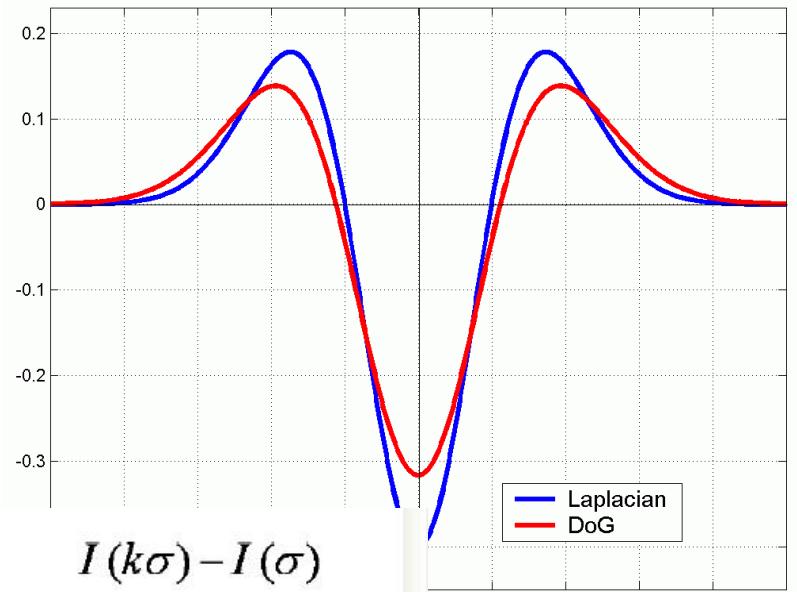
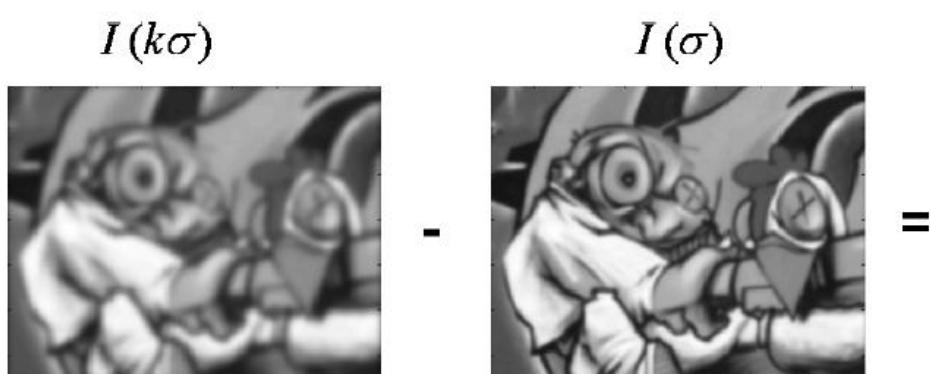
Can approximate the Laplacian with a difference of Gaussians; more efficient

$$L = \sigma^2 \left( G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right)$$

(Laplacian of Gaussian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)



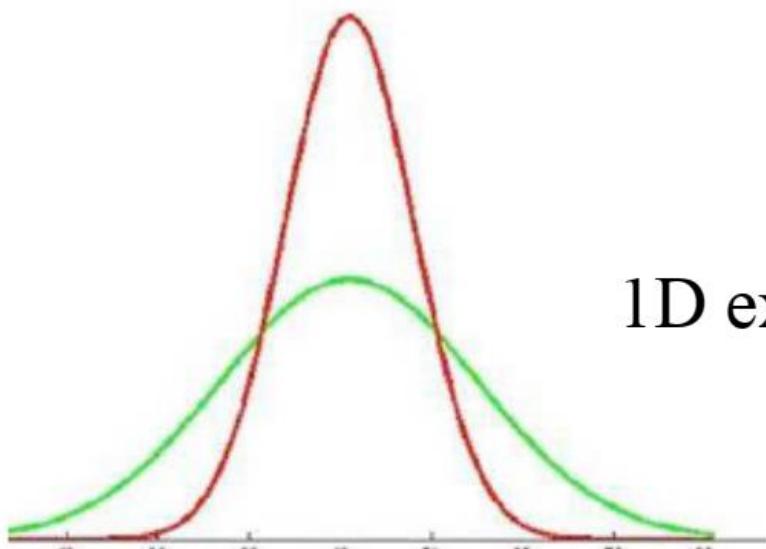
# Technical detail

---

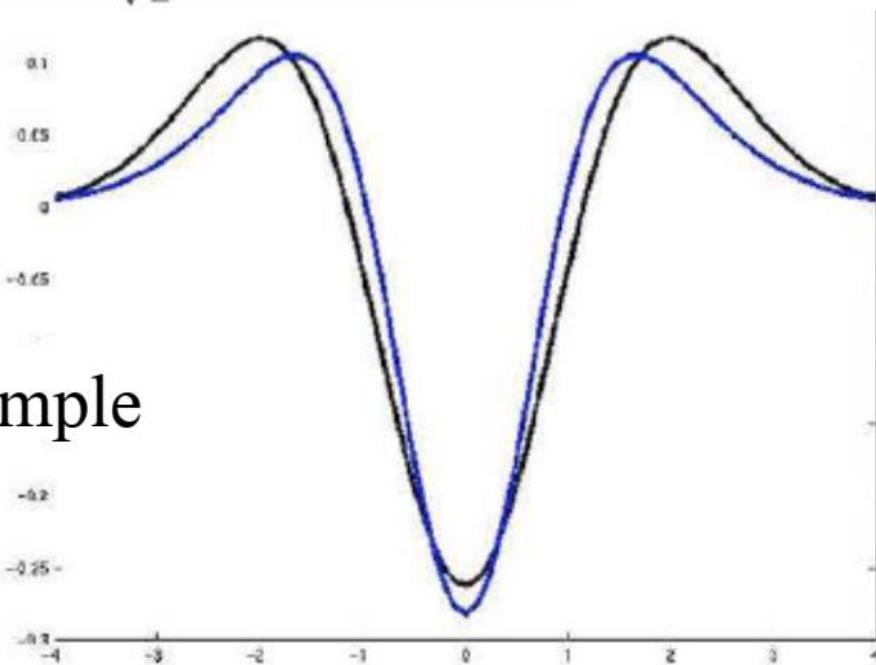
LoG can be approximate by a Difference of two Gaussians (DoG) at different scales

$$\nabla^2 G_\sigma \approx G_{\sigma_1} - G_{\sigma_2}$$

Best approximation when:  
 $\sigma_1 = \frac{\sigma}{\sqrt{2}}, \sigma_2 = \sqrt{2}\sigma$



1D example



# Scale Invariant Feature Transform (SIFT) descriptor [Lowe 2004]

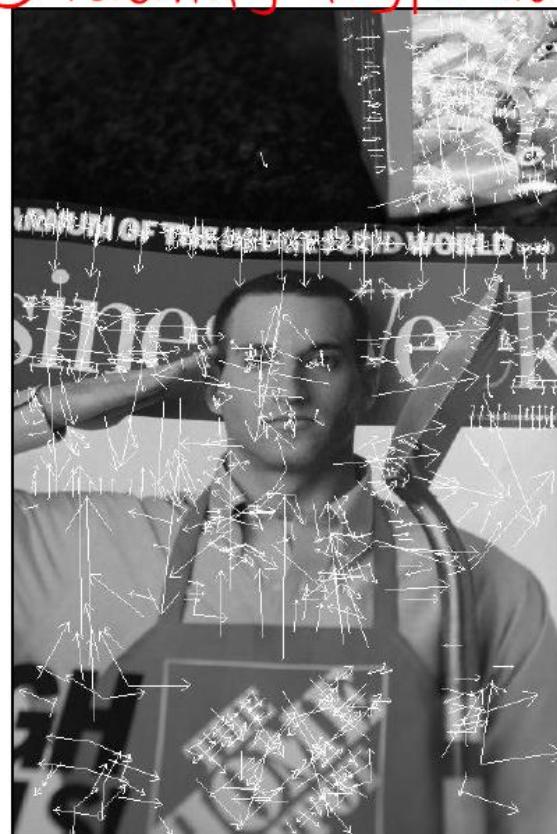
- introduction to the image matching problem
- image matching using SIFT features
- multi-scale interest-point detection
- **the SIFT feature detector**
- the SIFT descriptor

# Step 1: Compute a Set of Keypoints

Source Image I



① Identify keypoints

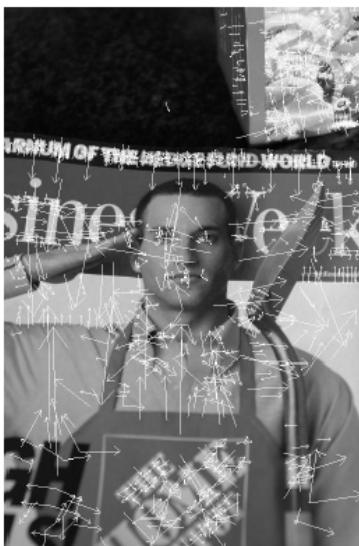


Goals :

- Identify distinctive image locations
- Assign scale & orientation to each keypoint
- Should be able to detect same keypoint in images that vary in magnification, brightness, etc

# Computing SIFT Keypoints: Basic Steps

Source image I



Gauss-pyramid

Step 1a  
Build pyramid  
of Gauss-  
smoothed  
images

DOG pyramid

Step 1b  
Build DOG  
pyramid

DOG extrema

Step 1c  
Locate extrema  
of DOG  
pyramid  
 $(x_i, y_i, p_i)$

Step 1f

Assign  
orientation  
to extrema  
 $p'_i = (x_i, y_i, p_i, \theta_i)$

Step 1e

Prune set of  
extrema  
Keypoints = {  
all remaining  
 $(x_i, y_i, p_i)$  }

Step 1d

Refine location  
of DOG  
extrema  
 $(x_i, y_i, p_i) \rightarrow (x_i, y_i, p'_i)$

Orientation assign

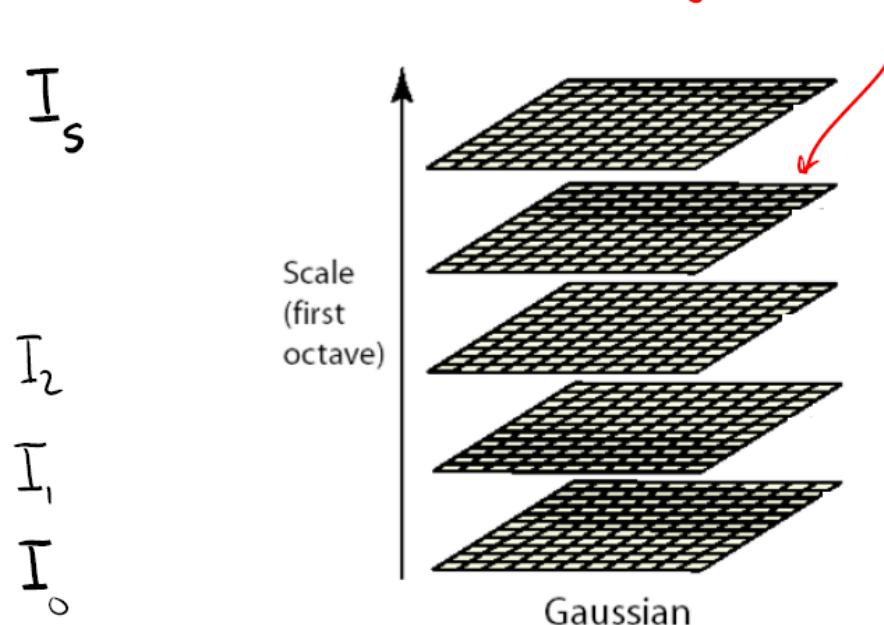
Extremum pruning

Location refinement

## Step 1a: Construct a Gauss-like Pyramid

Step 1a: Compute a pyramid of Gauss-filtered images, organized into octaves of  $s+1$  images

each image is smoothed by a factor of  $\kappa$  more than the image below it



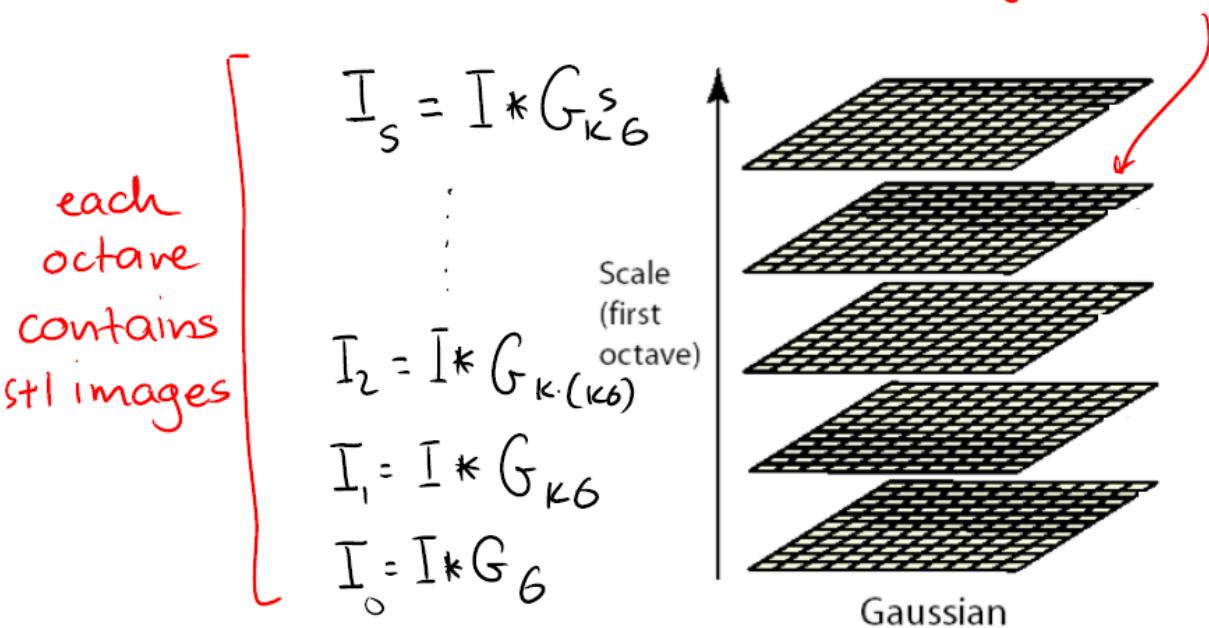
# Step 1a: Construct a Gauss-like Pyramid

SIFT

terminology-

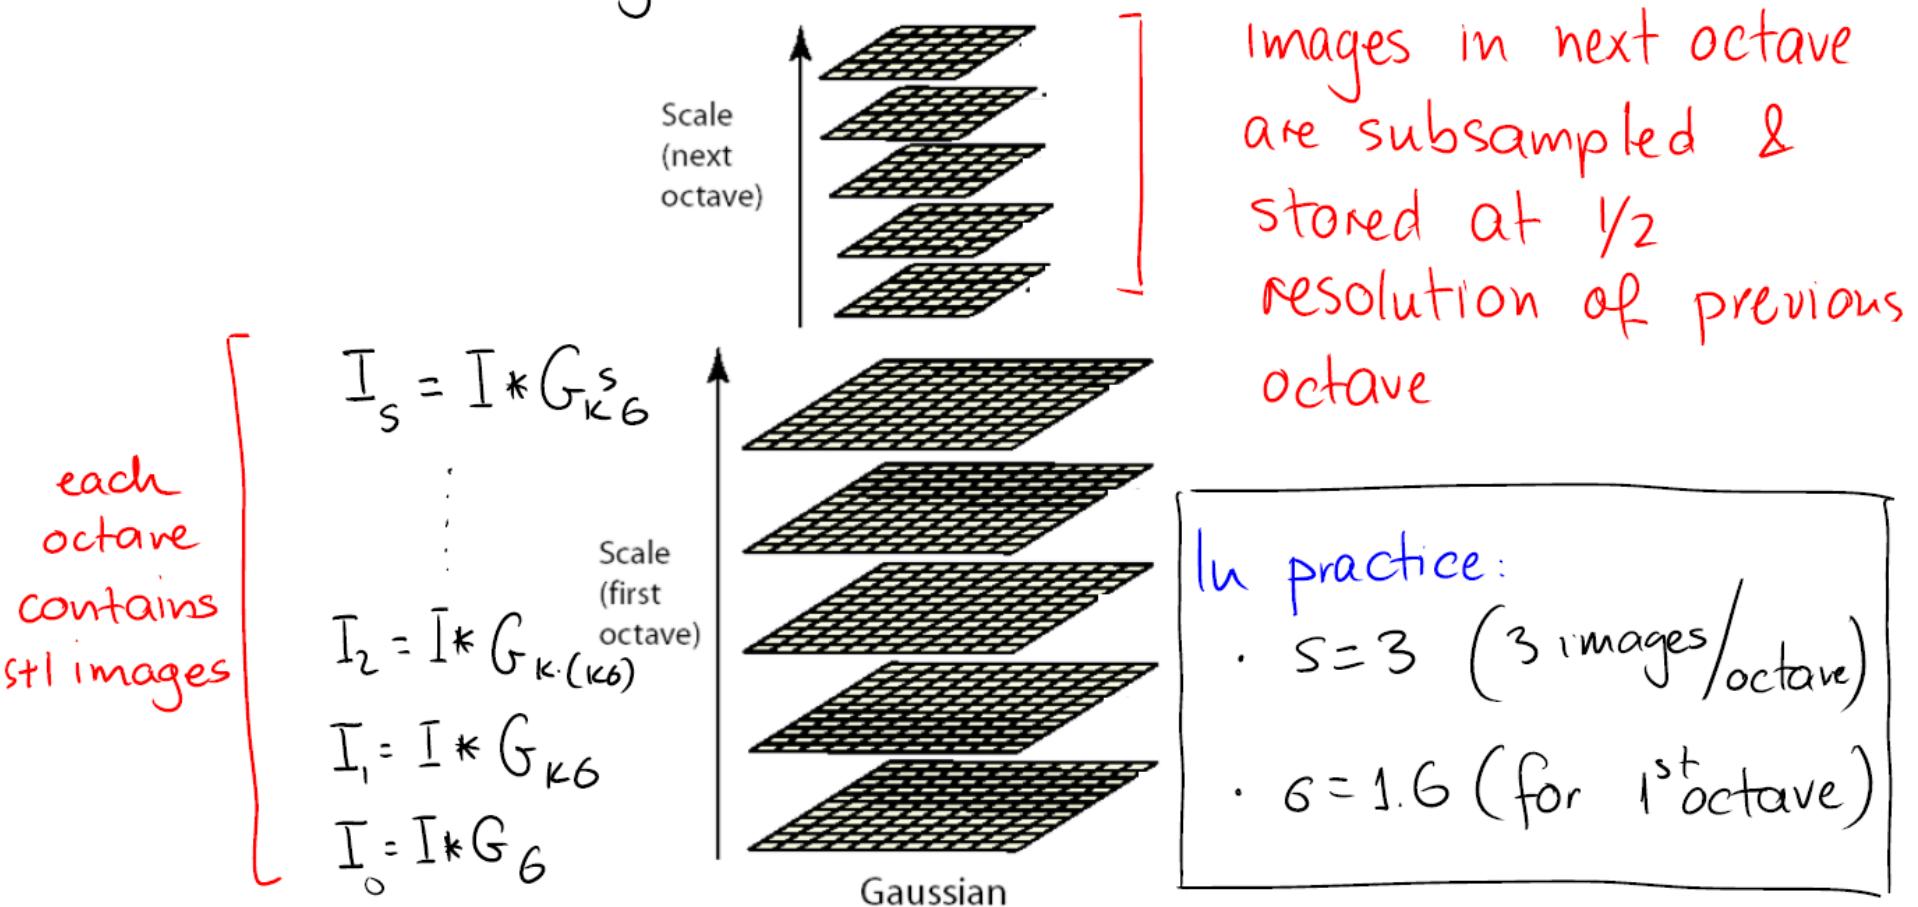
An octave is a set of Gauss-convolved images,  $I_1, \dots, I_s$  representing a doubling of the scale parameter  $\sigma$  between  $I_1$  and  $I_s$ .

each image is smoothed by a factor of  $\kappa$  more than the image below it



# Step 1a: Construct a Gauss-like Pyramid

Step 1a: Compute a pyramid of Gauss-filtered images, organized into octaves of  $S+1$  images



# Computing SIFT Keypoints: Basic Steps

Source image I



Gauss-pyramid

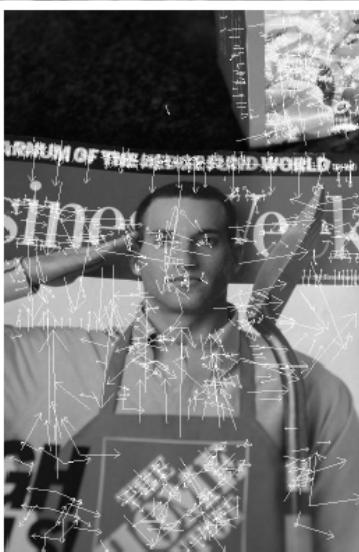
Step 1a  
Build pyramid  
of Gauss-  
smoothed  
images

DOG pyramid

Step 1b  
Build DOG  
pyramid

DOG extrema

Step 1c  
Locate extrema  
of DOG  
pyramid  
 $(x_i, y_i, p_i)$



Step 1f

Assign  
orientation  
to extrema  
 $P_i = (x_i, y_i, p_i, \theta_i)$

Step 1e

Prune set of  
extrema  
Keypoints = {  
all remaining  
 $(x_i, y_i, p_i)$  }

Step 1d

Refine location  
of DOG  
extrema  
 $(x_i, y_i, p_i) \rightarrow (x_i, y_i, p'_i)$

Orientation assign

Extremum pruning

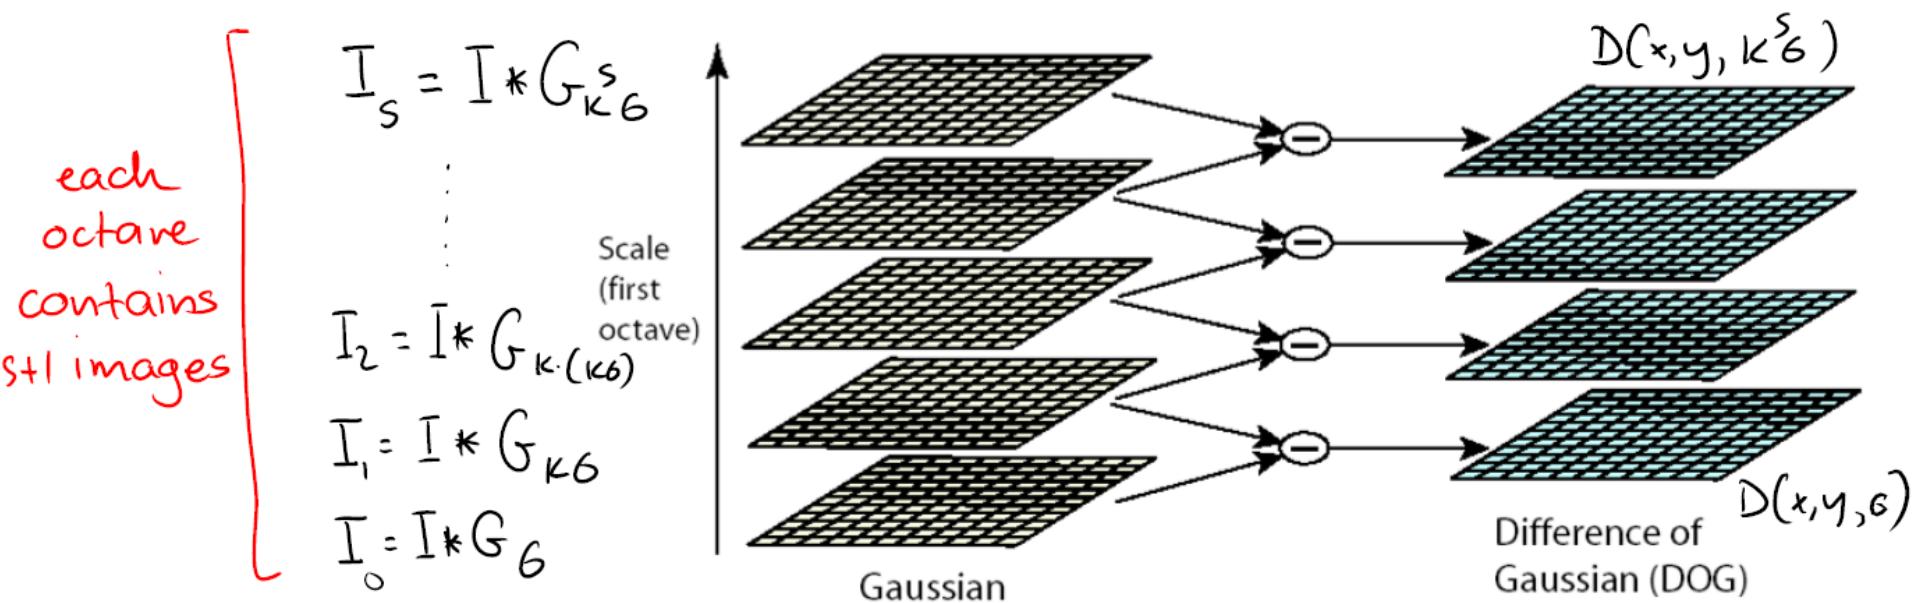
Location refinement

## Step 1b: Compute Pyramid of DOG Images

Step 1b: Compute a pyramid of DOG-filtered images

$$D(x,y,\rho) = I(x,y) * (G(x,y,k\rho) - G(x,y,\rho))$$

for  $\rho = 6, k6, k^26, \dots, k^{S-1}6$

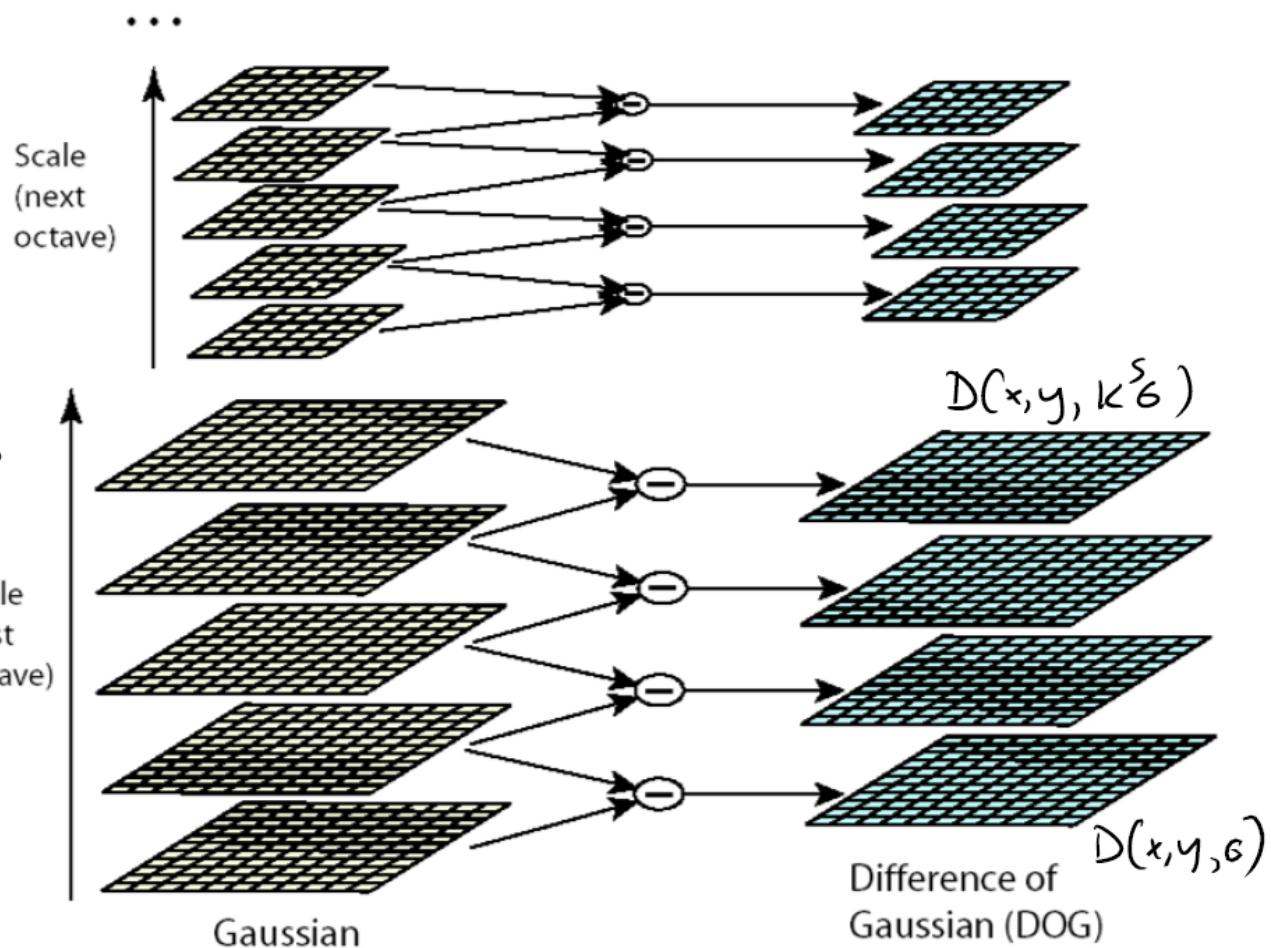


# Step 1b: Compute Pyramid of DOG Images

Step 1b: Compute a pyramid of DOG-filtered images

$$D(x,y,p) = I(x,y) * (G(x,y, \kappa p) - G(x,y, p))$$

for  $p = 6, \kappa 6, \kappa^2 6, \dots, \kappa^{S-1} 6$



# Reminder: Difference-of-Gaussian Filtering

---

$$I * G_\rho$$



# Reminder: Difference-of-Gaussian Filtering

---

$$I * G_{kp}$$



# Reminder: Difference-of-Gaussian Filtering

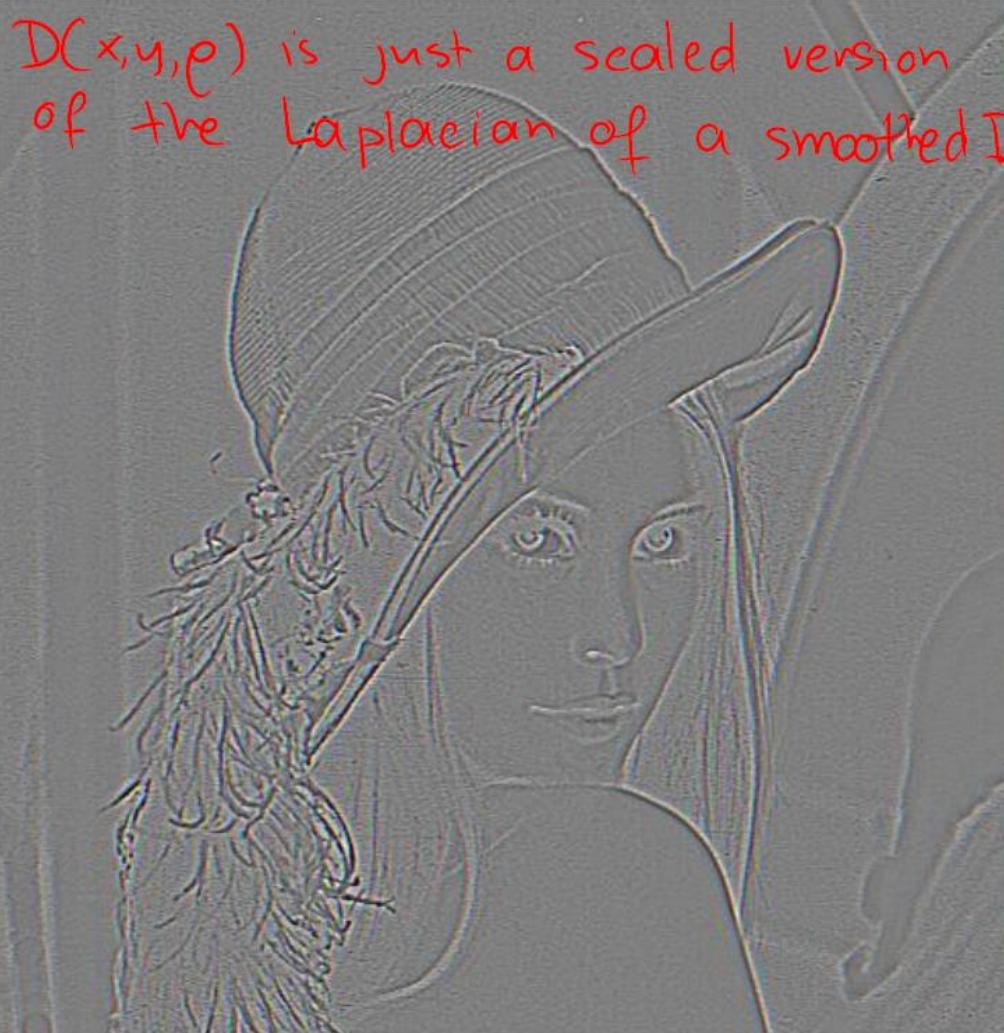
Difference of  
two Gaussian-  
smoothed versions  
of  $I$ :

$$I * G_{kp} -$$

$$I * G_p =$$

$$I * (G_{kp} - G_p)$$

(just the  
difference  
between two  
Gaussian masks)



But we know  
that

$$G_{kp} - G_p = kp(kp-p)\nabla^2 G_p$$

# Computing SIFT Keypoints: Basic Steps

Source image I



Gauss-pyramid

Step 1a

Build pyramid  
of Gauss-  
smoothed  
images

DOG pyramid

Step 1b

Build DOG  
pyramid

DOG extrema

Step 1c

Locate extrema  
of DOG  
pyramid  
 $(x_i, y_i, p_i)$

Step 1f

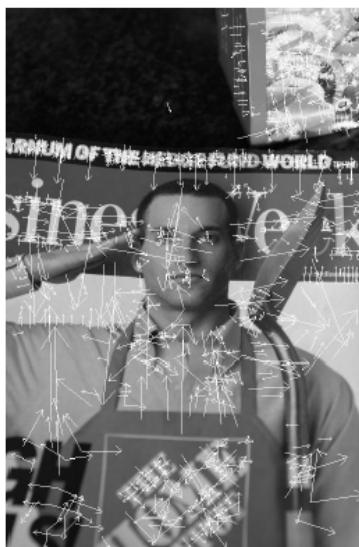
Assign  
orientation  
to extrema  
 $p'_i = (x_i, y_i, p_i, \theta_i)$

Step 1e

Prune set of  
extrema  
Keypoints = {  
all remaining  
 $(x_i, y_i, p_i)$  }

Step 1d

Refine location  
of DOG  
extrema  
 $(x_i, y_i, p_i) \rightarrow (x'_i, y'_i, p'_i)$



Orientation assign

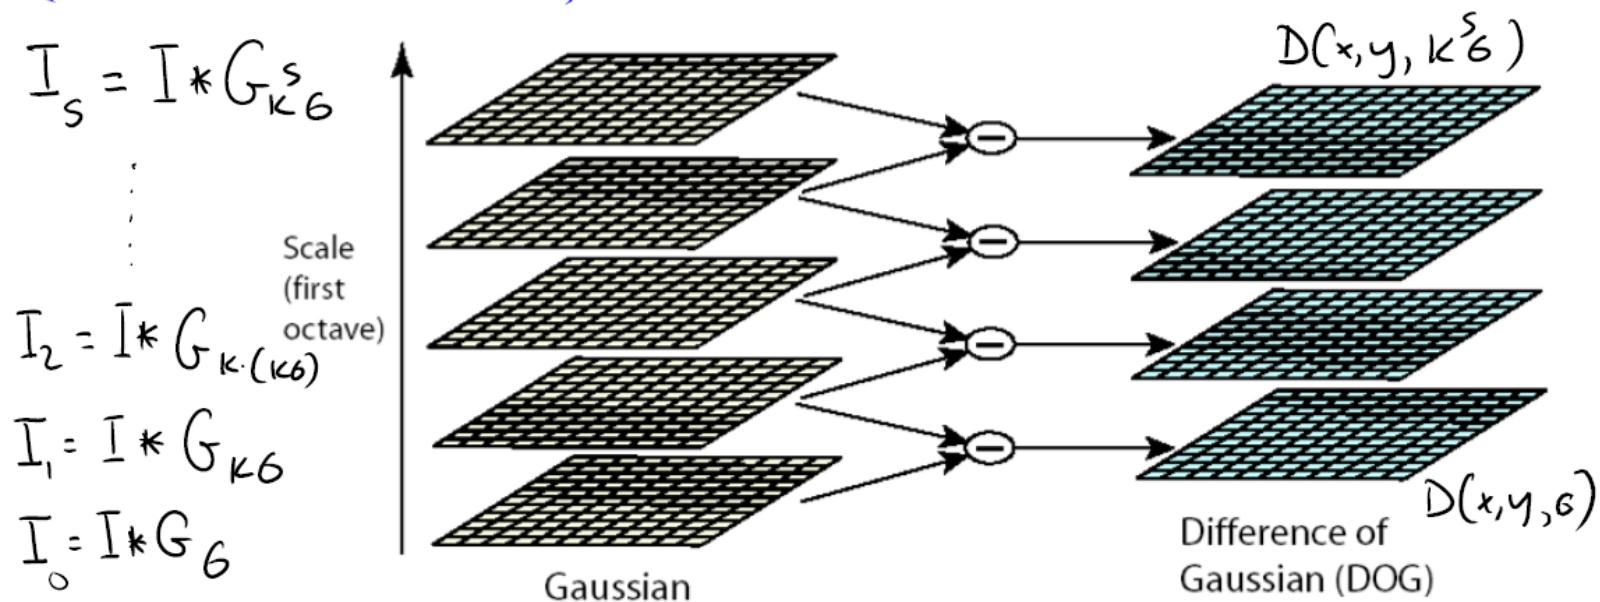
Extremum pruning

Location refinement

## Step 1c: Detecting DOG Extrema

Step 1c (Extremum detection): Find all pixels that correspond to extrema of  $D(x,y,\rho)$

- Extrema of the image Laplacian are readily distinguishable from their surroundings
- There are usually just a few in each pyramid (a few thousand)



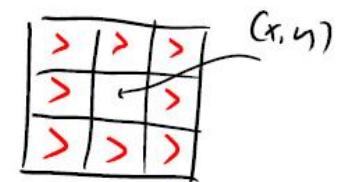
# The Difference-Of-Gaussians (DOG) Filter

Finding  
local extrema  
in a single  
image

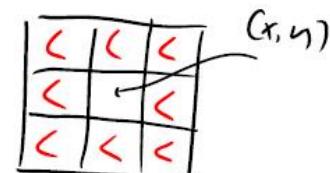
$$D(x, y, \rho)$$



- minimum at  $(x, y)$  if  $D(x, y, \rho) < \text{all neighbours}$



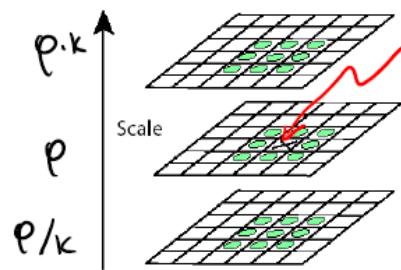
- maximum if  $D(x, y, \rho) > \text{all neighbours}$



## Step 1c: Detecting DOG Extrema

Step 1c (Extremum detection): Find all pixels that correspond to extrema of  $D(x,y,p)$

finding extrema  
in a "stack" of  
images



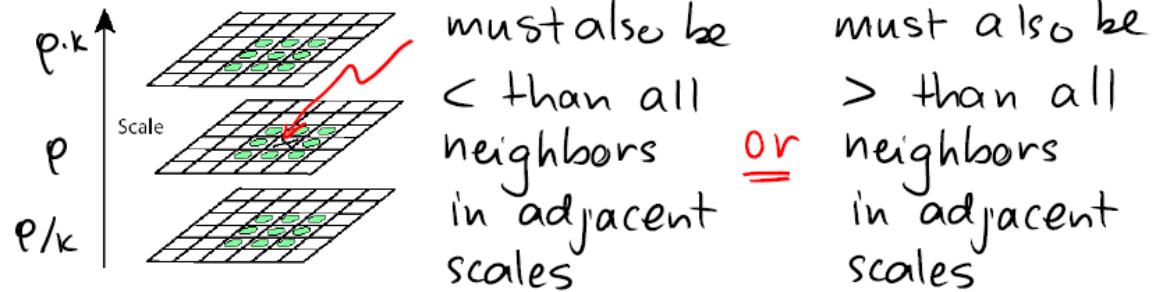
must also be  
< than all  
neighbors      or  
in adjacent  
scales

must also be  
> than all  
neighbors  
in adjacent  
scales

## Step 1c: Detecting DOG Extrema

Step 1c (Extremum detection): Find all pixels that correspond to extrema of  $D(x,y,p)$

finding extrema  
in a "stack" of  
images

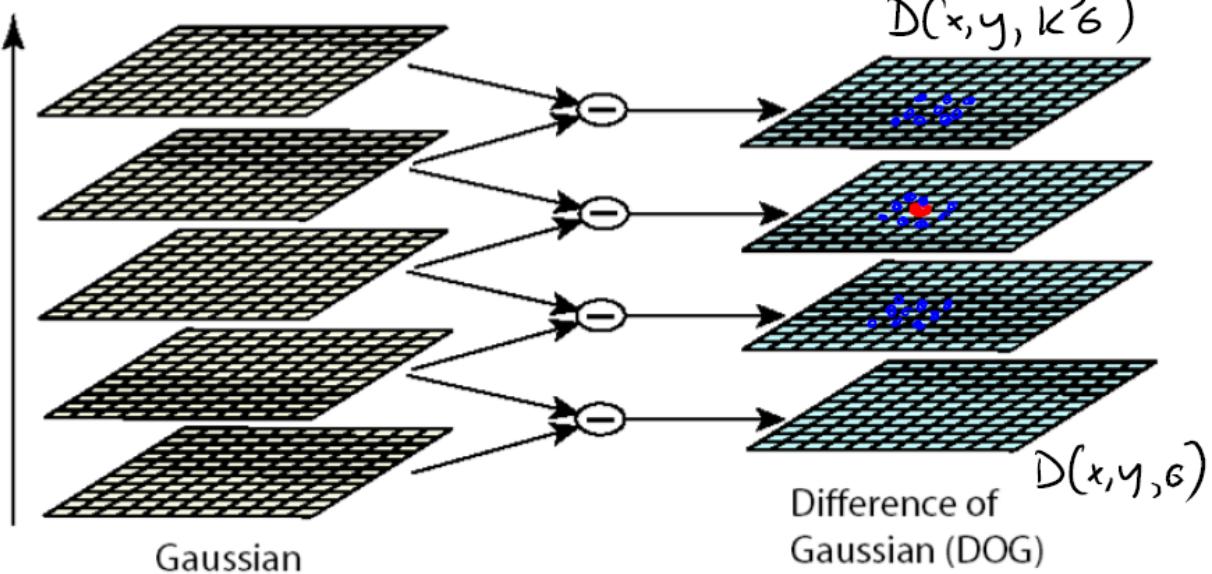


$$I_s = I * G_{K6}^s$$

$$I_2 = I * G_{K(K6)}$$

$$I_1 = I * G_{K6}$$

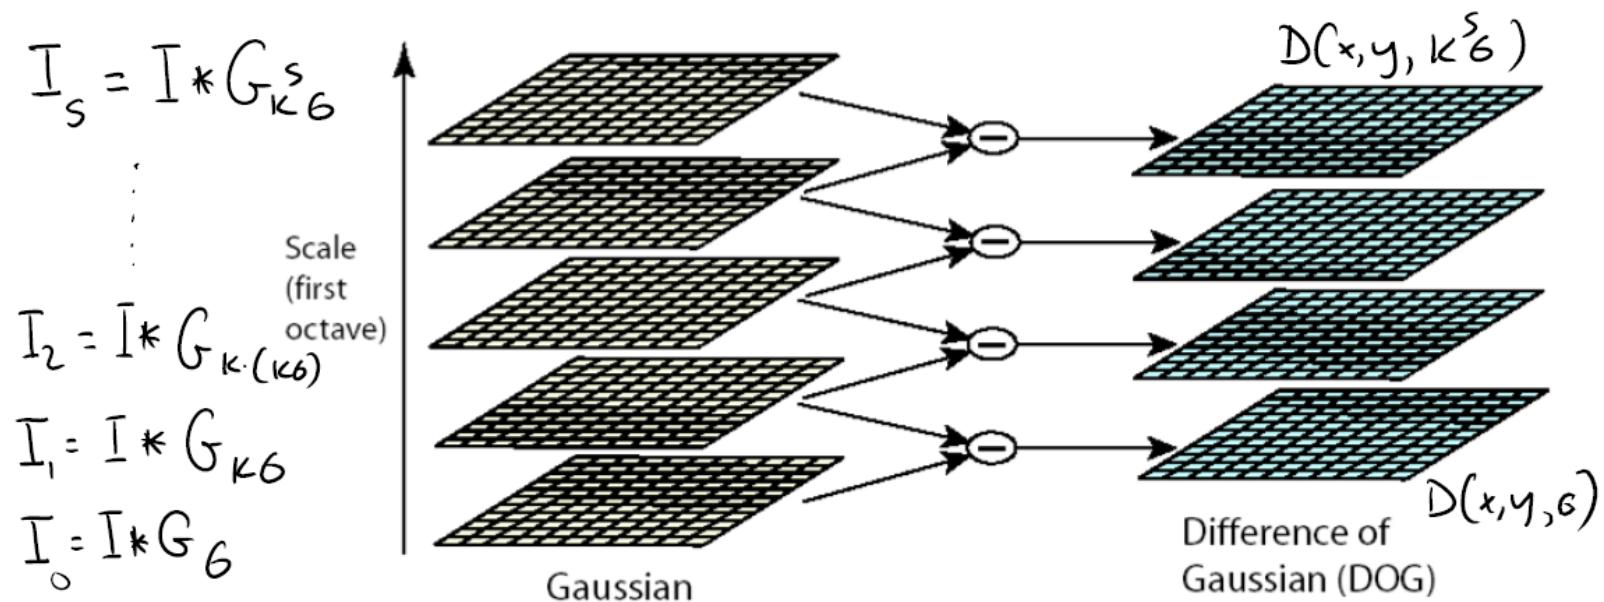
$$I_0 = I * G_6$$



## Step 1c: Detecting DOG Extrema: Algorithm

Step 1c (Extremum detection): Find all pixels that correspond to extrema of  $D(x,y,\rho)$

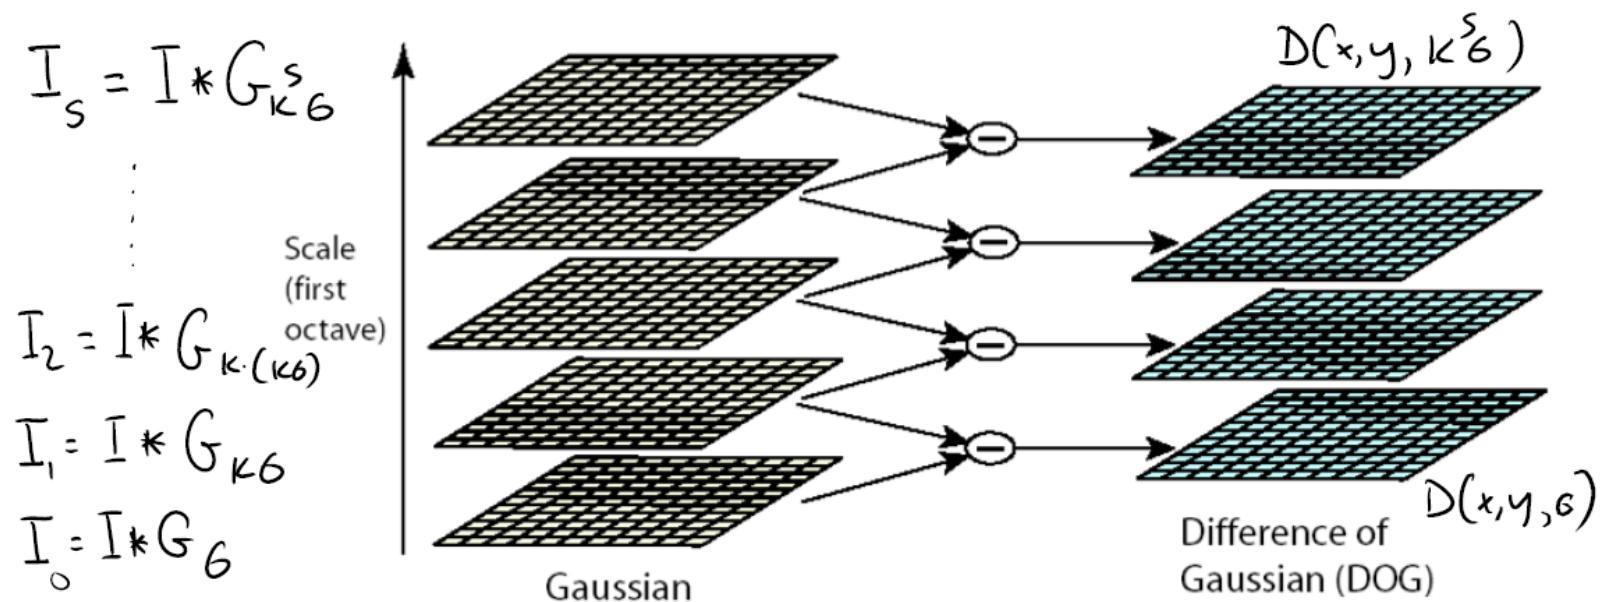
- for each  $(x,y,\rho)$ , check whether  $D(x,y,\rho)$  is greater than (or smaller than) all of its neighbours in current scale and adjacent scales above & below



## Step 1c: SIFT Keypoints = DOG Extrema

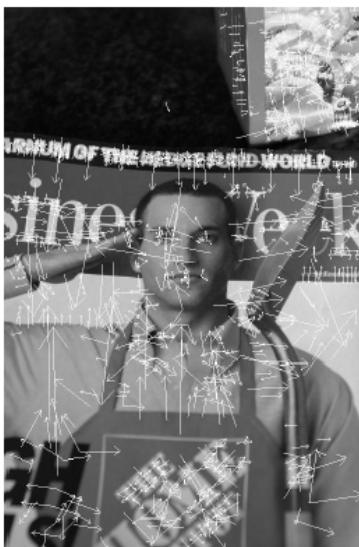
Step 1c (Extremum detection): Find all pixels that correspond to extrema of  $D(x,y,p)$

An extremum detected at  $D(x,y,p)$  defines the keypoint  $(x,y,p)$



# Computing SIFT Keypoints: Basic Steps

Source image I



Gauss-pyramid

Step 1a  
Build pyramid  
of Gauss-  
smoothed  
images

DOG pyramid

Step 1b  
Build DOG  
pyramid

DOG extrema

Step 1c  
Locate extrema  
of DOG  
pyramid  
 $(x_i, y_i, p_i)$

Step 1f

Assign  
orientation  
to extrema  
 $p'_i = (x_i, y_i, p_i, \theta_i)$

Step 1e

Prune set of  
extrema  
Keypoints = {  
all remaining  
 $(x_i, y_i, p_i)$  }

Step 1d

Refine location  
of DOG  
extrema  
 $(x_i, y_i, p_i) \rightarrow (x_i, y_i, p'_i)$

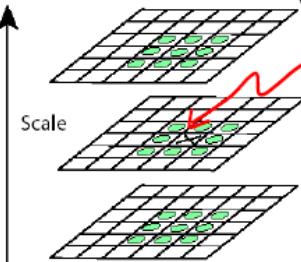
Orientation assign

Extremum pruning

Location refinement

## Step 1d: Refining Location of Extrema

Step 1d (Extremum localization) Refine the location of detected extrema through a quadratic least-sq fit



① 2<sup>nd</sup> order Taylor expansion of  $D$  at  $(x, y, p)$ :

$$D(\Delta x, \Delta y, \Delta p) = D(x, y, p) +$$

$$\frac{\partial D}{\partial x} \cdot \Delta x + \frac{\partial D}{\partial y} \cdot \Delta y + \frac{\partial D}{\partial p} \cdot \Delta p +$$

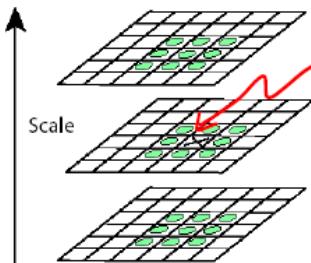
$$\frac{1}{2} [\Delta x \ \Delta y \ \Delta p] \begin{bmatrix} \frac{\partial^2 D}{\partial x^2} & \frac{\partial^2 D}{\partial x \partial y} & \frac{\partial^2 D}{\partial x \partial p} \\ \frac{\partial^2 D}{\partial x \partial y} & \frac{\partial^2 D}{\partial y^2} & \frac{\partial^2 D}{\partial y \partial p} \\ \frac{\partial^2 D}{\partial x \partial p} & \frac{\partial^2 D}{\partial y \partial p} & \frac{\partial^2 D}{\partial p^2} \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta p \end{bmatrix}$$

\* all derivatives approximated using finite differences, e.g.

$$\frac{\partial D}{\partial x} \approx D(x+1, y, p) - D(x, y, p)$$

# Step 1d: Refining Location of Extrema

①



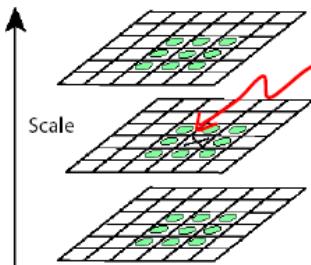
2<sup>nd</sup> order Taylor expansion of  $D$  at  $(x, y, p)$ :  
In vector notation

$$\vec{x} = \begin{bmatrix} x \\ y \\ p \end{bmatrix} \quad \Delta \vec{x} = \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta p \end{bmatrix}$$

$$D(\Delta \vec{x}) = D(\vec{x}) + \left( \frac{\partial D}{\partial \vec{x}} \right)^T \cdot \Delta \vec{x} \\ + \frac{1}{2} (\Delta \vec{x})^T \cdot \frac{\partial^2 D}{\partial \vec{x}^2} \cdot (\Delta \vec{x})$$

## Step 1d: Refining Location of Extrema

①



2<sup>nd</sup> order Taylor expansion of  $D$  at  $(x, y, p)$ :

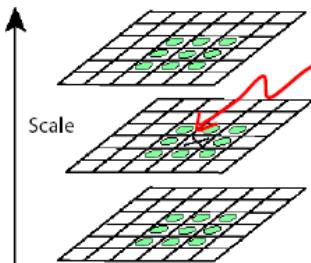
$$D(\Delta \vec{x}) = D(\vec{x}) + \left( \frac{\partial D}{\partial \vec{x}} \right)^T \cdot \Delta \vec{x} + \frac{1}{2} (\Delta \vec{x})^T \cdot \frac{\partial^2 D}{\partial \vec{x}^2} \cdot (\Delta \vec{x})$$

② Take derivatives with respect to  $\Delta \vec{x}$

$$\frac{\partial D}{\partial (\Delta \vec{x})} = \left( \frac{\partial D}{\partial \vec{x}} \right)^T + \left( \frac{\partial^2 D}{\partial \vec{x}^2} \right) (\Delta \vec{x})$$

## Step 1d: Refining Location of Extrema

①



2<sup>nd</sup> order Taylor expansion of  $D$  at  $(x, y, p)$ :

$$D(\Delta \vec{x}) = D(\vec{x}) + \left( \frac{\partial D}{\partial \vec{x}} \right)^T \cdot \Delta \vec{x} + \frac{1}{2} (\Delta \vec{x})^T \cdot \frac{\partial^2 D}{\partial \vec{x}^2} \cdot (\Delta \vec{x})$$

② Take derivatives with respect to  $\Delta \vec{x}$

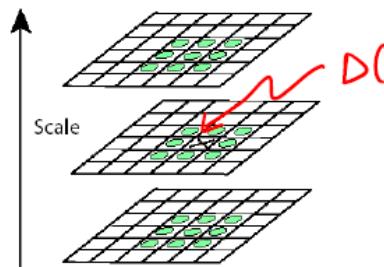
$$\frac{\partial D}{\partial (\Delta \vec{x})} = \left( \frac{\partial D}{\partial \vec{x}} \right) + \left( \frac{\partial^2 D}{\partial \vec{x}^2} \right) (\Delta \vec{x})$$

③ Extremum  $\Leftrightarrow$  derivative is zero  $\Rightarrow$  solve for  $\frac{\partial D}{\partial \Delta \vec{x}} = 0$

$$(\Delta \vec{x}) = - \left( \frac{\partial^2 D}{\partial \vec{x}^2} \right)^{-1} \cdot \left( \frac{\partial D}{\partial \vec{x}} \right)$$

## Step 1d: Refining Location of Extrema

Step 1d (Extremum localization) Refine the location of detected extrema through a quadratic least-sq fit



Location refinement:

extremum  $(x, y, p)$

moved to

$(x + \Delta x, y + \Delta y, p + \Delta p)$

where

$$(\vec{\Delta x}) = \left( \frac{\partial^2 D}{\partial \vec{x}^2} \right)^{-1} \cdot \left( \frac{\partial D}{\partial \vec{x}} \right)$$

$$\frac{\partial D}{\partial \vec{x}^2} = \begin{bmatrix} \frac{\partial^2 D}{\partial x^2} & \frac{\partial^2 D}{\partial xy} & \frac{\partial^2 D}{\partial xz} \\ \frac{\partial^2 D}{\partial xy} & \frac{\partial^2 D}{\partial y^2} & \frac{\partial^2 D}{\partial yz} \\ \frac{\partial^2 D}{\partial xz} & \frac{\partial^2 D}{\partial yz} & \frac{\partial^2 D}{\partial z^2} \end{bmatrix}$$

$$\vec{\Delta x} = \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta p \end{bmatrix}$$

$$\frac{\partial D}{\partial \vec{x}} = \begin{bmatrix} \frac{\partial D}{\partial x} \\ \frac{\partial D}{\partial y} \\ \frac{\partial D}{\partial p} \end{bmatrix}$$

# Computing SIFT Keypoints: Basic Steps

Source image I



Gauss-pyramid

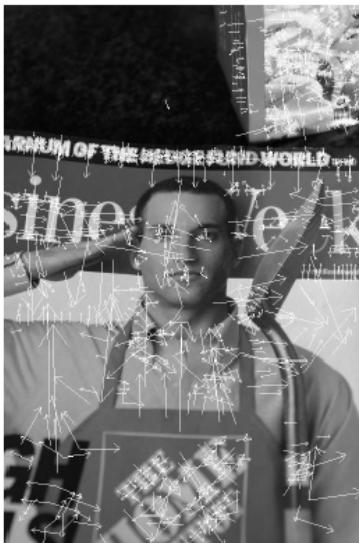
Step 1a  
Build pyramid  
of Gauss-  
smoothed  
images

DOG pyramid

Step 1b  
Build DOG  
pyramid

DOG extrema

Step 1c  
Locate extrema  
of DOG  
pyramid  
 $(x_i, y_i, p_i)$



Step 1f

Assign  
orientation  
to extrema  
 $p_i = (x_i, y_i, p_i, \theta_i)$

Step 1e

Prune set of  
extrema  
Keypoints = {  
all remaining  
 $(x_i, y_i, p_i)$  }

Step 1d

Refine location  
of DOG  
extrema  
 $(x_i, y_i, p_i) \rightarrow (x_i, y_i, p'_i)$

Orientation assign

Extremum pruning

Location refinement

## Step 1e: Pruning “Insignificant” Extrema

Step 1e (Extremum pruning): Prune all extrema that are weak or that correspond to edges

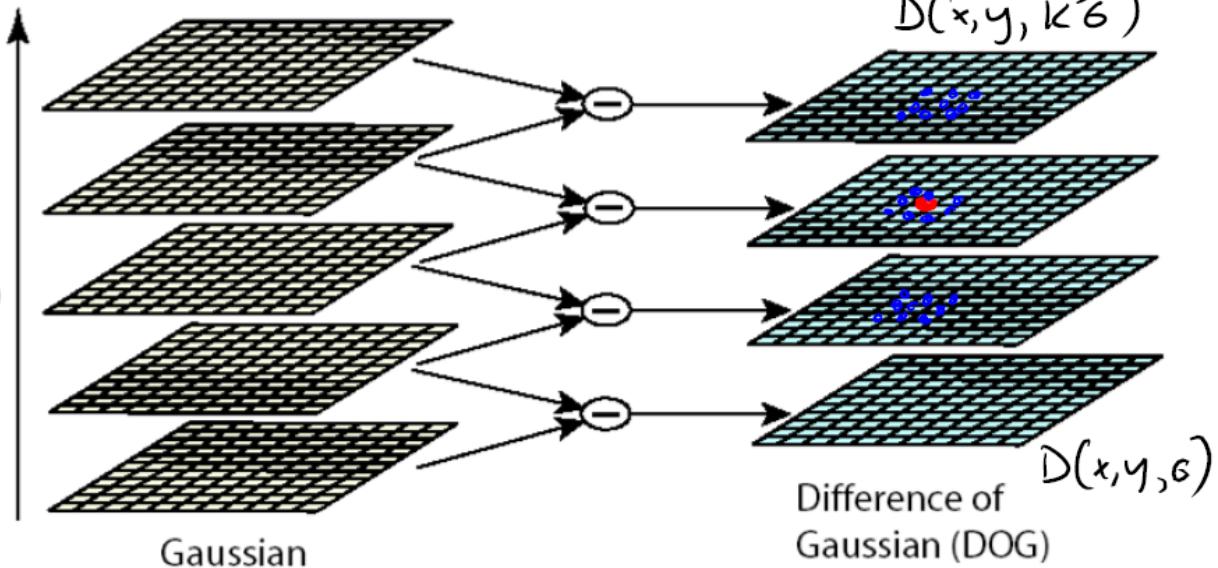
Condition for detecting a “strong” extremum  $(x'_i, y'_i, p'_i)$ :

$$|D(x'_i, y'_i, p'_i)| = \text{large}$$

in practice  $> 0.03$

assumes image  $J$  has pixel intensities in the range  $[0,1]$

$$I_s = I * G_{K_6}^s$$



$$I_2 = I * G_{K_6(K_6)}$$

$$I_1 = I * G_{K_6}$$

$$I_0 = I * G_6$$

Scale  
(first  
octave)

Gaussian

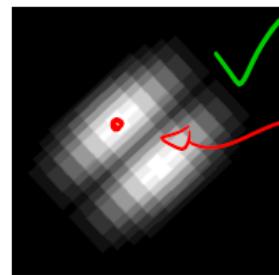
Difference of  
Gaussian (DOG)

# Step 1e: Pruning “Insignificant” Extrema

Step 1e (Extremum pruning): Prune all extrema that are weak or that correspond to edges

corner-like extremum

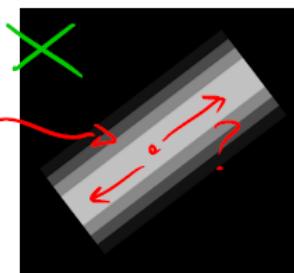
position is  
is well-  
constrained



keep

edge-like extremum

Prune X



Position along  
edge is not  
constrained

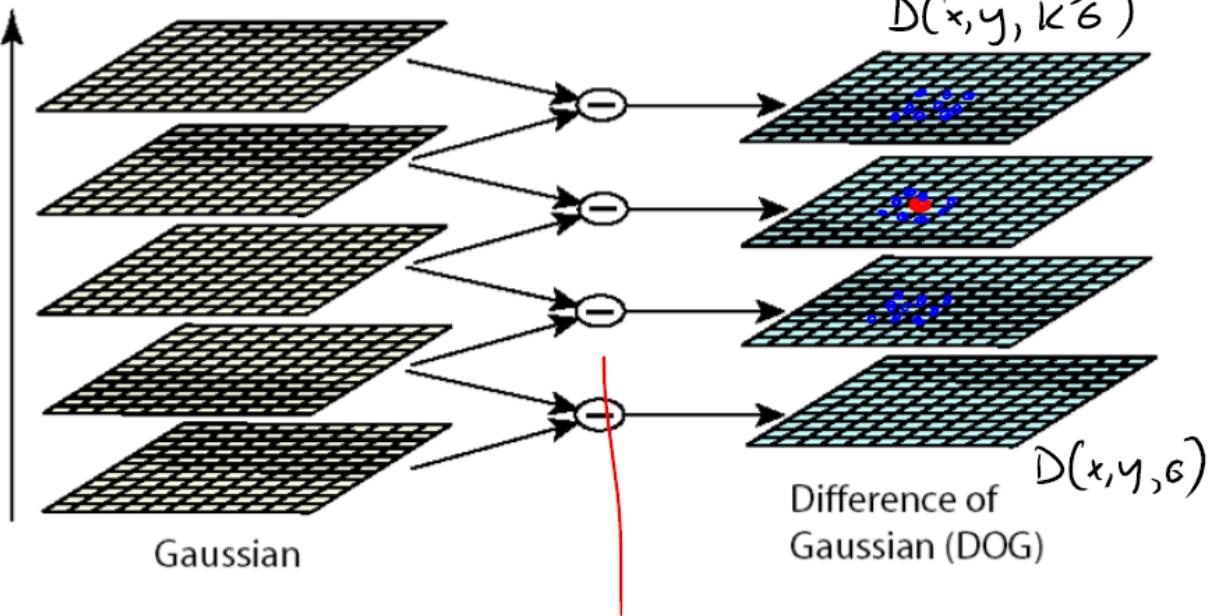
$$I_s = I * G_{K_6}^s$$

:

$$I_2 = I * G_{K_6(K_6)}$$

$$I_1 = I * G_{K_6}$$

$$I_0 = I * G_6$$

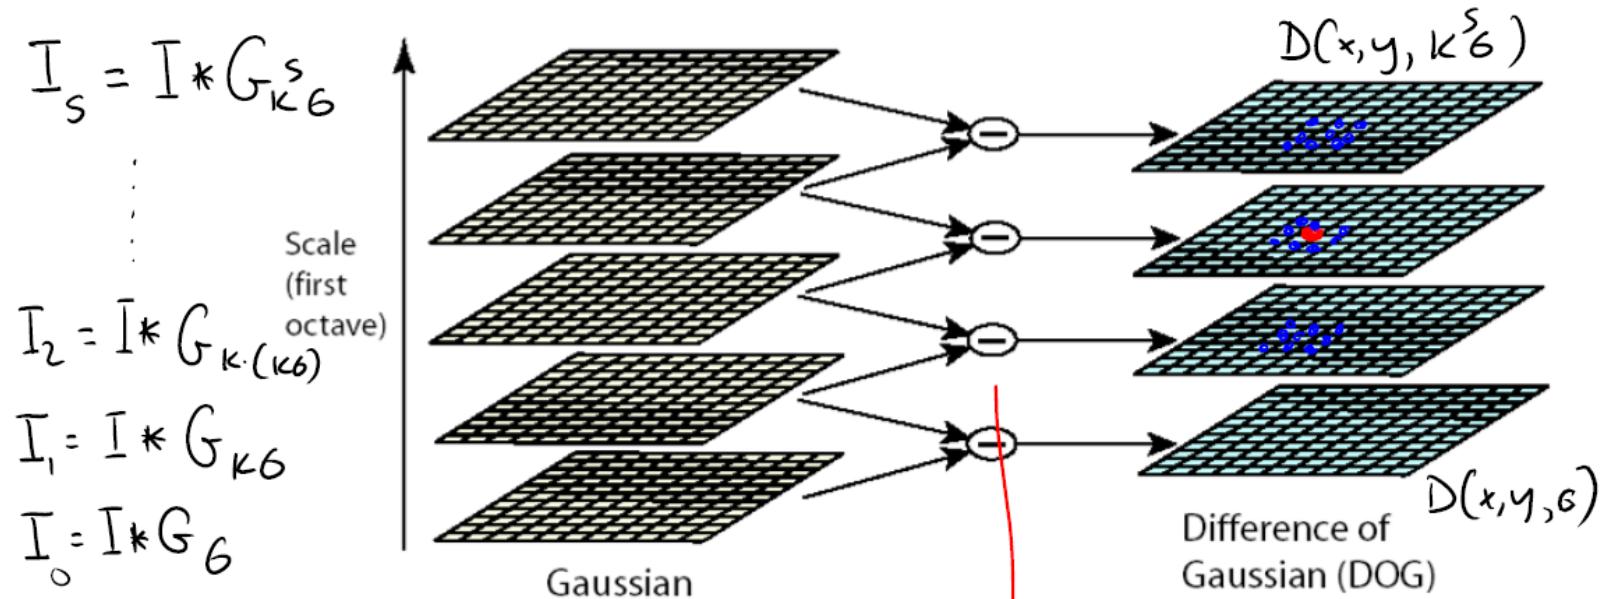


# Step 1e: Pruning “Insignificant” Extrema

Step 1e (Extremum pruning): Prune all extrema that are weak or that correspond to edges

$$H = \begin{bmatrix} \frac{\partial^2 D}{\partial x^2} & \frac{\partial^2 D}{\partial x \partial y} \\ \frac{\partial^2 D}{\partial x \partial y} & \frac{\partial^2 D}{\partial y^2} \end{bmatrix}$$

- Compute hessian  $H$  of  $D(x, y, e_i')$  at  $(x, y) = (x_i', y_i')$
- Prune if  $\frac{\text{Tr}(H)}{\text{Det}(H)} > \left(\frac{11}{10}\right)^2$



# Computing SIFT Keypoints: Basic Steps

Source image I



Gauss-pyramid

Step 1a

Build pyramid  
of Gauss-  
smoothed  
images



DOG pyramid

Step 1b

Build DOG  
pyramid



DOG extrema

Step 1c

Locate extrema  
of DOG  
pyramid  
 $(x_i, y_i, p_i)$



Step 1f

Assign  
orientation  
to extrema

$$P_i = (x_i, y_i, p_i, \theta_i)$$



Step 1e

Prune set of  
extrema

$$\text{Keypoints} = \{ \text{all remaining } (x_i, y_i, p_i) \}$$



Step 1d

Refine location  
of DOG  
extrema  
 $(x_i, y_i, p_i) \rightarrow (x_i, y_i, p'_i)$



Orientation assign

Extremum pruning

Location refinement

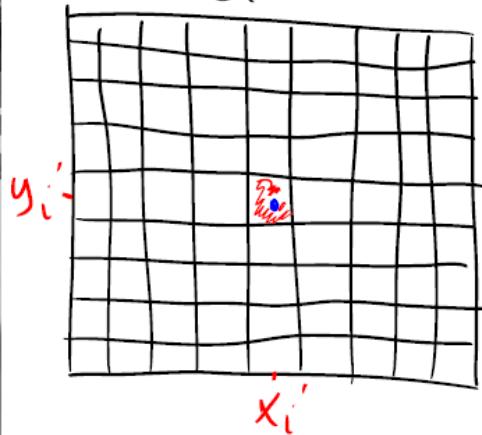
# Step 1f: Keypoint Orientation Assignment

Assigning an orientation  $\theta_i$  to keypoint  $(x'_i, y'_i, p'_i)$ :

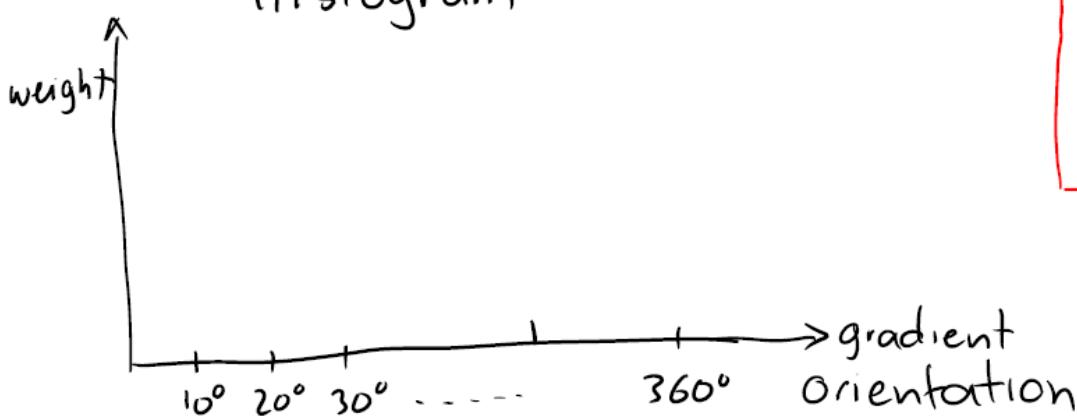
$$I * G_{p'_i}$$



Gradient direction



Histogram

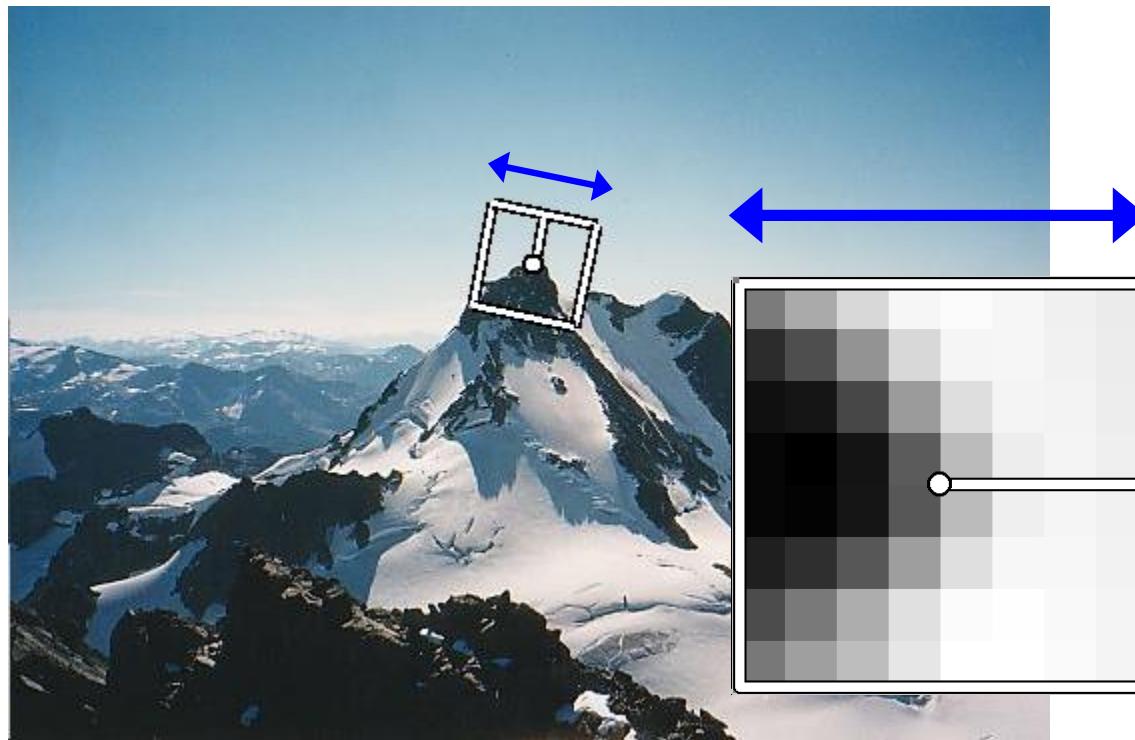


- A. Compute smoothed image  $I * G_{p'_i}$
- B. Compute gradient magnitude & orientation in neighborhood of  $(x'_i, y'_i)$  in  $I * G_{p'_i}$

- C. Compute histogram of orientations

- D. Assigned orientation  $\theta_i$  = highest peak in hist

# Making descriptor rotation invariant

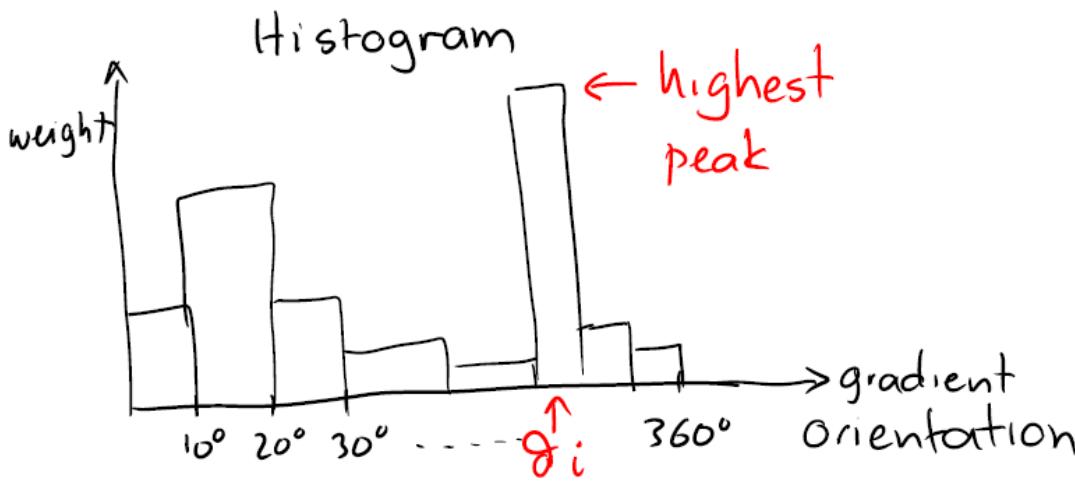
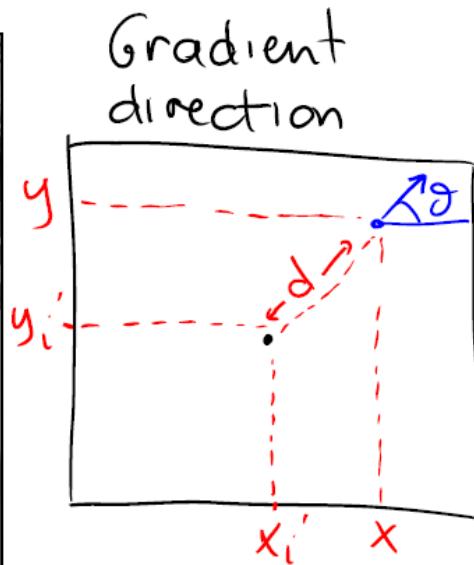


- Rotate patch according to its dominant gradient orientation
- This puts the patches into a canonical orientation.

# Step 1f: Keypoint Orientation Assignment

## Computing Histogram of Orientations

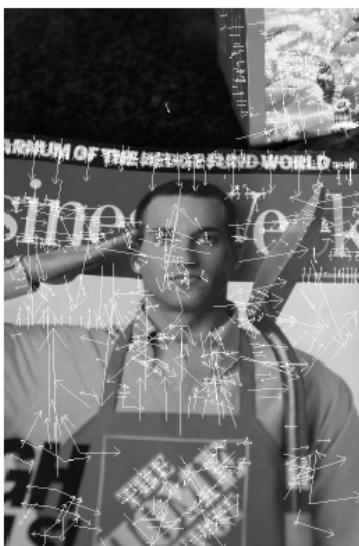
$I * G_{P_i}$



- Orientations divided into 36 bins (one every 10 degrees)
- Pixel  $(x,y)$  contributes to the bin corresponding to the gradient orientation  $\theta$  at  $(x,y)$
- Contribution to the bin is equal to  $|\nabla I(x,y)| \cdot G_{I,s,p_i}(d)$
- Total bin weight = sum of contributions from all pixels

# Computing SIFT Keypoints: Basic Steps

Source image I



Gauss-pyramid

Step 1a  
Build pyramid  
of Gauss-  
smoothed  
images

DOG pyramid

Step 1b  
Build DOG  
pyramid

DOG extrema

Step 1c  
Locate extrema  
of DOG  
pyramid  
 $(x_i, y_i, p_i)$

Step 1f

= Assign  
orientation  
to extrema  
 $p'_i = (x_i, y_i, p_i, \theta_i)$

Step 1e

Prune set of  
extrema  
Keypoints = {  
all remaining  
 $(x_i, y_i, p_i)$  }

Step 1d

Refine location  
of DOG  
extrema  
 $(x_i, y_i, p_i) \rightarrow (x_i, y_i, p'_i)$

Orientation assign

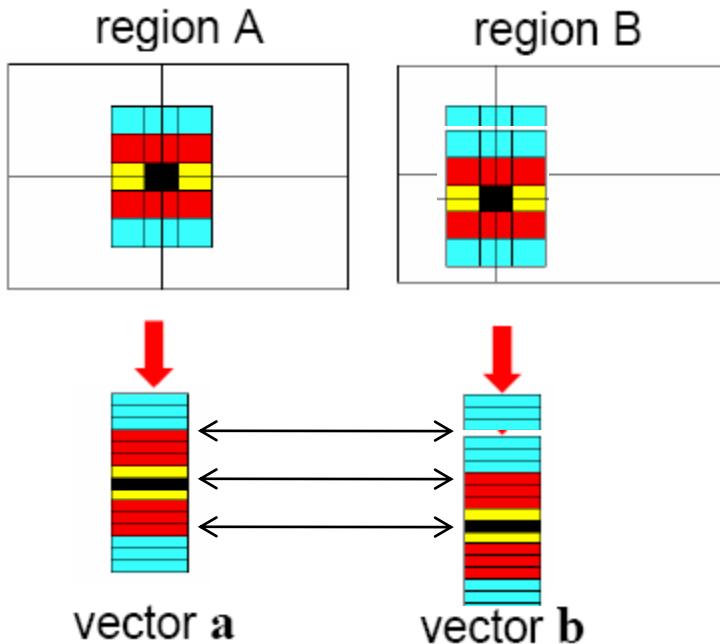
Extremum pruning

Location refinement

# Scale Invariant Feature Transform (SIFT) descriptor [Lowe 2004]

- introduction to the image matching problem
- image matching using SIFT features
- multi-scale interest-point detection
- the SIFT feature detector
- **the SIFT descriptor**

# Raw patches as local descriptors

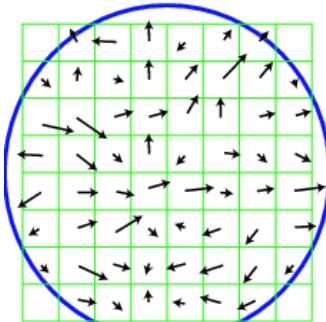


The simplest way to describe the neighborhood around an interest point is to write down the list of intensities to form a feature vector.

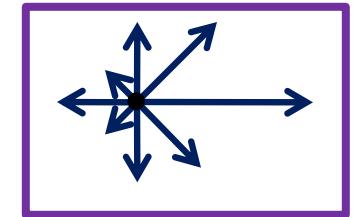
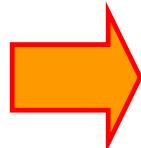
But this is very sensitive to even small shifts, rotations.

# Scale Invariant Feature Transform (SIFT) descriptor [Lowe 2004]

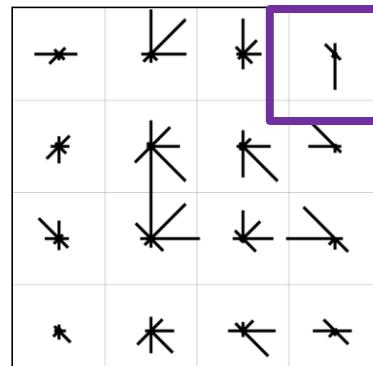
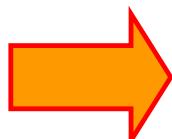
- Use histograms to bin pixels within sub-patches according to their orientation.



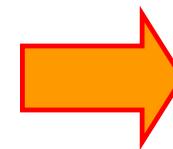
gradients



subdivided local patch



histogram per grid cell



Final descriptor =  
concatenation of all  
histograms, normalized

# The SIFT Keypoints

**Keypoint:** A location  $(x, y)$  in the source image, with an associated orientation & scale, that is "visually distinct" from its surroundings

Input:

Image  $I$

Output:

A set of  $k$  SIFT  
keypoints

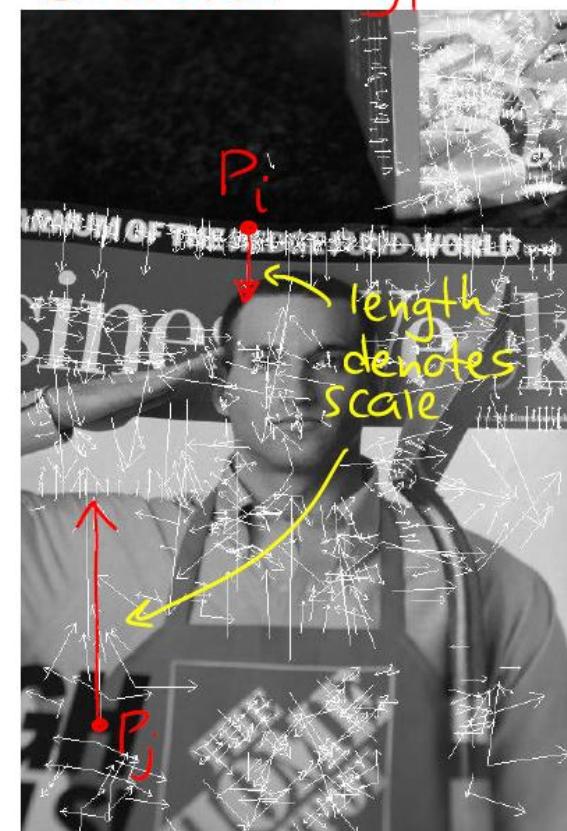
$$P_i = (x_i, y_i, g_i, \theta_i)$$

location  
scale  
orientation

"Source" Image I



Detected keypoints



# The SIFT Feature Vectors

Feature vector (of a keypoint  $P_i$ ): A vector of fixed length that represents the image patch centered at  $P_i$ .

Input:

Image  $I$

Output:

A set of  $K$   
keypoints

$$P_i = (x_i, y_i, g_i, \theta_i)$$

location  
scale  
orientation

A set of  $K$   
feature vectors

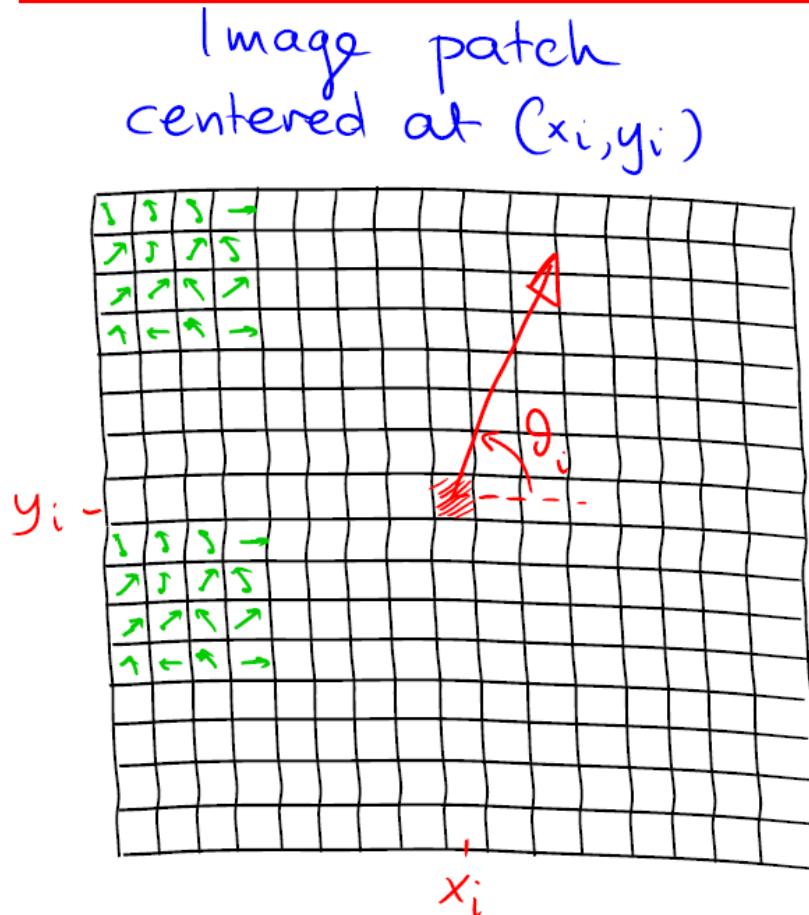
$$f_i = [ \quad | \quad \dots | \quad ]$$

describes the  
Patch centered at  $P_i$

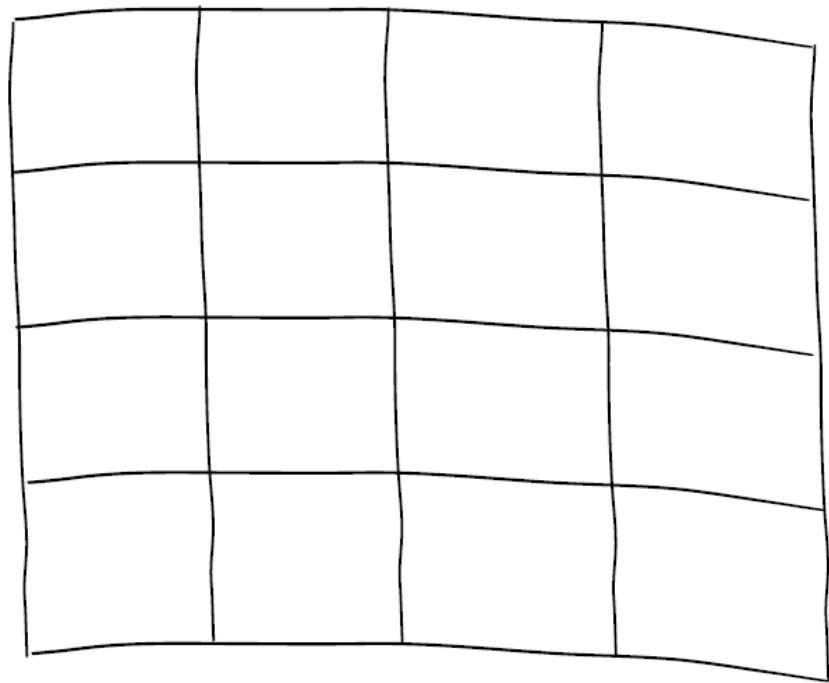
Detected keypoints



# Building the SIFT Descriptor



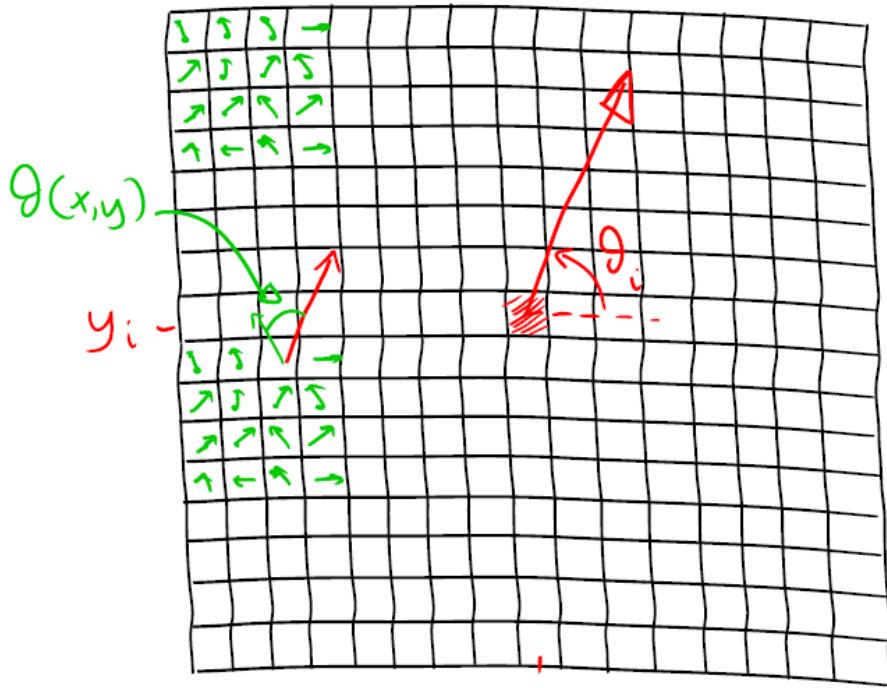
SIFT  
Descriptor



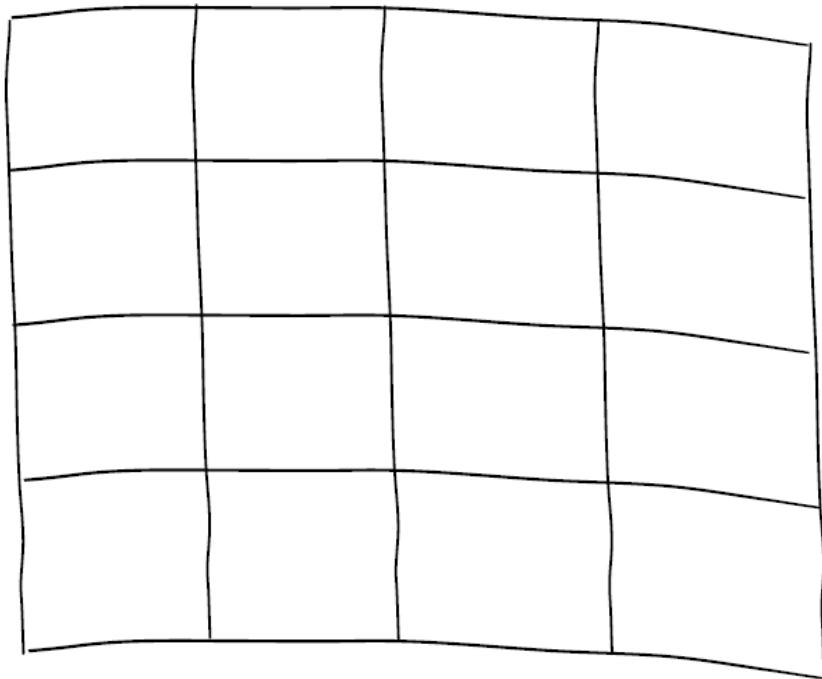
1. Compute gradients in  
16x16 pixel patch of  
image  $I * G_{6i} \leftarrow$  Gaussian-smoothed image  
at scale of the keypoint  
centered at  $(x_i, y_i)$

# Building the SIFT Descriptor

Image patch  
centered at  $(x_i, y_i)$



SIFT  
Descriptor

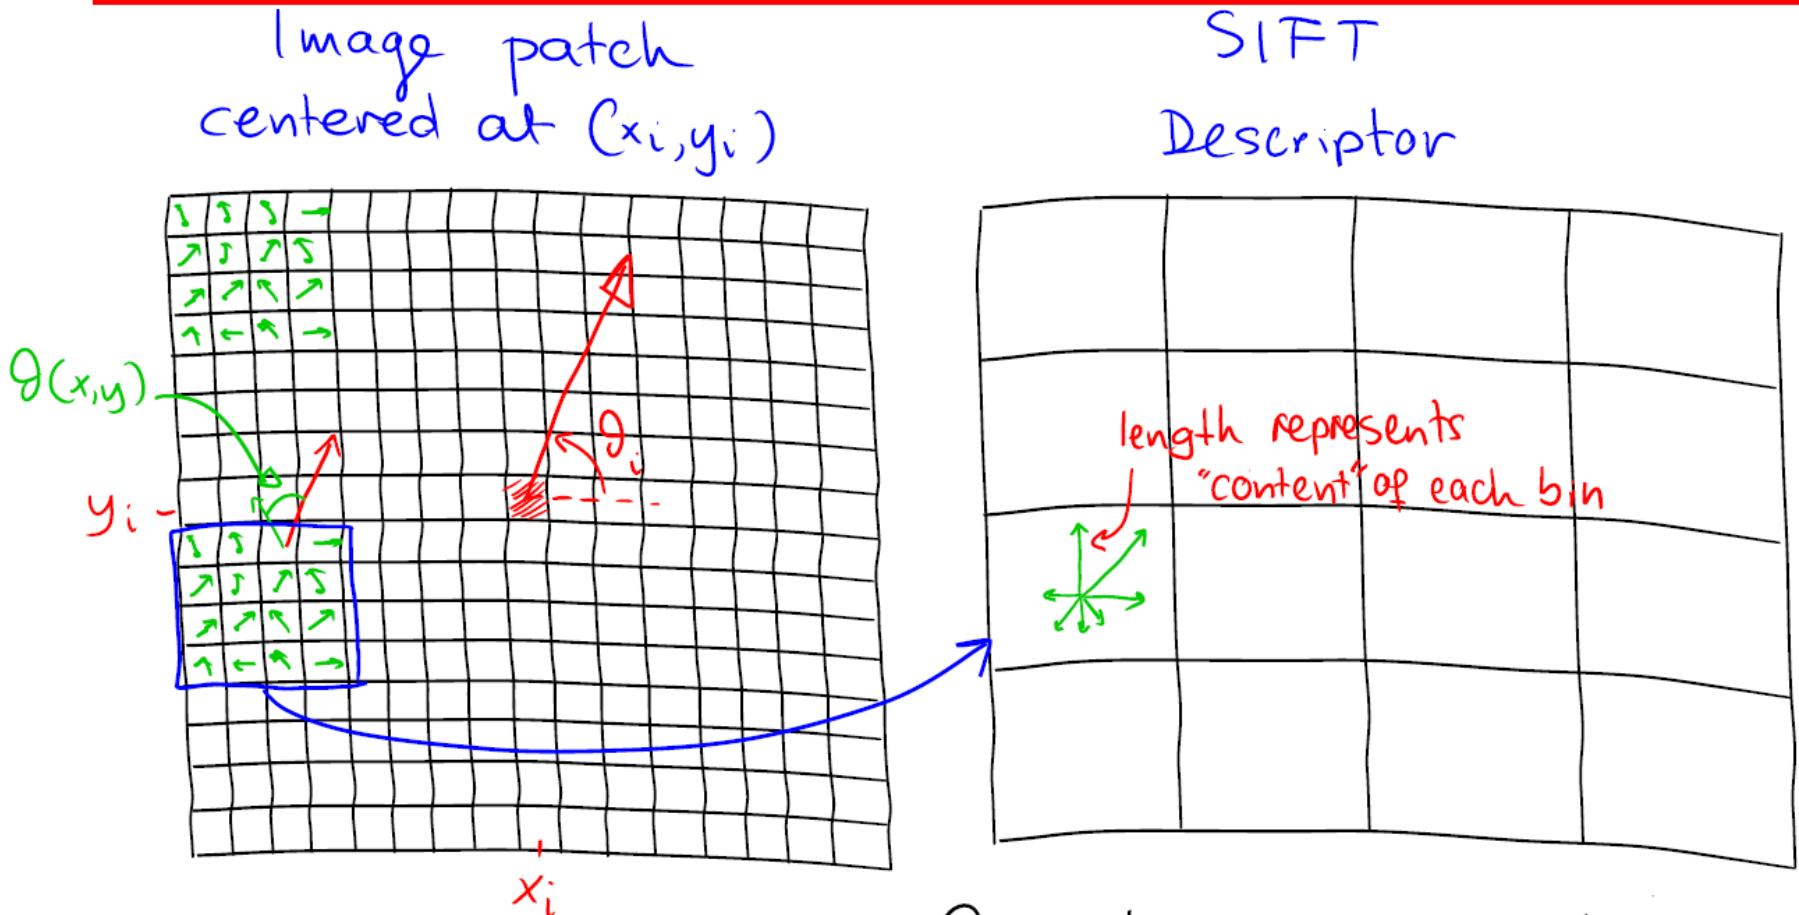


1. Compute gradients in 16x16 pixel patch of image  $I * G_{6i}$  centered at  $(x_i, y_i)$

2. Compute gradient orientation relative to keypoint orientation

$$\theta(x, y) = \tan^{-1} \left[ \frac{\frac{\partial(I * G_{6i})}{\partial y}}{\frac{\partial(I * G_{6i})}{\partial x}} \right] - g_i$$

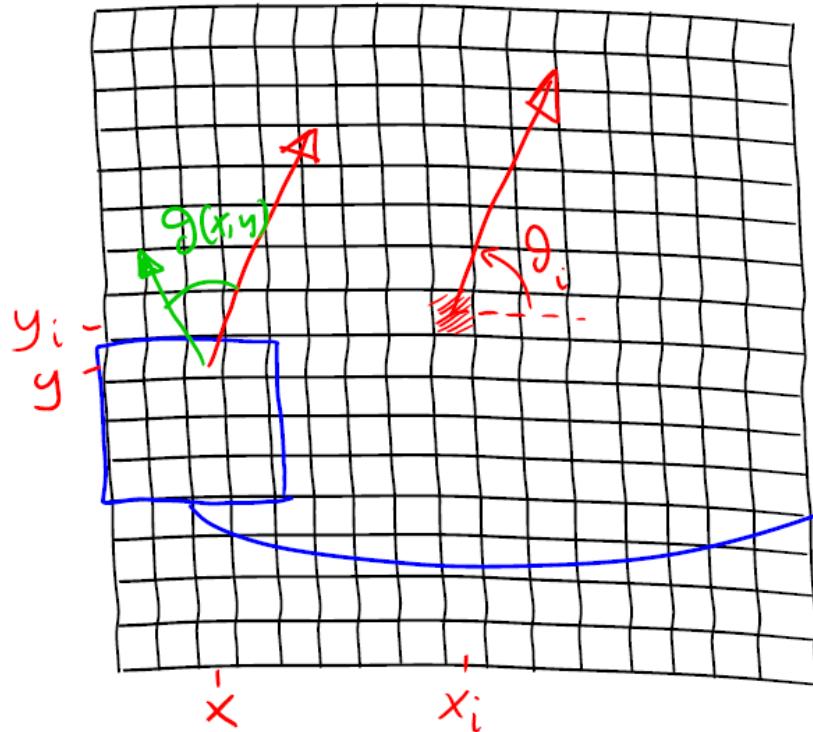
# Building the SIFT Descriptor



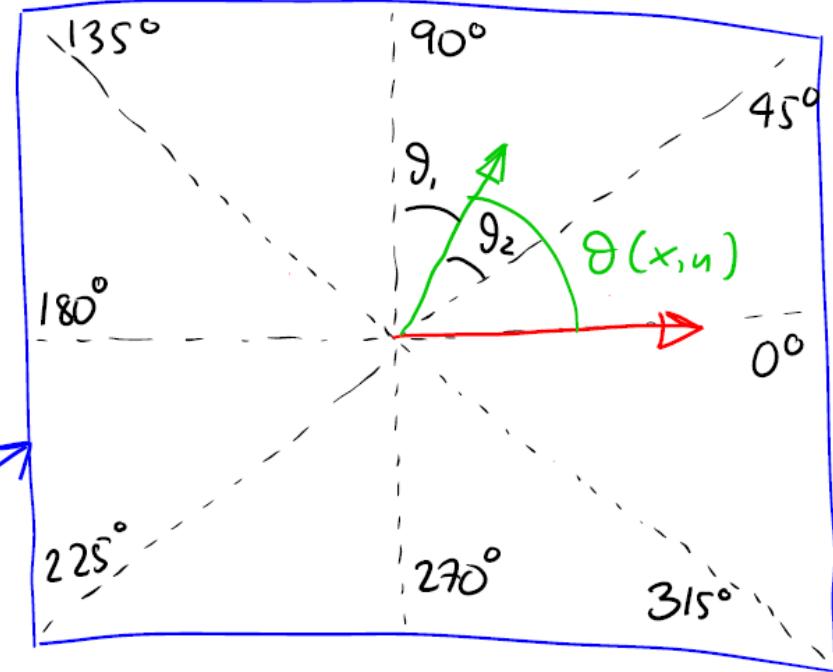
1. Compute gradients
2. Compute relative gradient orientations
3. Compute orientation histogram of each  $4 \times 4$  pixel block  
histogram contains 8 bins, each covering  $45^\circ$

# The 4x4 Orientation Histogram

Image patch  
centered at  $(x_i, y_i)$



The Orientation Histogram  
of a 4x4 pixel block



- Weight of  $(x, y)$ :

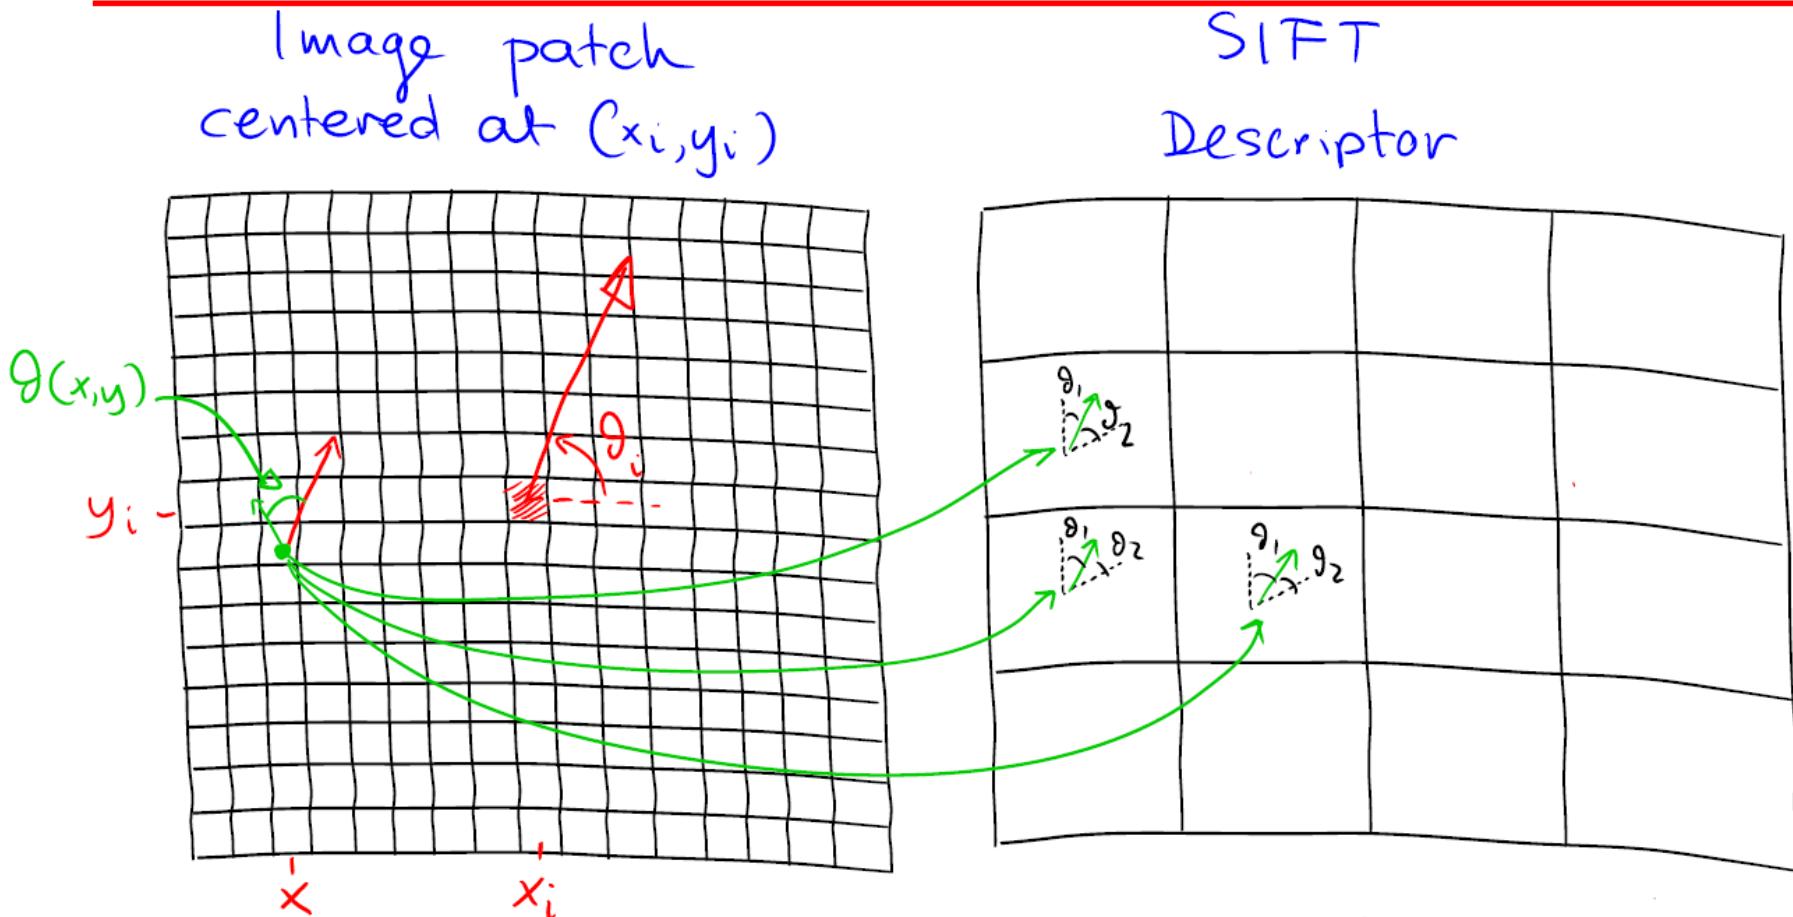
$$w = G_{\frac{G_i}{2}}( \|(x - x_i, y - y_i)\| )$$

⇒ pixels closer to keypoint center have higher weight

- Total contribution of  $(x, y)$  to the orientation histograms:

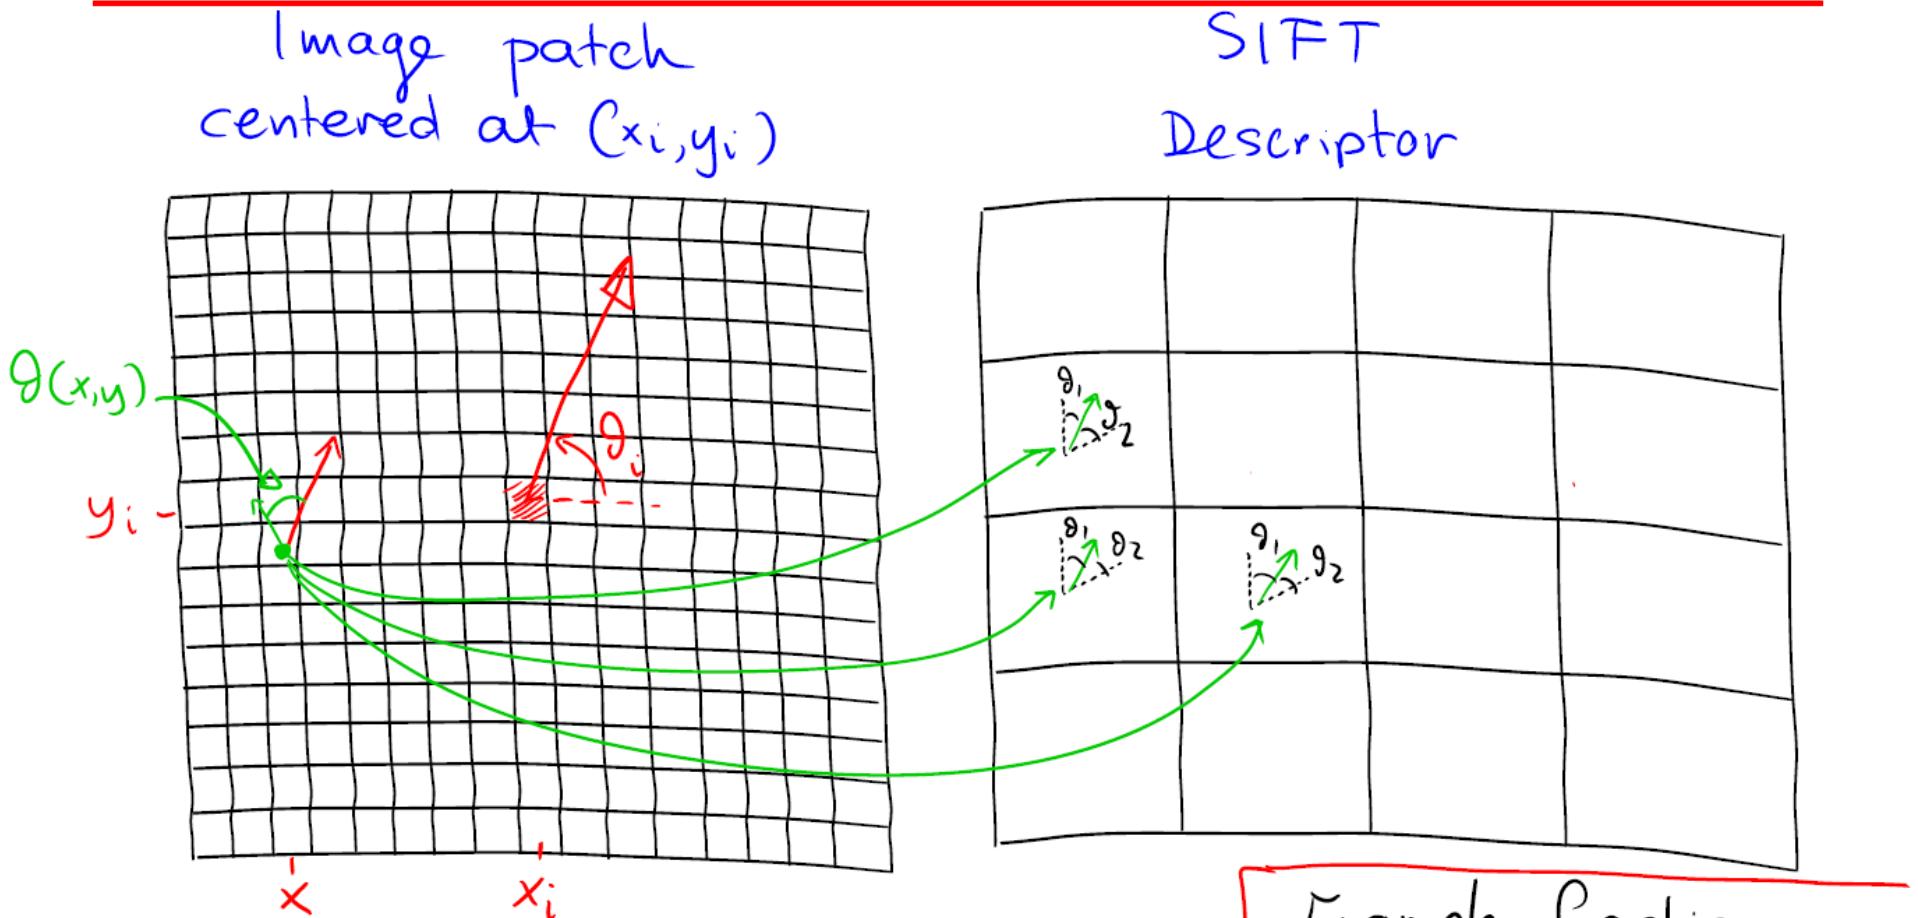
$$C(x, y) = w \cdot \underbrace{\|\nabla I * G_{G_i}(x, y)\|}_{\text{gradient magnitude of smoothed image}}$$

# Building the SIFT Descriptor



- Contribution spread across 2 closest orientations & 3 closest histograms
- Total contribution of  $(x,y)$  to the orientation histograms:  
$$C(x,y) = \omega \cdot \|\nabla I * G_{\sigma_i}(x,y)\|$$
- ⇒ no abrupt changes in histogram if keypoint center displaced by 3-4 pixels

# Building the SIFT Descriptor

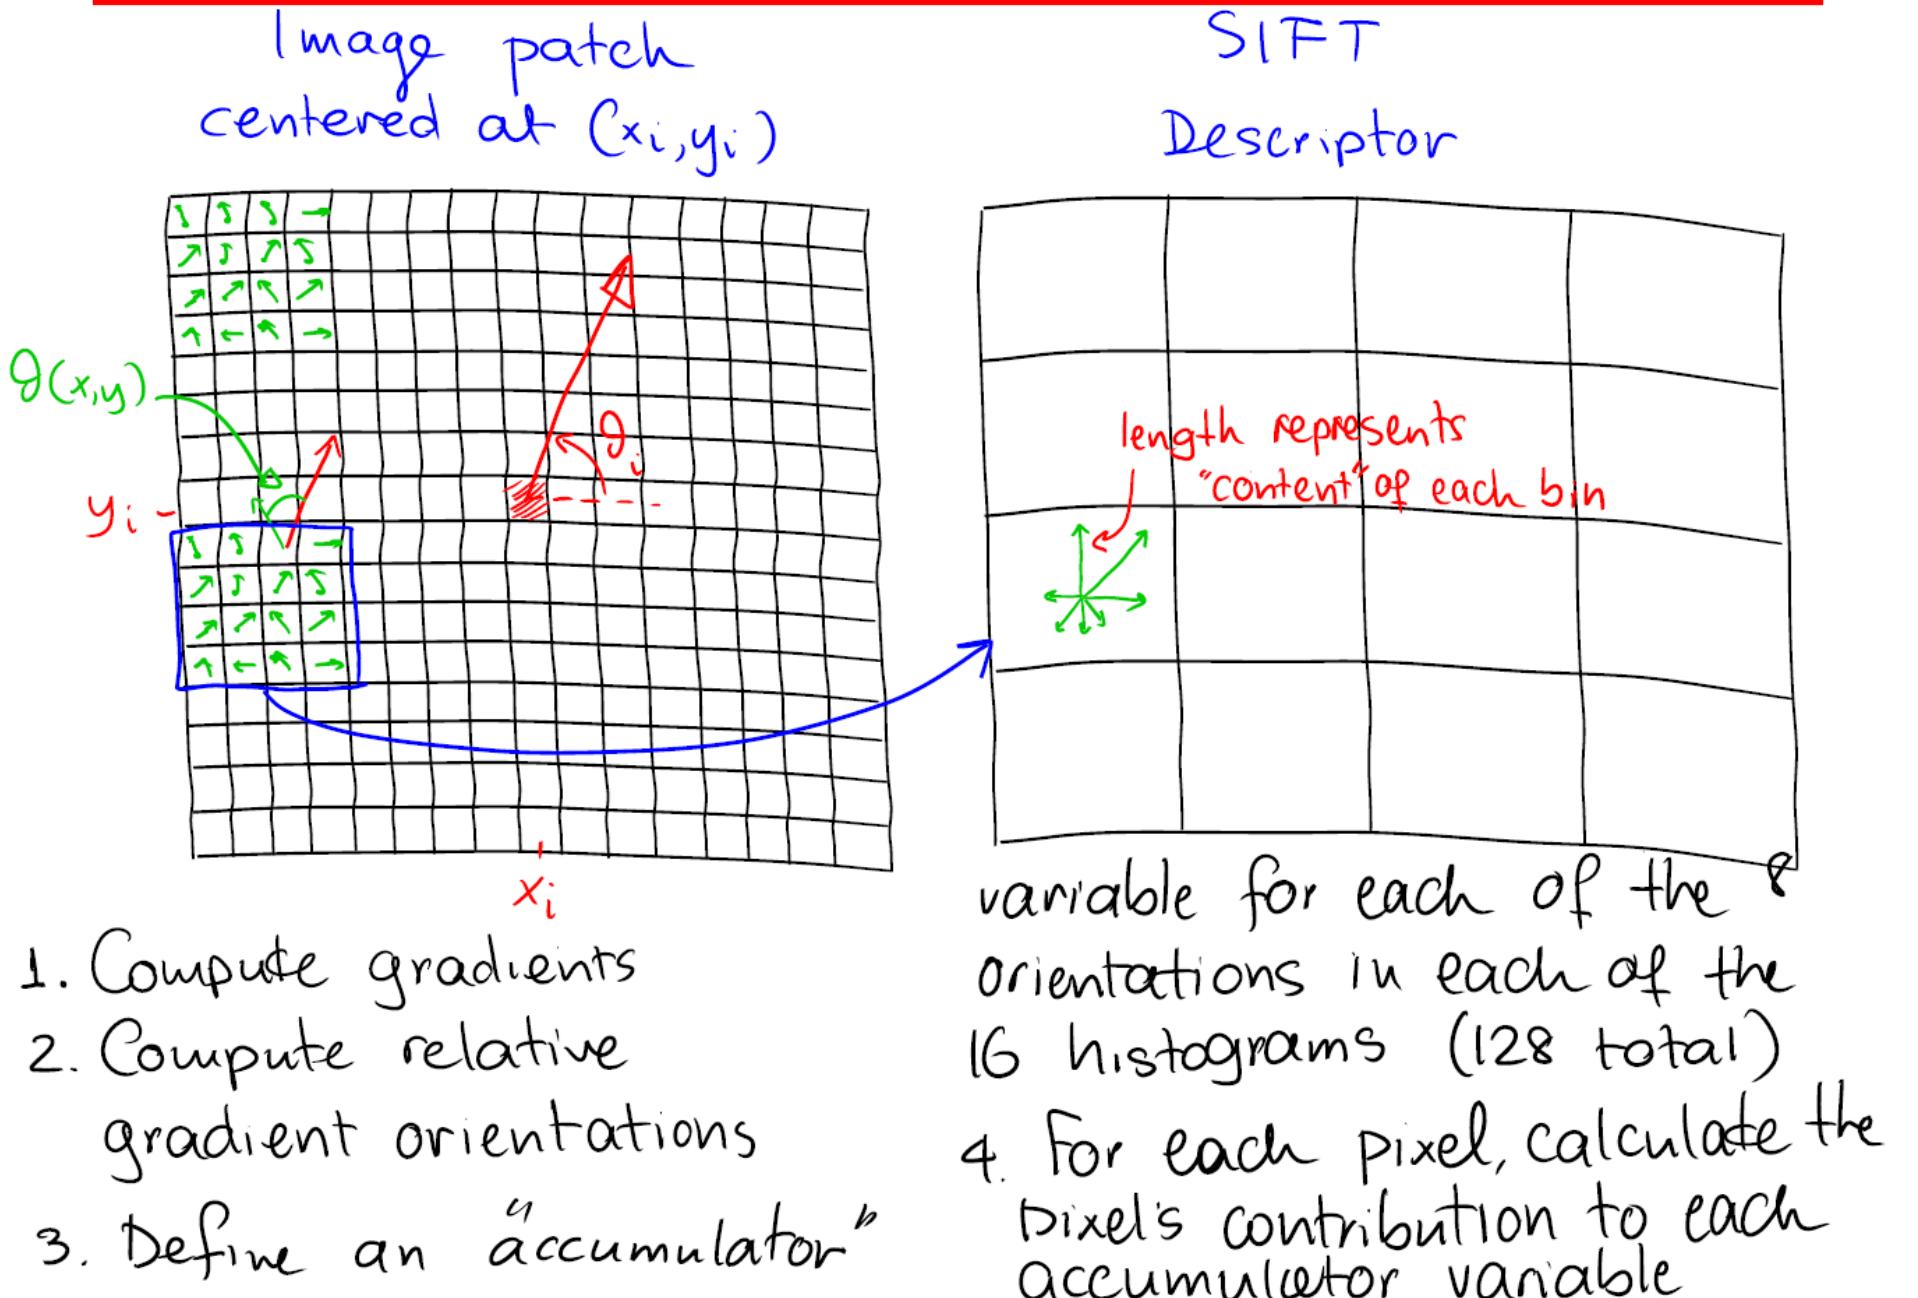


- Contribution spread linearly across 2 closest orientations & 3 closest histograms
- ⇒ no abrupt changes in histogram if keypoint center displaced by 3-4 pixels

Example: fraction allocated to orientation ①:  $\frac{\theta_1}{\theta_1 + \theta_2}$

Orientation ②:  $\frac{\theta_2}{\theta_1 + \theta_2}$

# Building the SIFT Descriptor: Complete Algorithm



# Converting SIFT Descriptors to 128-dim Vectors

Post processing:

- ① Normalize  $f_i$ :

$$f_i \rightarrow \frac{f_i}{\|f_i\|}$$

⇒ gives invariance to linear lighting variations across images, i.e. when matching image  $I$  and image  $a \cdot I + b$  (because  $f_i$  will be the same in both images)

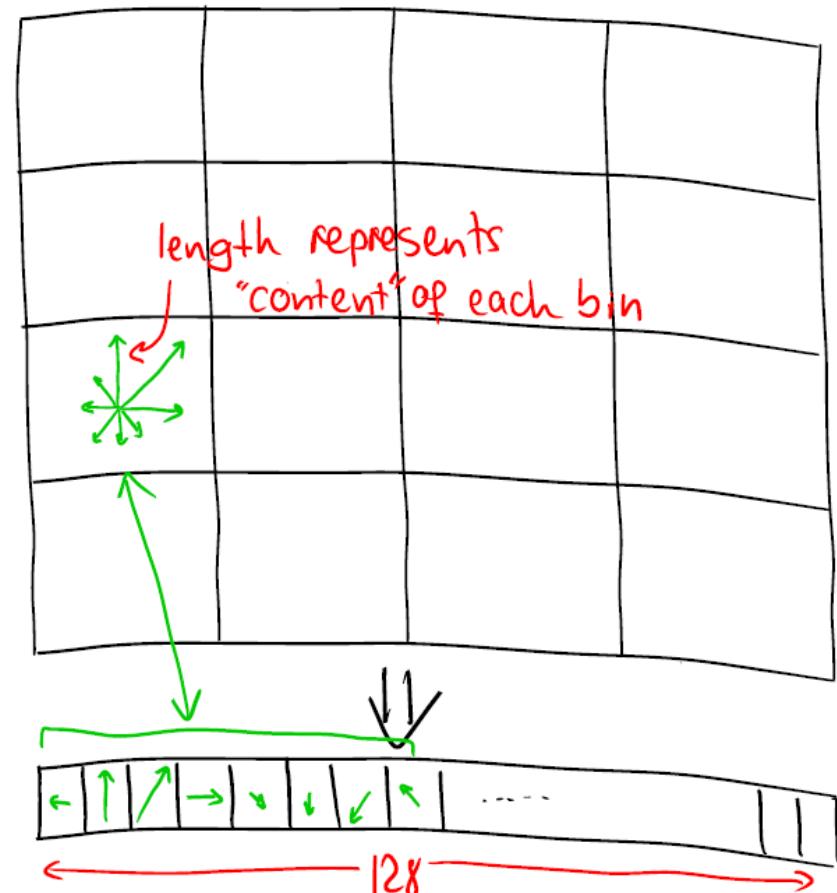
- ② Clamp  $f_i$

Clamp all elements of  $f_i$  at 0.2

⇒ gives less weight to very large gradient magnitudes

SIFT

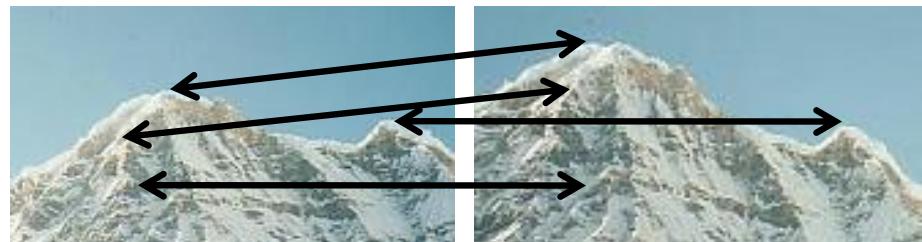
Descriptor



- ③ Re-normalize

# Local features: main components

- 1) Detection: Identify the interest points
- 2) Description: Extract vector feature descriptor surrounding each interest point.
- 3) Matching: Determine correspondence between descriptors in two views



# Matching 2 Images Using SIFT Features

Image 1

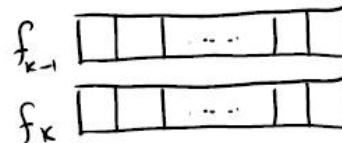
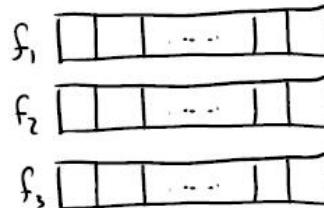
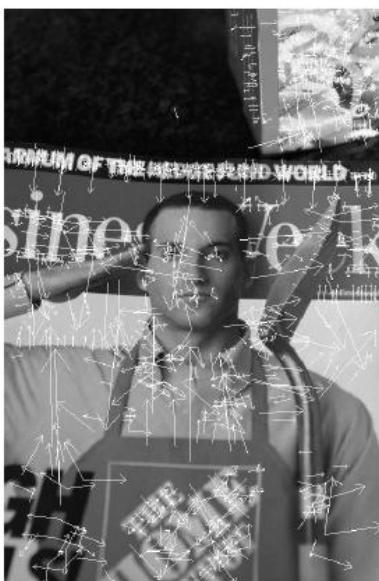
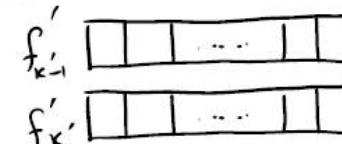
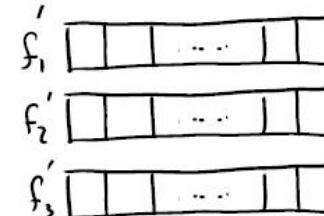


Image 2



② Build feature vectors

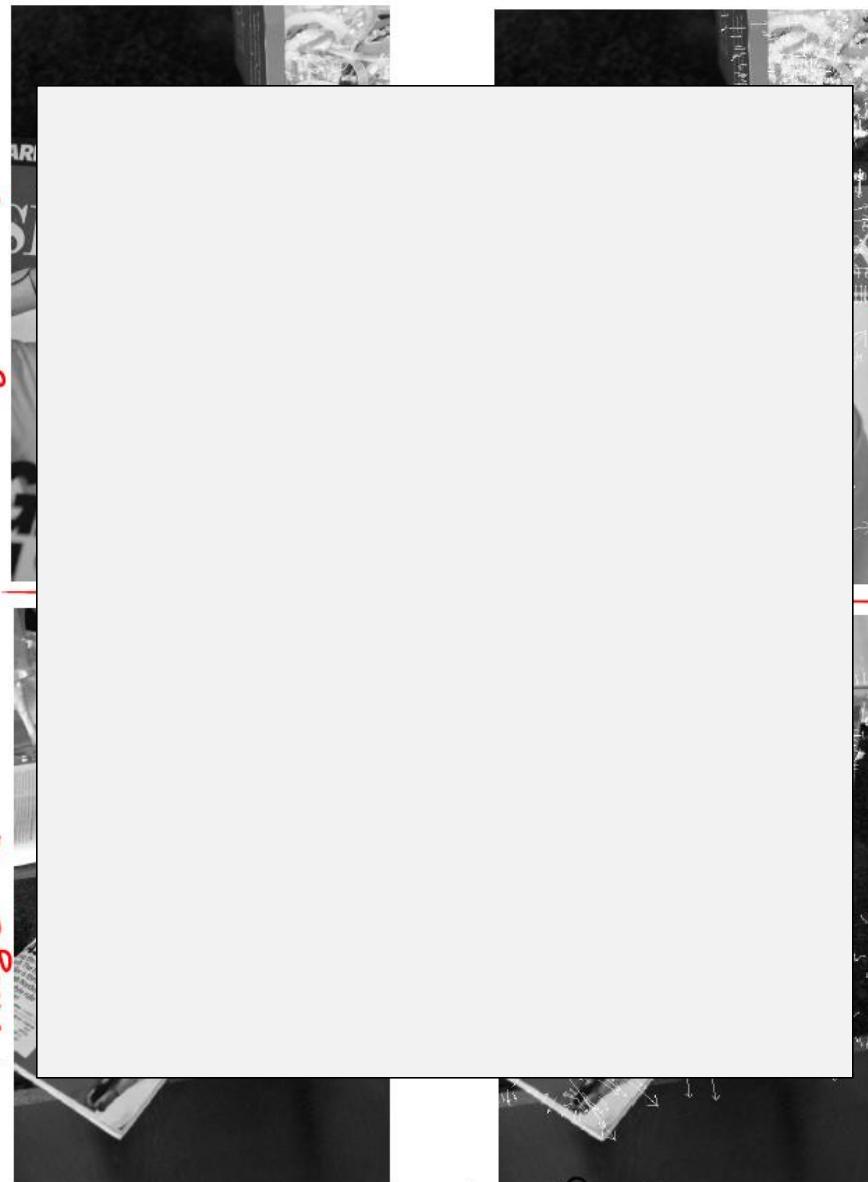
③ Match  
feature vectors  
in the two  
sets

$$\{f_1, f_2, \dots, f_k\}$$

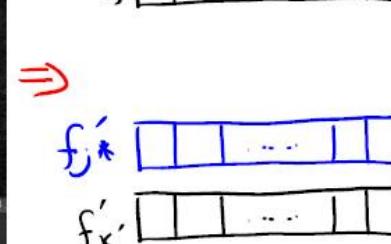
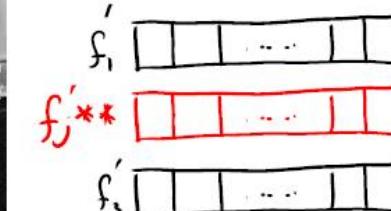
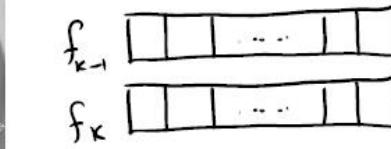
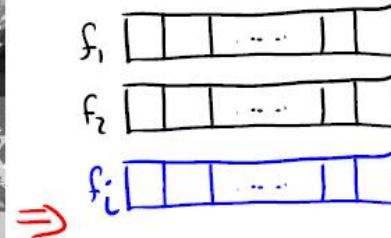
and

$$\{f'_1, f'_2, \dots, f'_{k'}\}$$

# SIFT Feature Matching Algorithm



① Identify keypoints



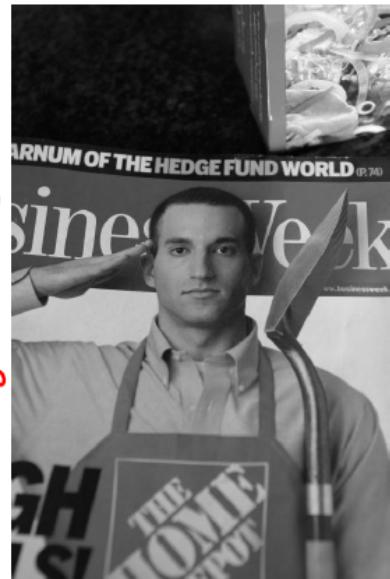
② Build feature vectors

③ Match  $f_i$

- Compute  $\|f_i - f_j'\|$  for all  $j$
- Compute fraction  $\phi = \frac{\|f_i - f_j^*\|}{\|f_i - f_{j''}\|}$  where  $f_j^*$  is closest descriptor in  $I'$  and  $f_{j''}$  is 2nd-closest
- Match  $f_i$  to  $f_j^*$  if  $\phi < 0.8$

# Matching 2 Images Using SIFT Features

I

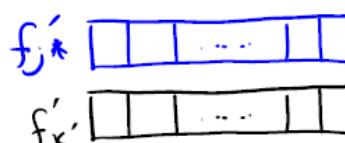
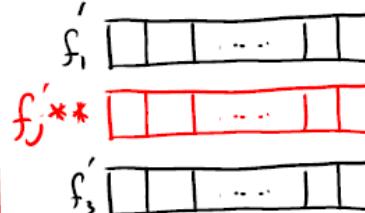
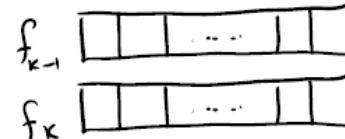
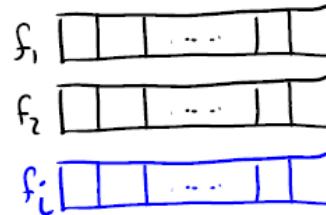


I'



Intuition for matching algorithm:  
match established only if it is deemed reliable, i.e. if there is only one very similar feature in image I'

① Identify keypoints



② Build feature vectors

③ Match  $f_i$

a. Compute  $\|f_i - f_j'\|$  for all  $j$

b. Compute fraction

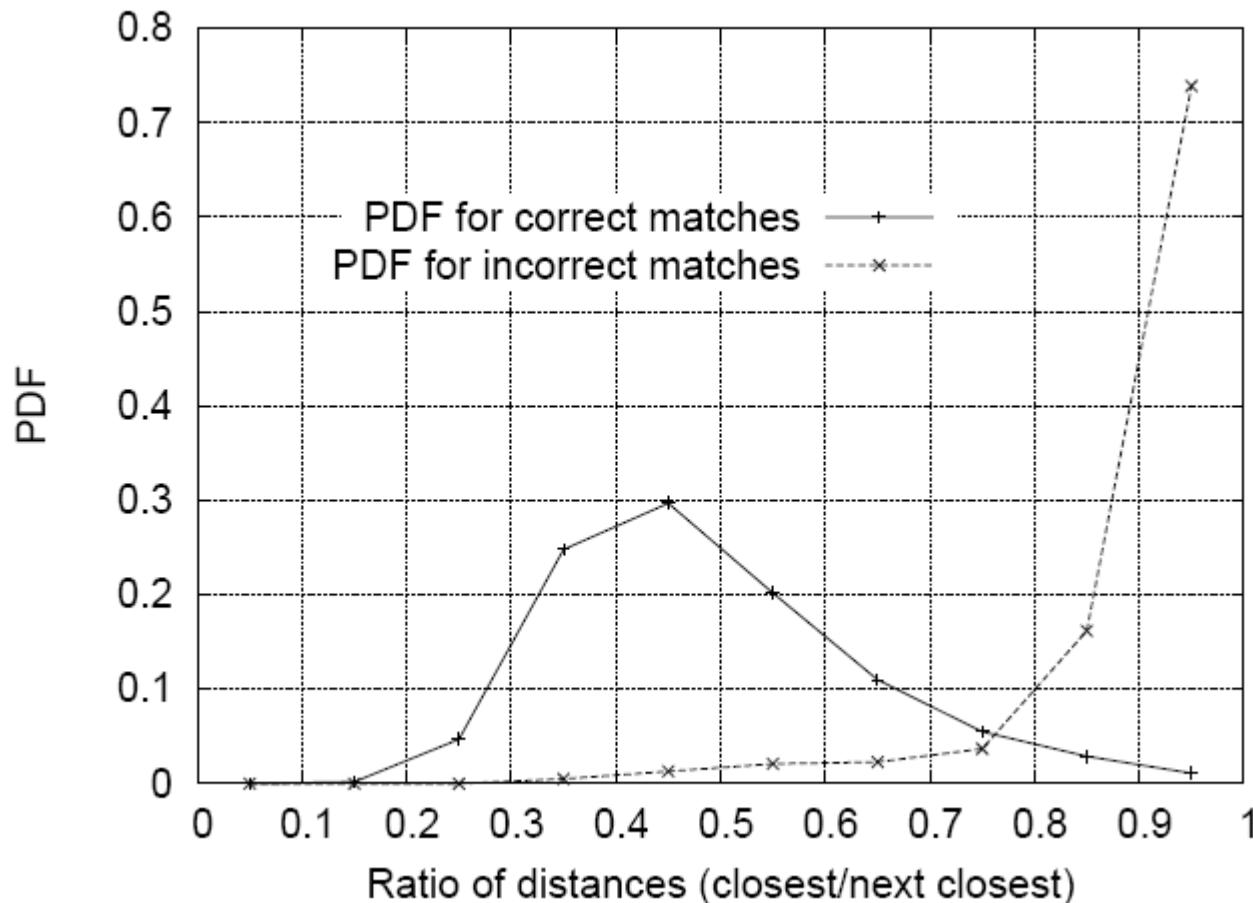
$$\phi = \frac{\|f_i - f_j^*\|}{\|f_i - f_j^{**}\|}$$

where  $f_j^*$  is closest descriptor in I' and  $f_j^{**}$  is 2nd-closest

c. Match  $f_i$  to  $f_j^*$  if  $\phi < 0.8$

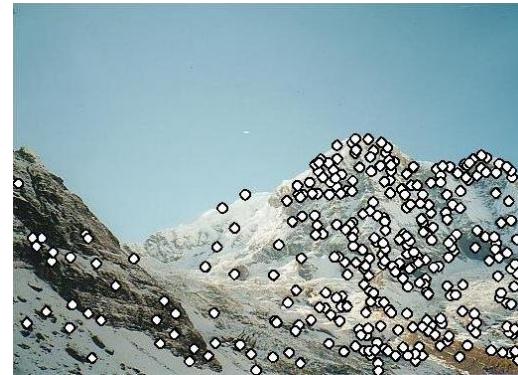
# Matching SIFT Descriptors

- Nearest neighbor (Euclidean distance)
- Threshold ratio of nearest to 2<sup>nd</sup> nearest descriptor



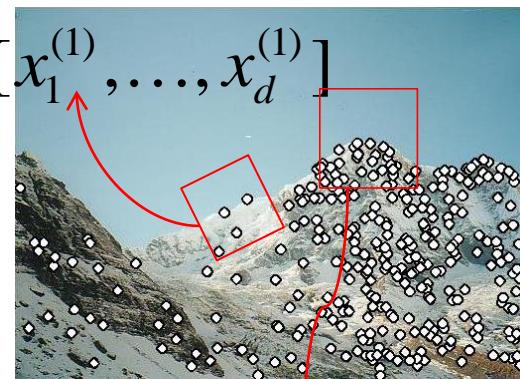
# Local features: main components

- 1) Detection: Identify the interest points



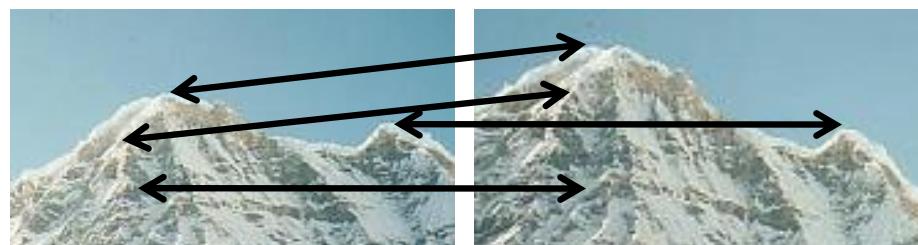
- 2) Description: Extract vector feature descriptor surrounding each interest point.

$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$



- 3) Matching: Determine correspondence between descriptors in two views

$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$



# Value of local (invariant) features

- Complexity reduction via selection of distinctive points
- Describe images, objects, parts without requiring segmentation
  - Local character means robustness to clutter, occlusion
- Robustness: similar descriptors in spite of noise, blur, etc.

# Applications of local invariant features

- Wide baseline stereo
- Motion tracking
- Panoramas
- Mobile robot navigation
- 3D reconstruction
- Recognition
- ...

# Summary

- Interest point detection
  - Harris corner detector
  - Laplacian of Gaussian, automatic scale selection
- Invariant descriptors
  - Rotation according to dominant gradient direction
  - Histograms for robustness to small shifts and translations (SIFT descriptor)